

---

# BAYESIAN INFERENCE WITH LATENT SPACE SAMPLING

NIKOLAJ TAKATA MÜCKE – [NIKOLAJ.MUCKE@CWI.NL](mailto:NIKOLAJ.MUCKE@CWI.NL)

POSTDOC AT CWI

14/11/2024

JOINT WORK WITH:

CORNELIS OOSTERLEE, SANDER BOHTÉ, AND BENJAMIN SANDERSE



# WHAT IS THE PROBLEM?

- Bayesian inference is great!
- But it is typically too slow for real-time problems
- ... Especially when problems are high-dimensional and non-Gaussian

# OVERVIEW

- Static Bayesian inference
  - Markov Chain Monte Carlo
  - Generative Adversarial Networks (GANs)
- Dynamic Bayesian inference (filtering)
  - Particle filter
  - Wasserstein Autoencoders

# STATIC BAYESIAN INFERENCE

# BAYESIAN INFERENCE

- **Framework** for solving inverse problems in a probabilistic setting
- **Result:** posterior distribution over parameters and/or state

$$\rho_{u|y}(\mathbf{u}|\mathbf{y}) = \frac{\rho_{y|u}(\mathbf{y}|\mathbf{u})\rho_0(\mathbf{u})}{\int_{\mathbb{R}^{N_u}} \rho_{y|u}(\mathbf{y}|\mathbf{u})\rho_0(\mathbf{u}) \, d\mathbf{u}}$$

Posterior

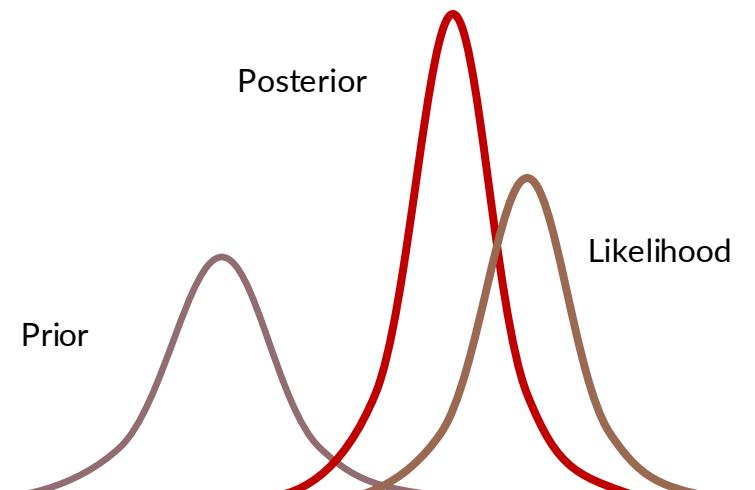
Likelihood

Prior

State + parameters

Observations

Evidence



# HOW DO WE COMPUTE THE POSTERIOR?

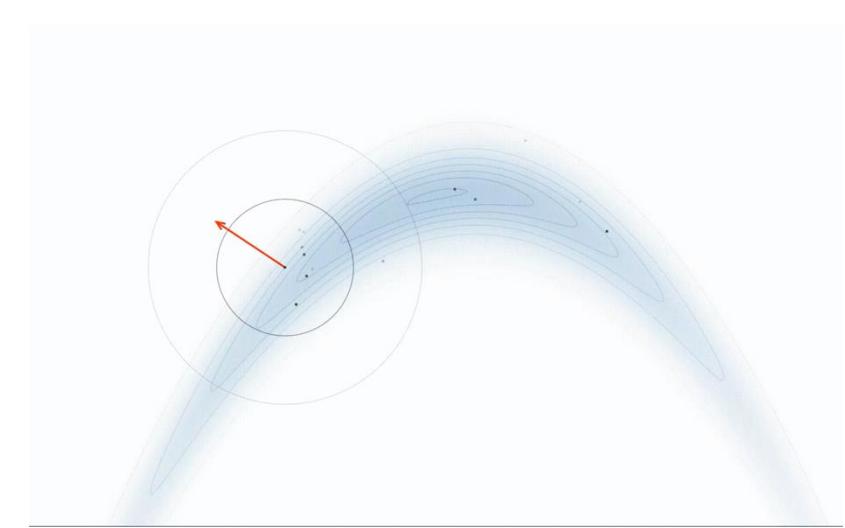
- Posterior is not analytically available
- We need numerical techniques!
- Sampling?

$$\rho_{u|y}(\mathbf{u}|\mathbf{y}) = \frac{\rho_{y|u}(\mathbf{y}|\mathbf{u})\rho_0(\mathbf{u})}{\int_{\mathbb{R}^{N_u}} \rho_{y|u}(\mathbf{y}|\mathbf{u})\rho_0(\mathbf{u}) \, d\mathbf{u}}$$

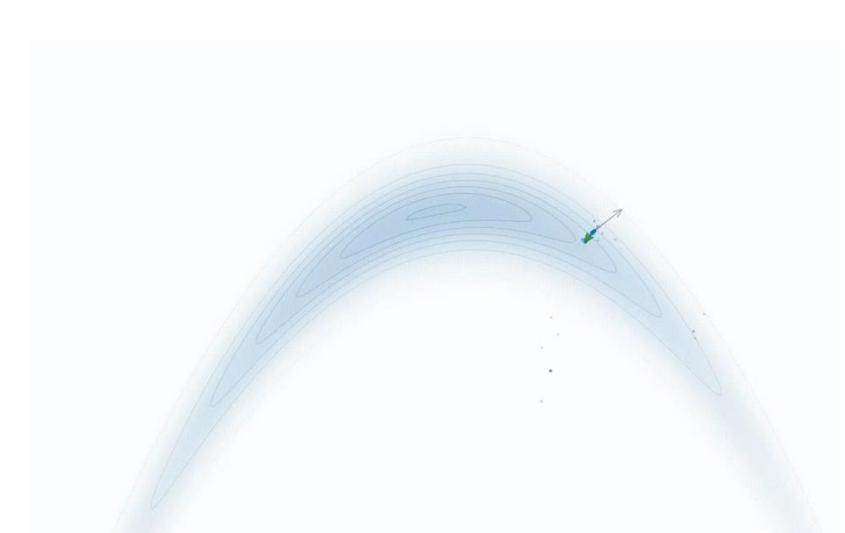
# MARKOV CHAIN MONTE CARLO (MCMC)

- ✗ Sample from posterior distribution by creating Markov chain
- ✗ Accept/reject proposed samples
- ✗ Converges weakly towards true distribution

Metropolis-Hastings



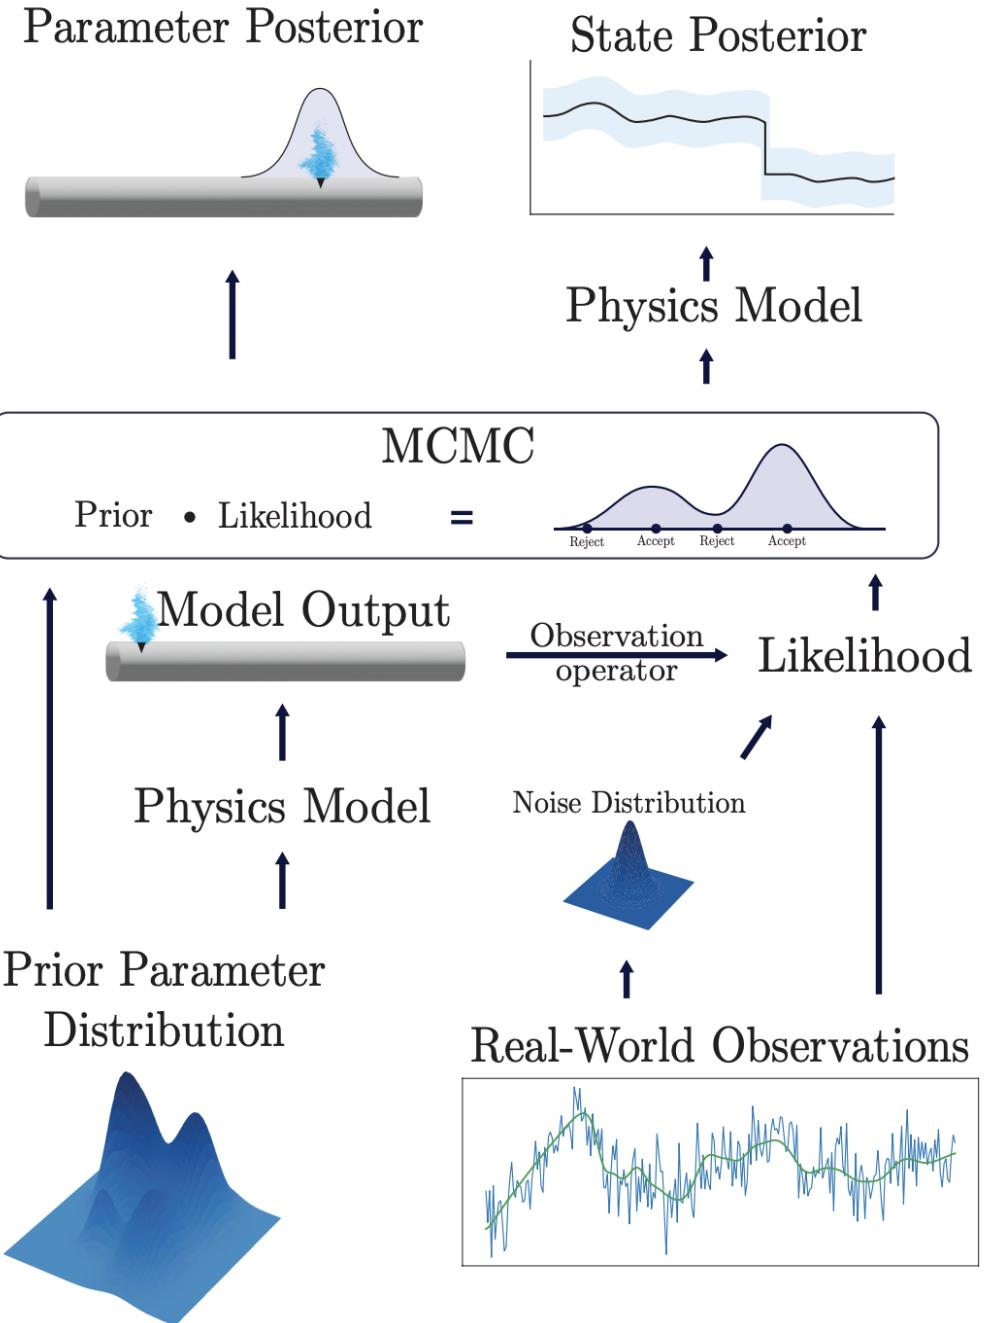
HMC with NUTS



# BAYESIAN INFERENCE WITH MCMC

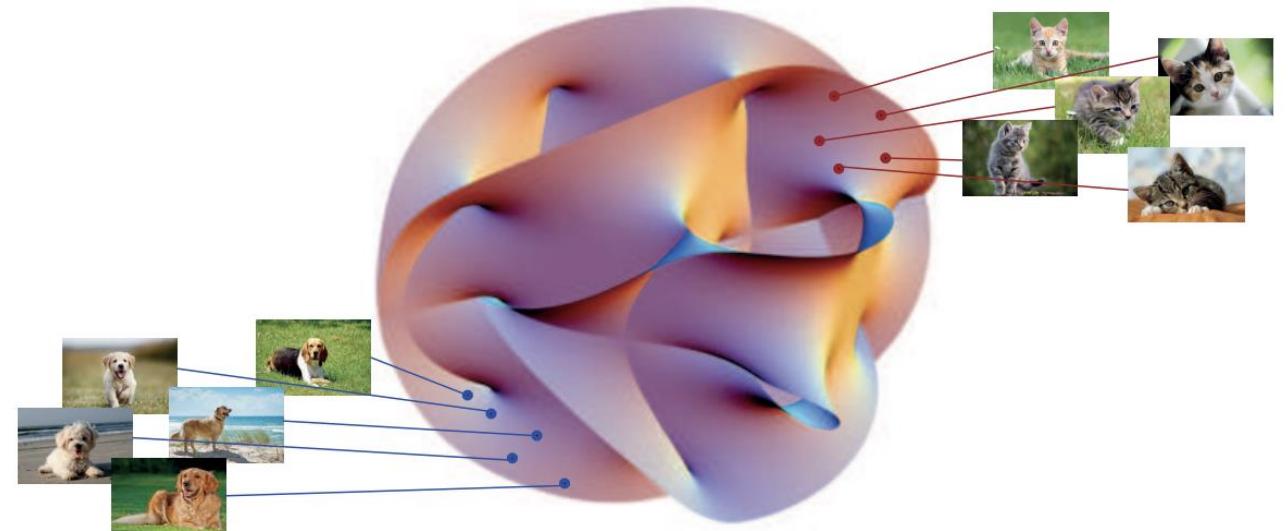
1. Sample from prior distribution
2. Evaluate forward model
3. Compute likelihood
4. Accept/reject with MCMC

What is a good prior?  
Expensive!  
Requires many samples



# GENERATIVE DEEP LEARNING?

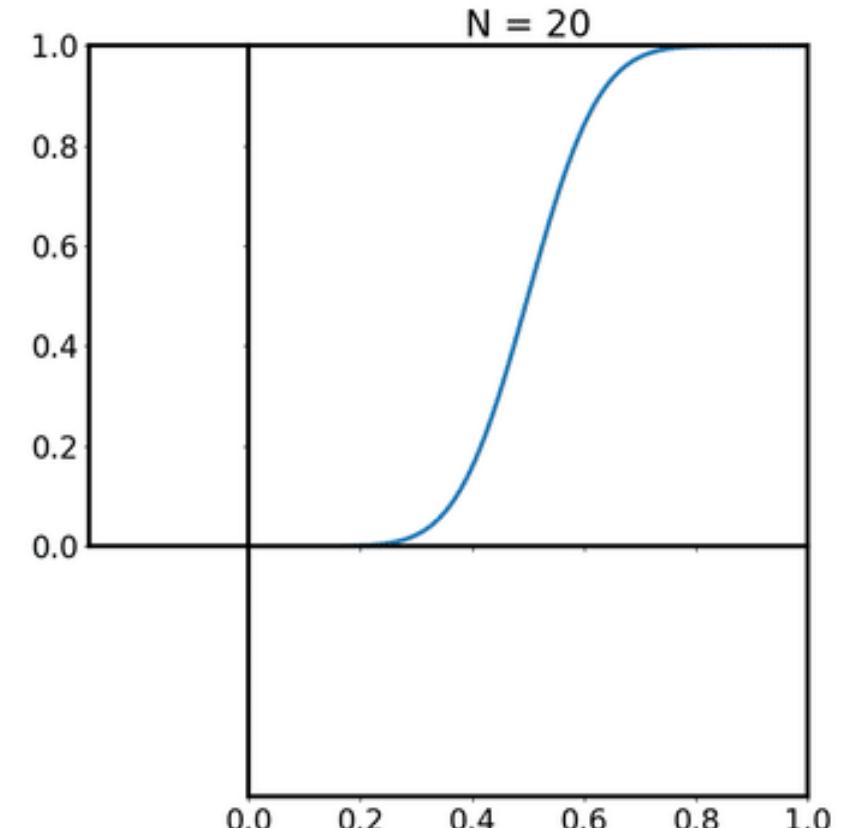
- Efficient sampling of unknown distribution
- Dimensionality reduction



# GENERATING DISTRIBUTIONS VIA PUSH-FORWARD MAP

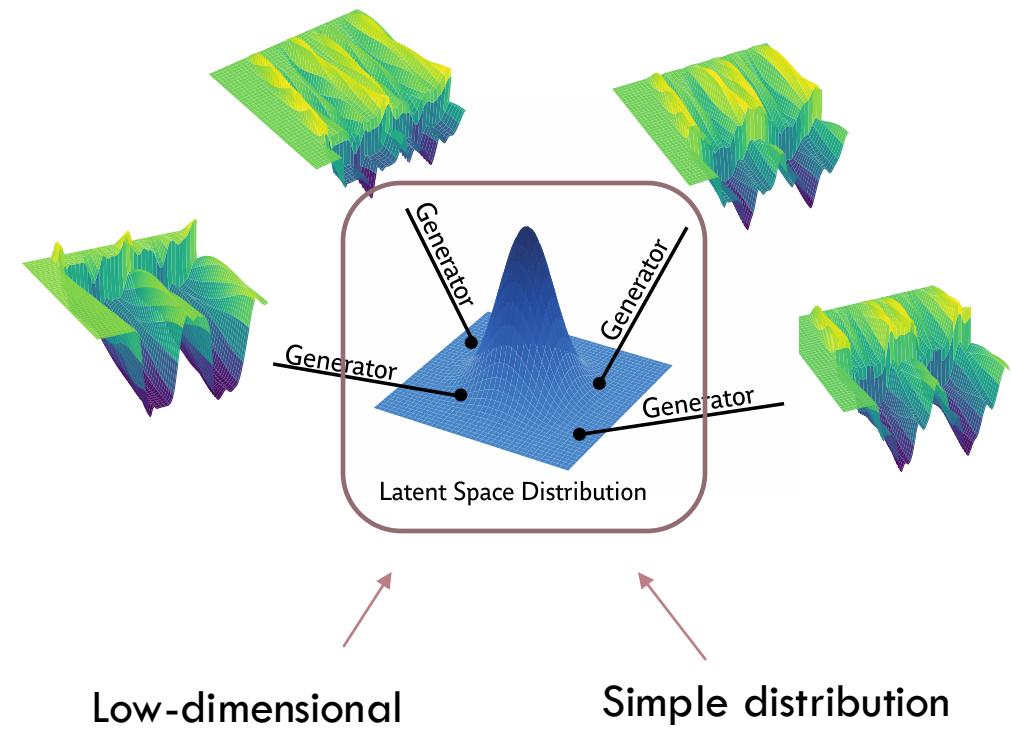
1. Sample from simple distribution
2. Pass sample through function (push-forward map)
  - Generator or decoder
3. Get data sample

Inverse transform sampling



# GENERATING DISTRIBUTIONS VIA PUSH-FORWARD MAP

1. Sample from latent distribution
2. Pass sample through function (push-forward map)
  - Generator or decoder
3. Get data sample

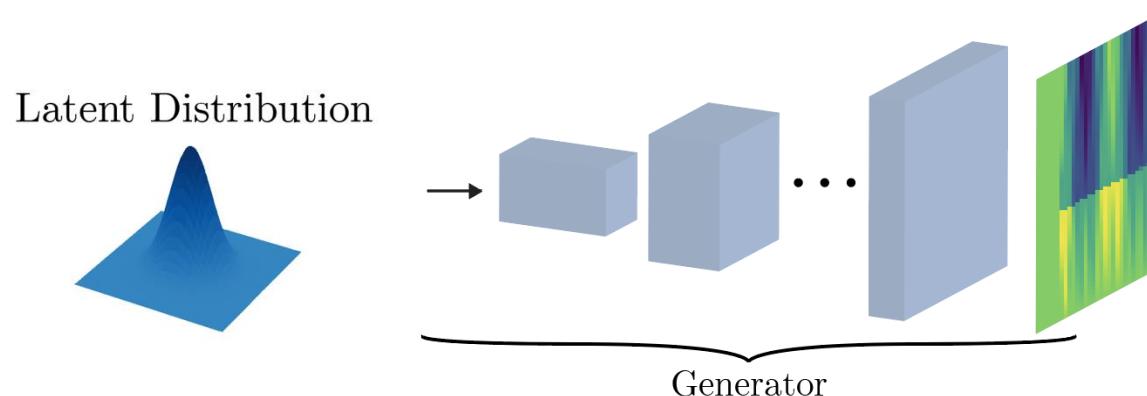


# GENERATOR DISTRIBUTION

- Generator maps latent samples to high-fidelity samples:  $U = G_\theta(Z)$
- Generator pushes forward the latent distribution:

$$U \sim P_u^g = G_\theta \# P_z^g$$

$$\rho_u^g = \rho_z^g \circ G^{-1}$$



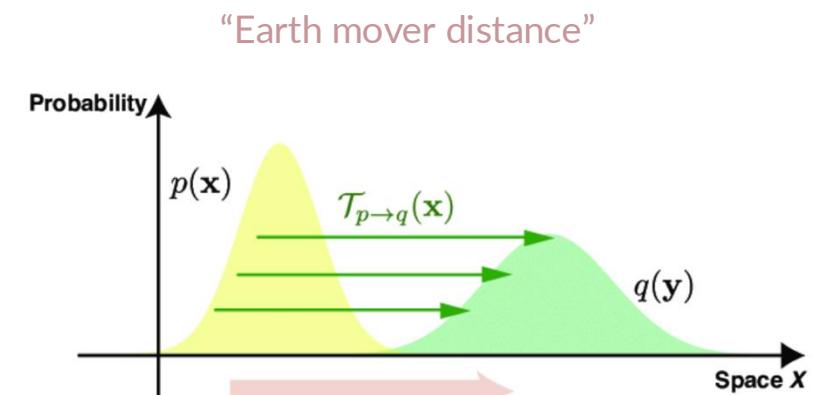
WASSERSTEIN  
GANS

# WASSERSTEIN DISTANCE

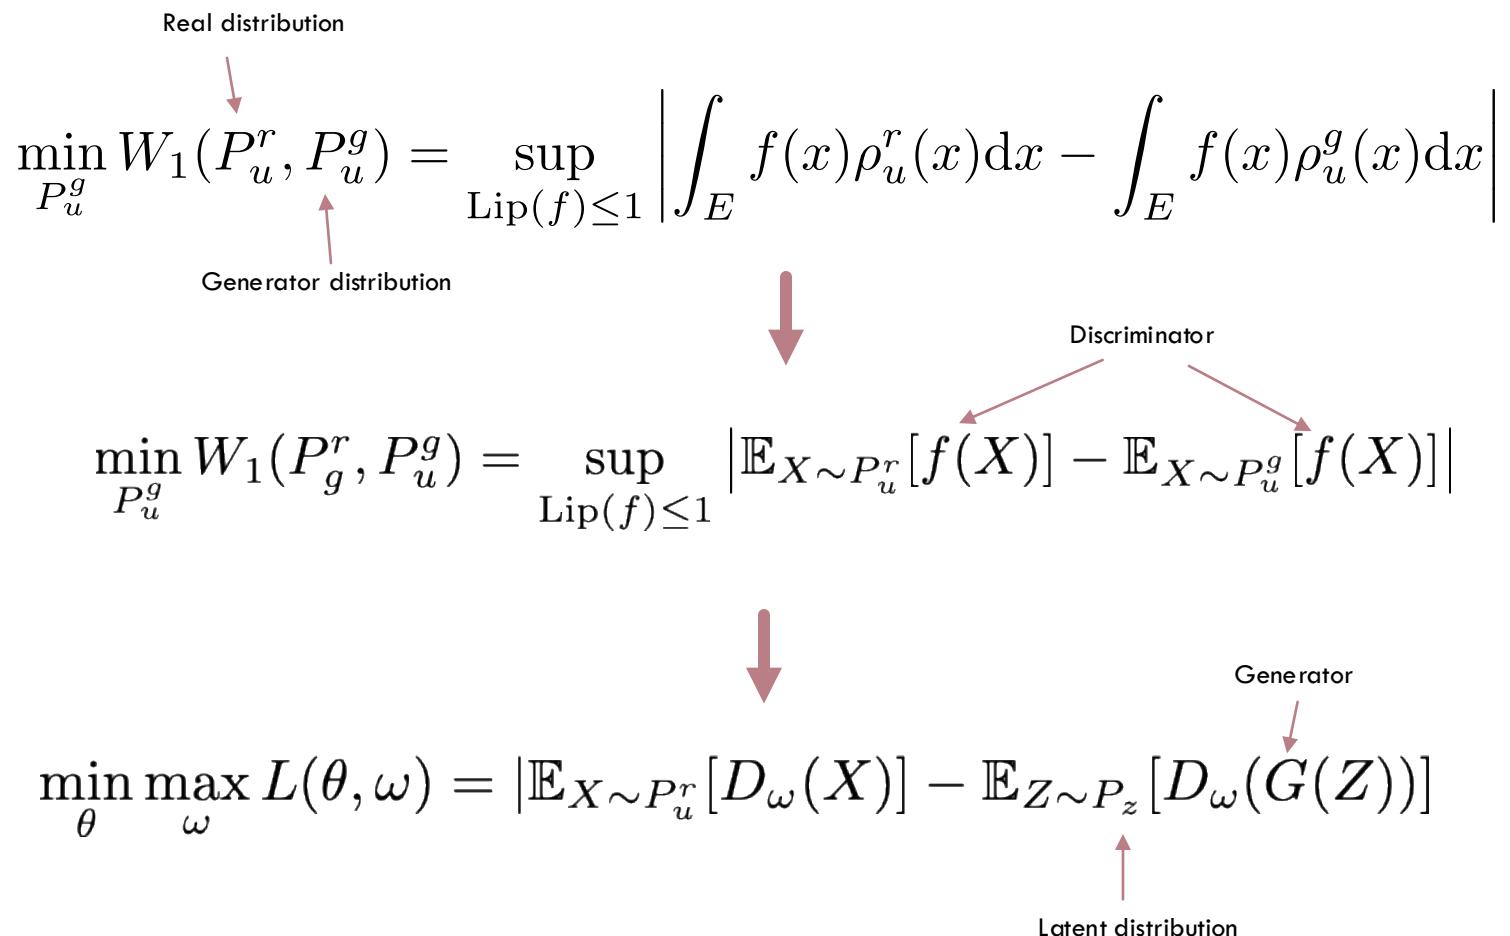
$$W_1(P_1, P_2) = \inf_{\gamma \in \Gamma(P_1, P_2)} \left| \int_E \int_E d_E(x, y) \gamma(x, y) dx dy \right|$$

↓  
Kantorovich–Rubinstein duality

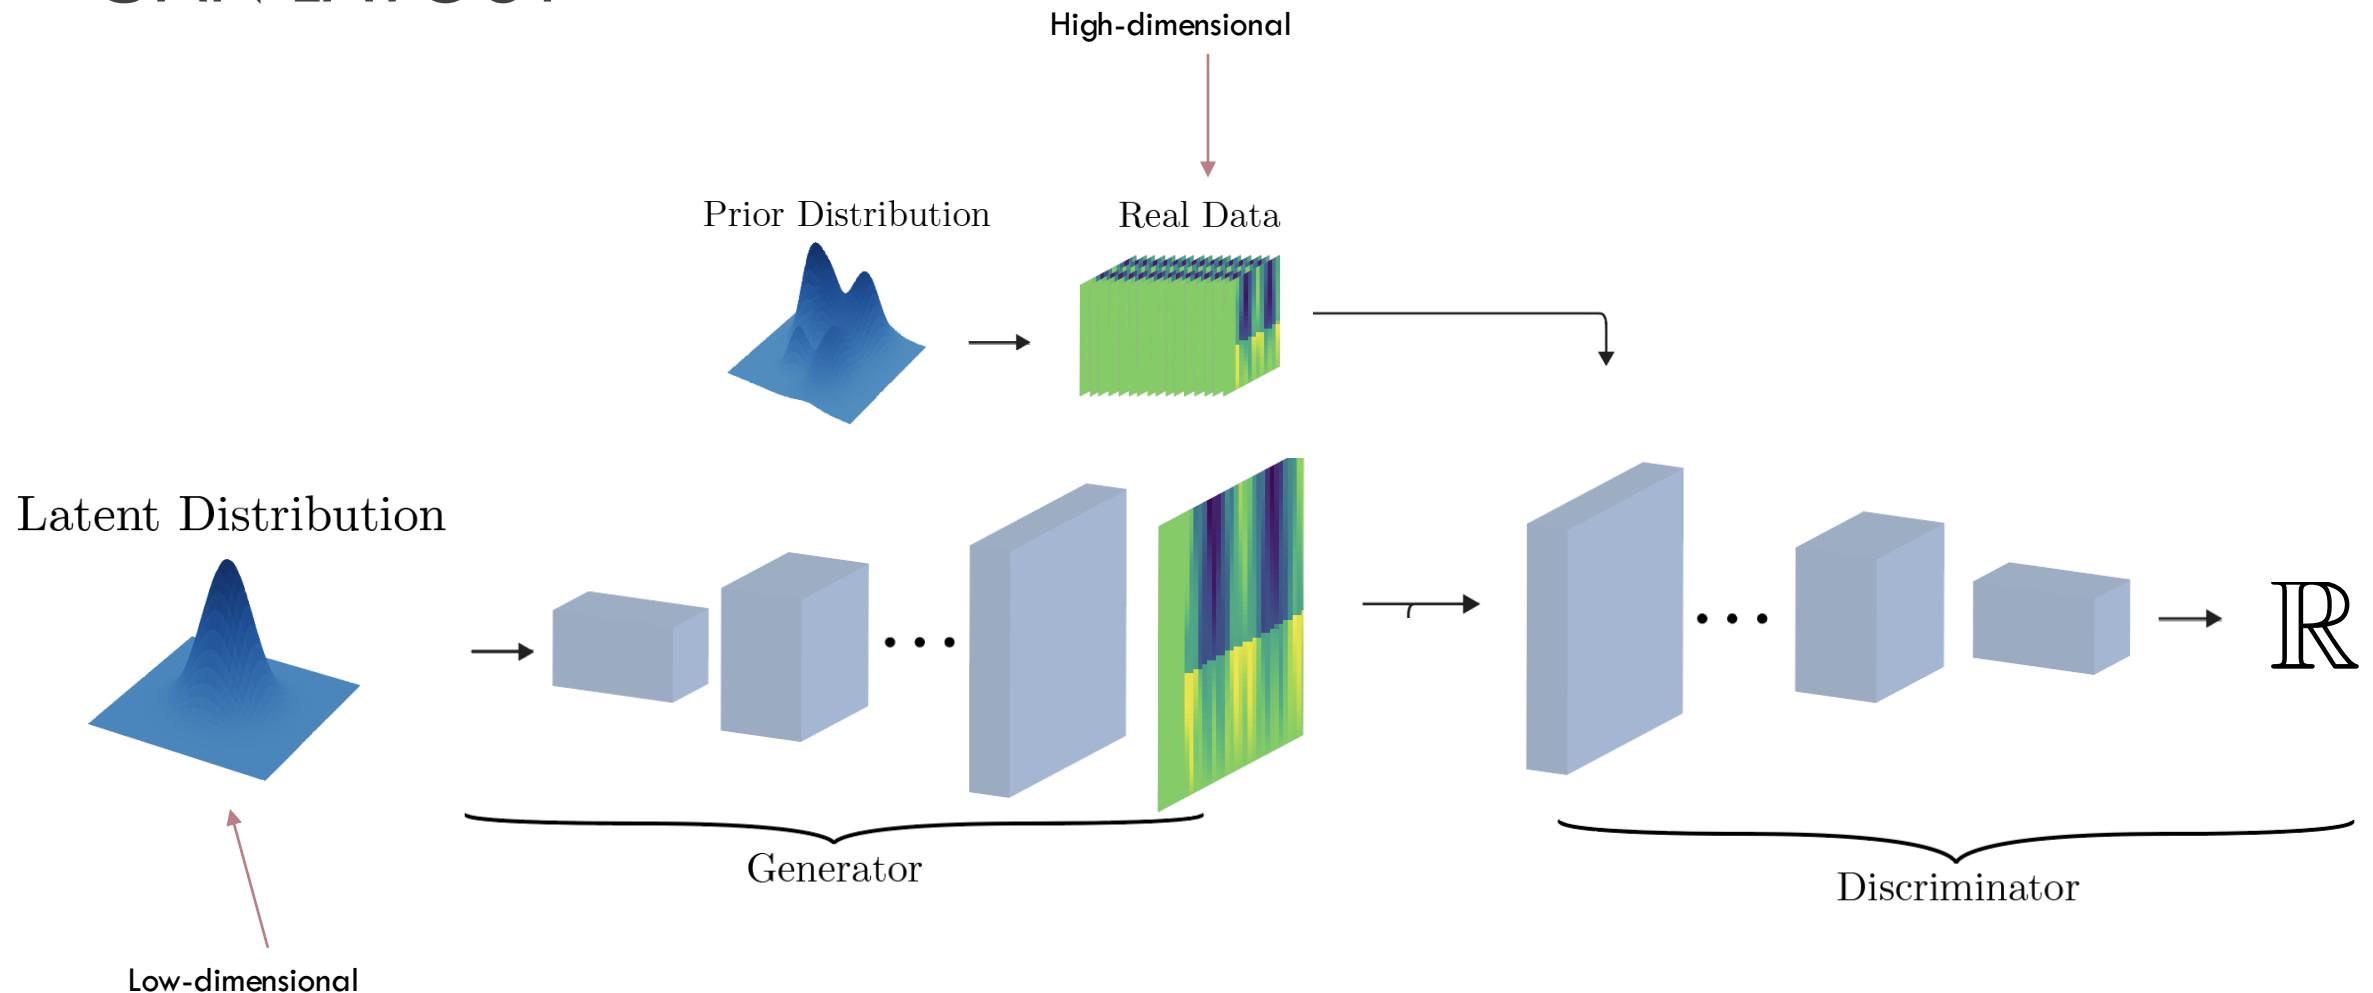
$$W_1(P_1, P_2) = \sup_{\text{Lip}(f) \leq 1} \left| \int_E f(x) \rho_1(x) dx - \int_E f(x) \rho_2(x) dx \right|$$



# TRAINING OBJECTIVE



# GAN LAYOUT



# LATENT SPACE POSTERIOR

# LATENT SPACE POSTERIOR

- Train GAN on prior samples of parameters and/or state
- Replace prior distribution with latent distribution
- Compute posterior in latent space!
  - Maximum a-posteriori is now (almost) free
  - Sample Markov chain in latent space

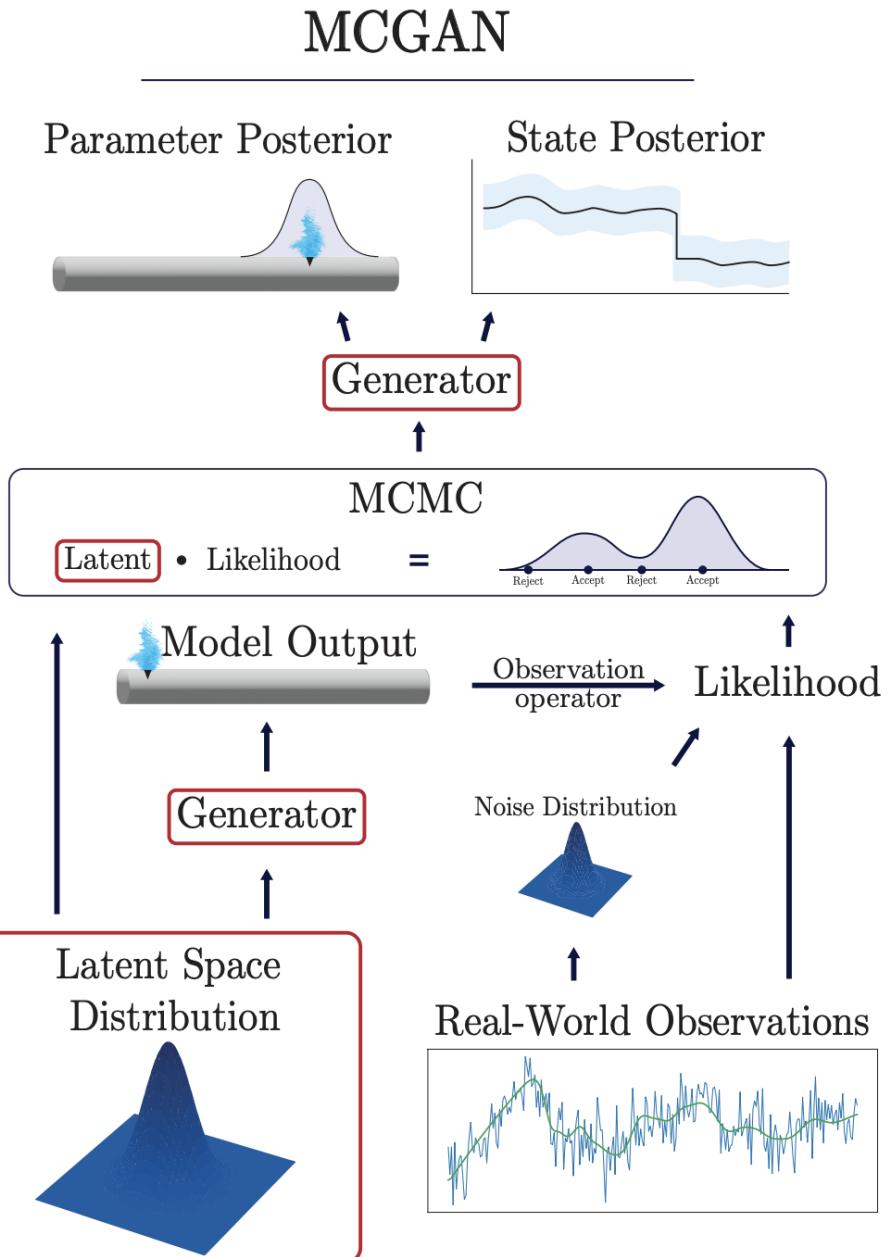
$$\rho_{u|y}(\mathbf{u}|\mathbf{y}) = \frac{\rho_{y|u}(\mathbf{y}|\mathbf{u})\rho_0(\mathbf{u})}{\int_{\mathbb{R}^{N_u}} \rho_{y|u}(\mathbf{y}|\mathbf{u})\rho_0(\mathbf{u}) \, d\mathbf{u}}$$



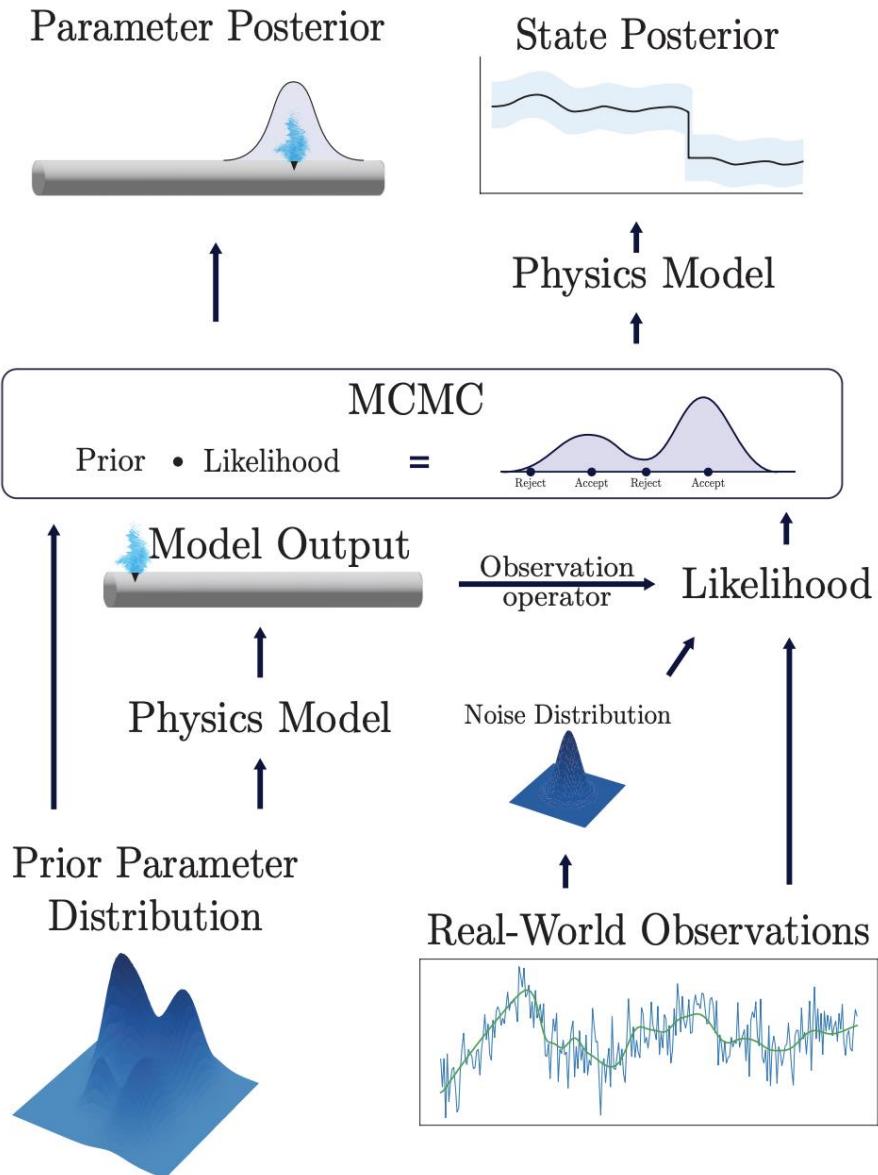
$$\rho_{z|y}^g(\mathbf{z}|\mathbf{y}) = \frac{\rho_{y|u}^g(\mathbf{y}|G_\theta(\mathbf{z}))\rho_z^g(\mathbf{z})}{\int_{\mathbb{R}^{N_z}} \rho_{y|u}^g(\mathbf{y}|G_\theta(\mathbf{z}))\rho_z^g(\mathbf{z}) \, d\mathbf{z}}$$

# MARKOV CHAIN GAN (MCGAN)

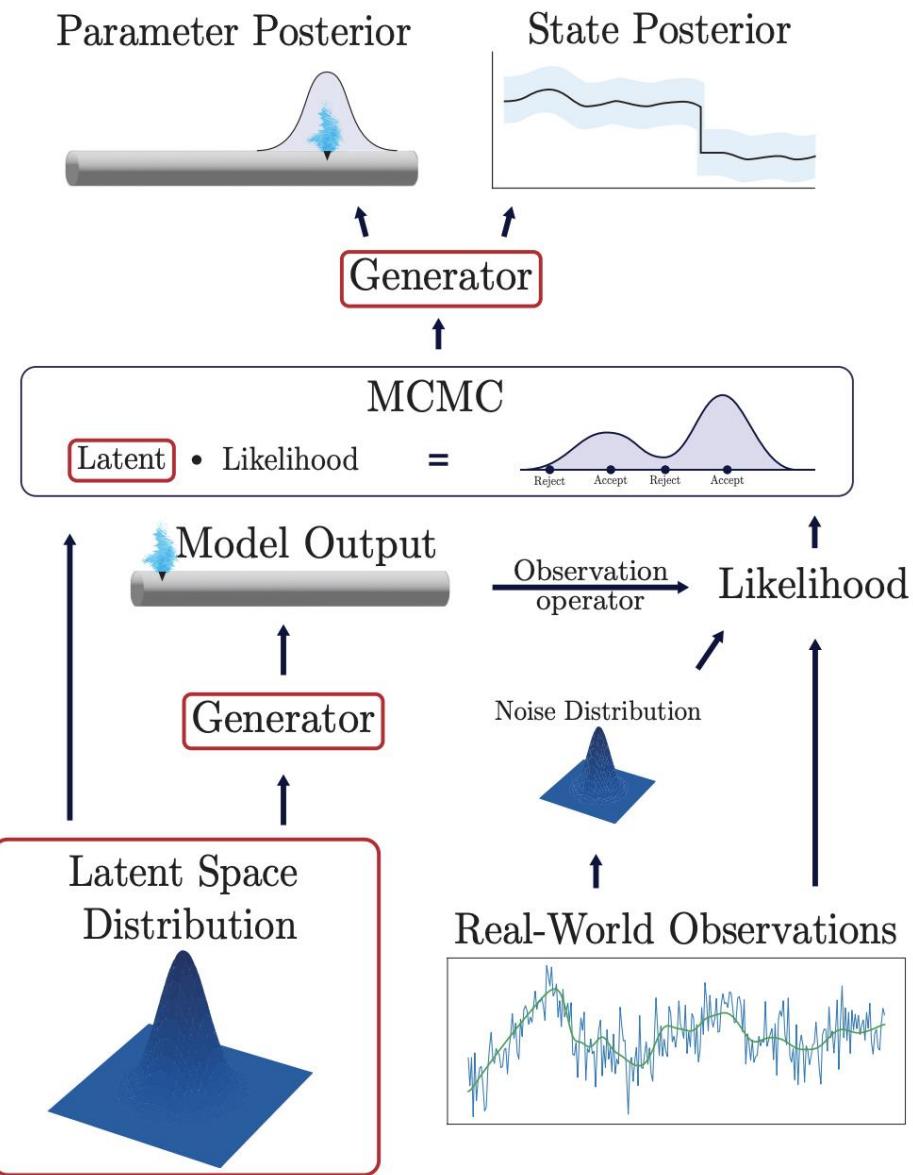
1. Sample from latent distribution
2. Push forward using generator
3. Compute likelihood
4. Accept/reject with MCMC
5. Push forward accepted latent samples using generator



# Bayesian Inference with MCMC



# MCGAN



# LATENT SPACE SAMPLING

- Sampling latent posterior and pushing forward is weakly the same as sampling high-fidelity posterior

$$\mathbb{E}_{U \sim P_{u|y}^g} [f(U)] = \mathbb{E}_{Z \sim P_{z|y}^g} [f(G(Z))]$$

- Change of variable formula

$$\mathbb{E}_{U \sim P_u} [f(U)] = \int_E f(\mathbf{u}) \rho_u(\mathbf{u}) d\mathbf{u} = \int_{G^{-1}(E)} f(G(\mathbf{z})) \rho_z(\mathbf{z}) d\mathbf{z} = \mathbb{E}_{Z \sim P_z} [f(G(Z))]$$

# CONVERGENCE OF POSTERIOR

- GAN converges to true prior in Wasserstein 1 metric  $\rightarrow$  Posterior also converges

$$W_1(P_{u|y}^r, P_{u|y}^g) \leq C_1 W_1(P_0^r, P_0^r) + C_2 \|l^r(\mathbf{u}) - l^g(\mathbf{u})\|_{L_{\rho_0^g}^1} + C_3 \|l^r(\mathbf{u}) - l^g(\mathbf{u})\|_{L_{\rho_0^g}^2}$$

$$C_1 = \frac{(1 + D\text{Lip}(\Phi^r))}{Q_u^r(\mathbf{y})}, \quad C_2 = \frac{\max(\Phi^r, \Phi^g)}{Q_u^r(\mathbf{y})Q_u^g(\mathbf{y})} (1 + D\text{Lip}(\Phi^r)) |P_0^g|_{\mathcal{W}_1}, \quad C_3 = \frac{\max(\Phi^r, \Phi^g)}{Q_u^g(\mathbf{y})} |P_0^g|_{\mathcal{W}_2}$$

# DARCY FLOW

$$\mathbf{v} + k \nabla p = 0,$$

$$\nabla \cdot \mathbf{v} = 0,$$

$$p = 1,$$

$$p = 0,$$

$$\mathbf{v} \cdot \mathbf{n} = 0, \quad \mathbf{x} \in [0, 1] \times \{0, 1\}.$$

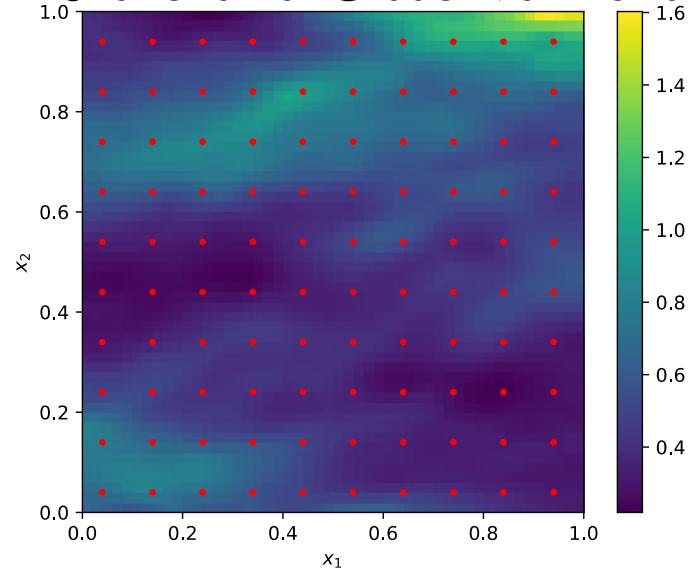
$$\mathbf{x} \in [0, 1]^2,$$

$$\mathbf{x} \in [0, 1]^2,$$

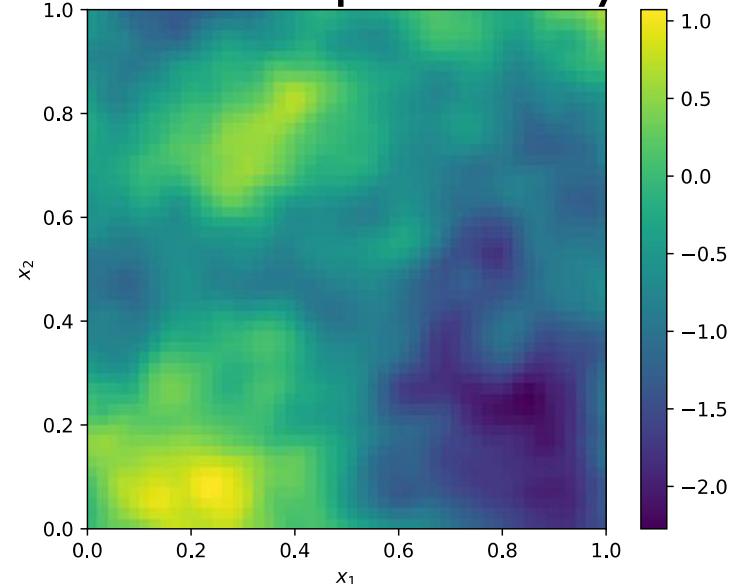
$$\mathbf{x} \in 0 \times [0, 1],$$

$$\mathbf{x} \in 1 \times [0, 1],$$

State and Observartions

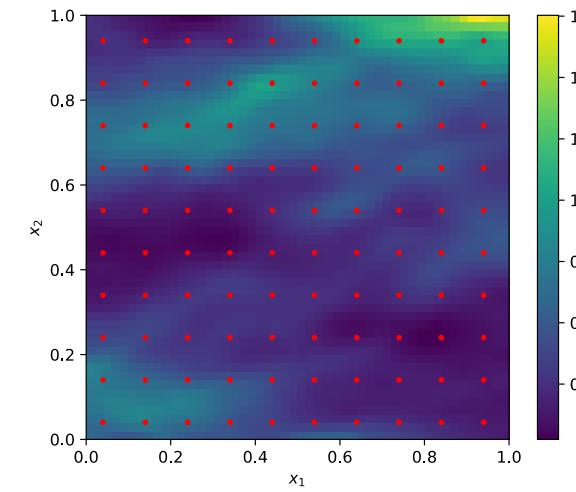


Parameters – permeability field

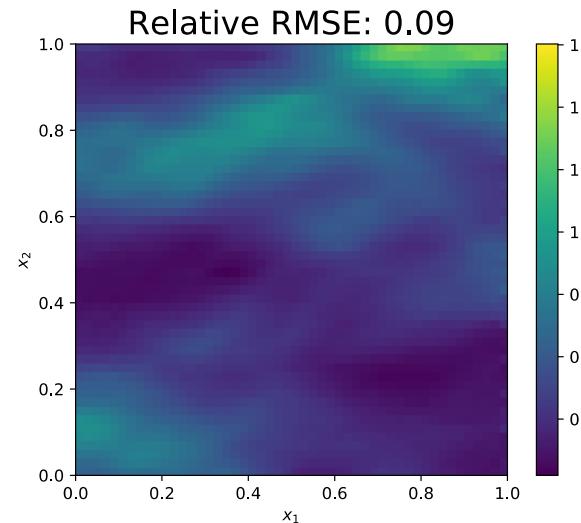


# RECONSTRUCTION

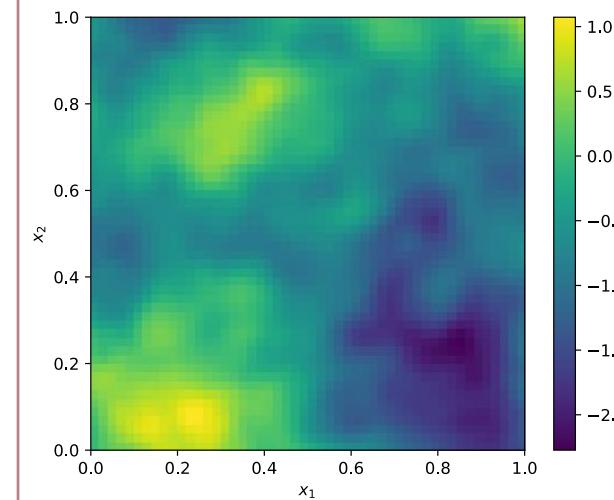
True State



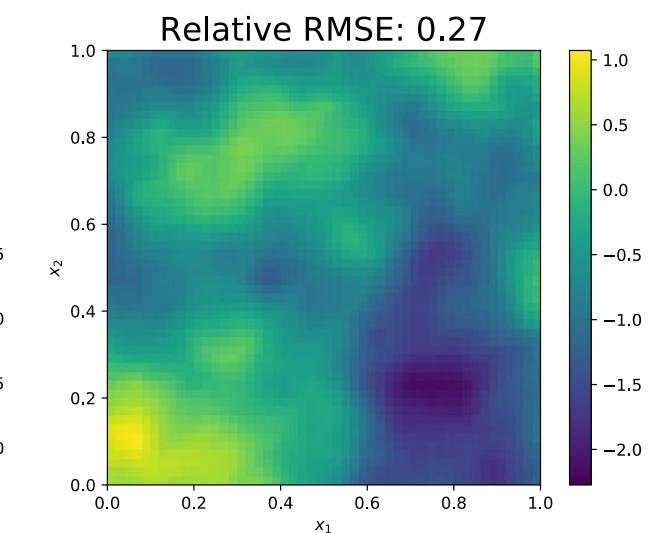
MCGAN



True Parameters



MCGAN Parameters

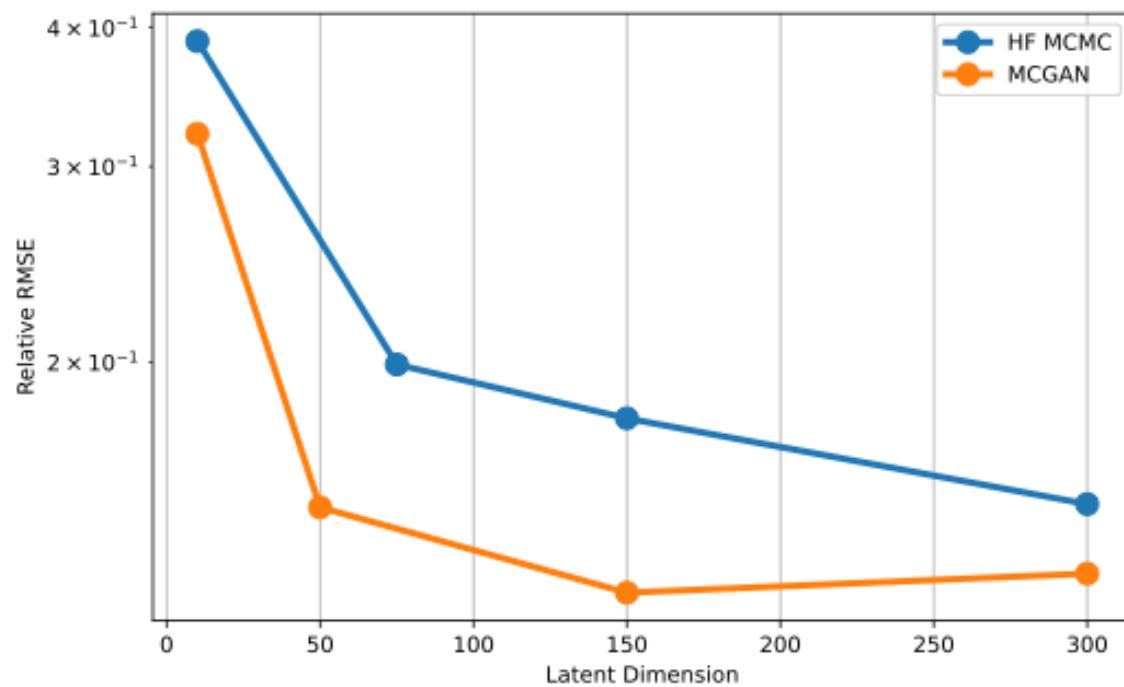


$\dim(u) = 25,793$

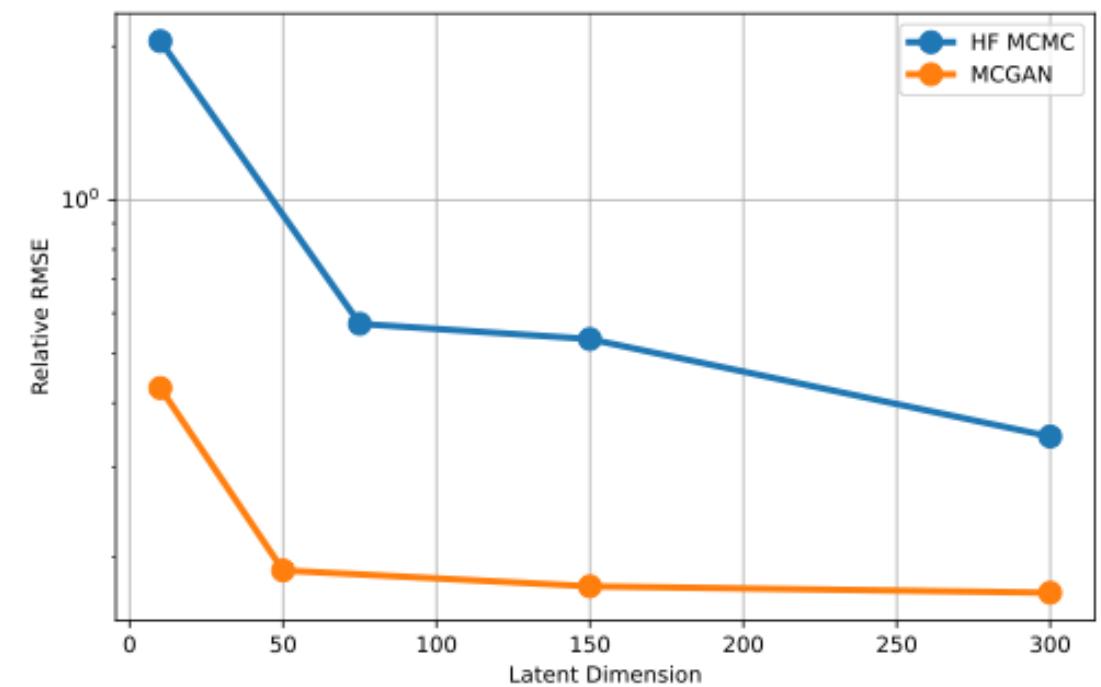
$\dim(z) = 150$

# COMPARISON WITH KARHUNEN-LOEVE EXPANSION

State

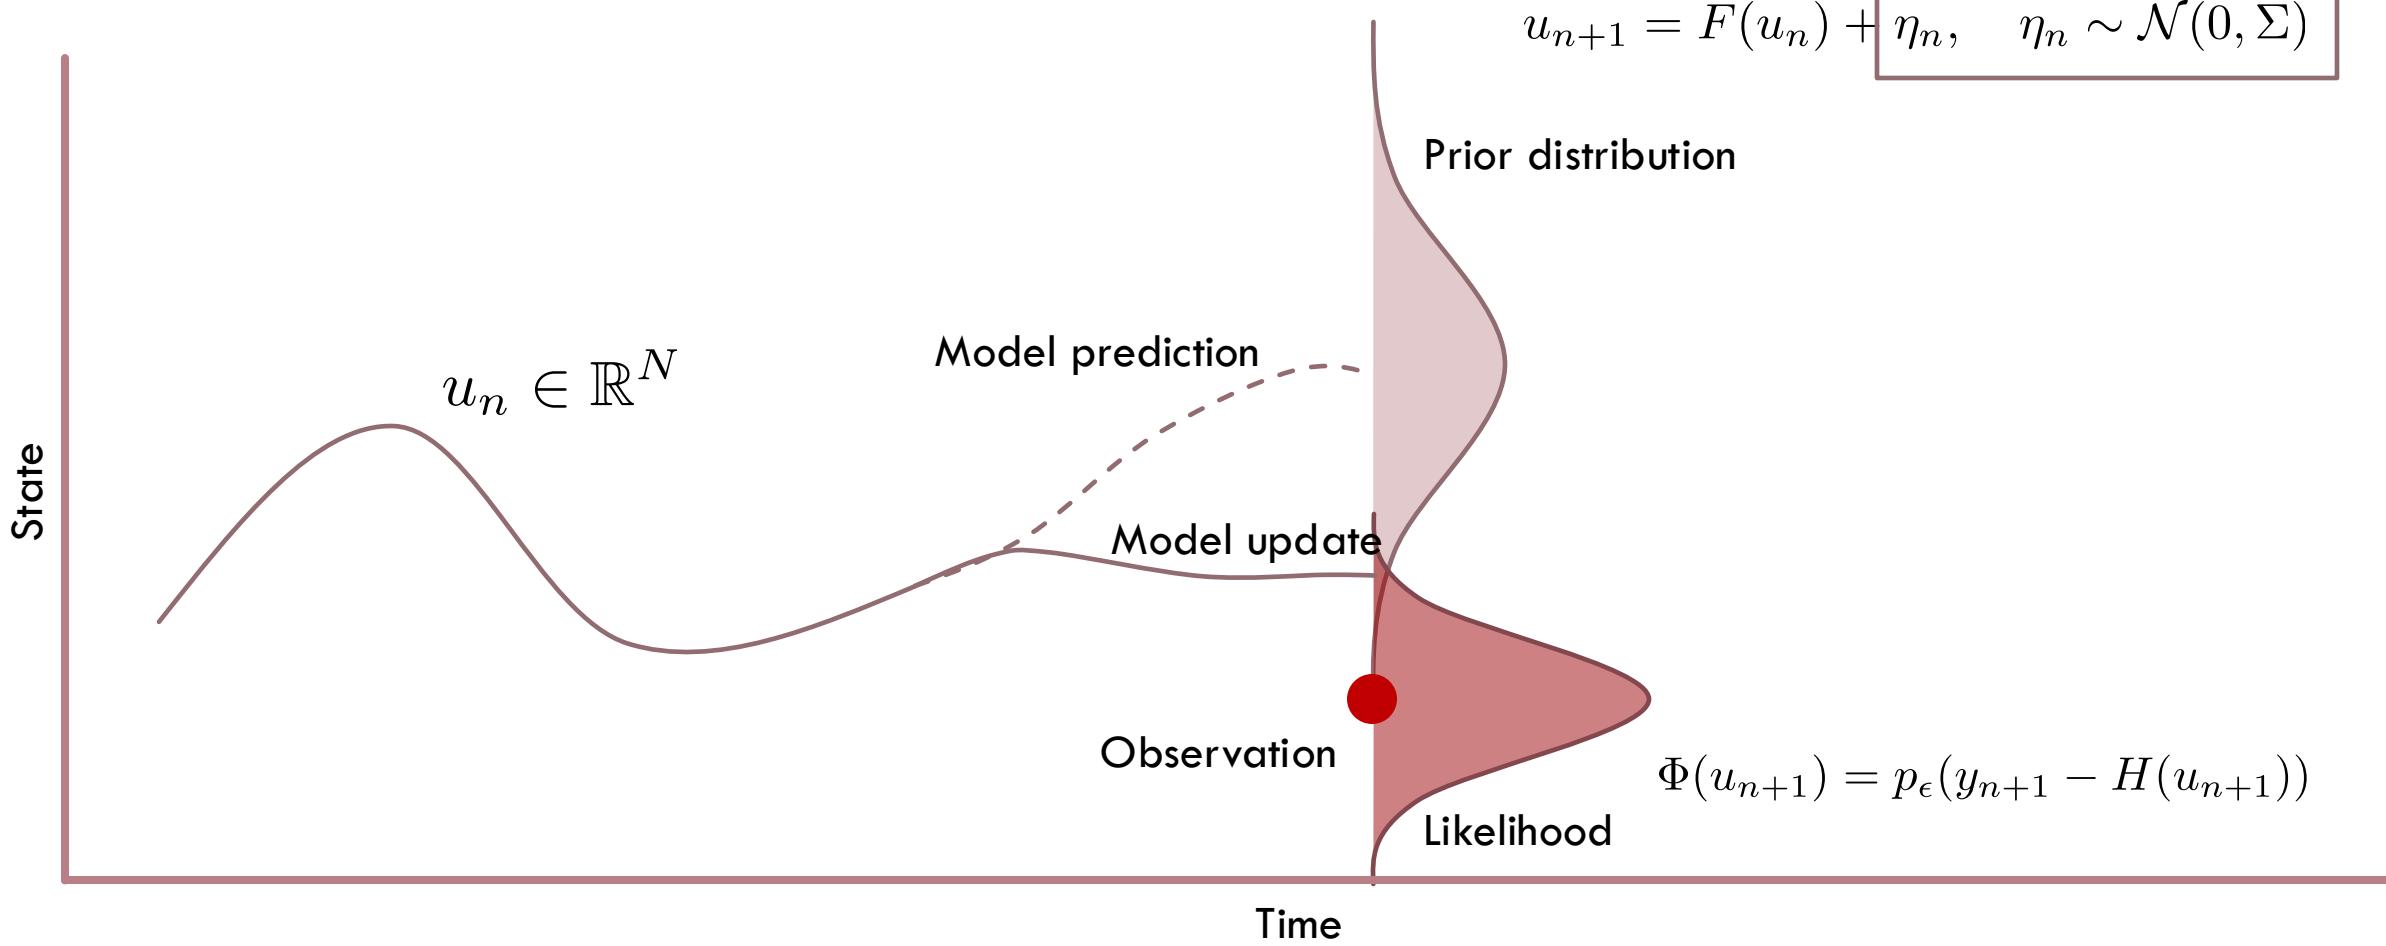


Parameters

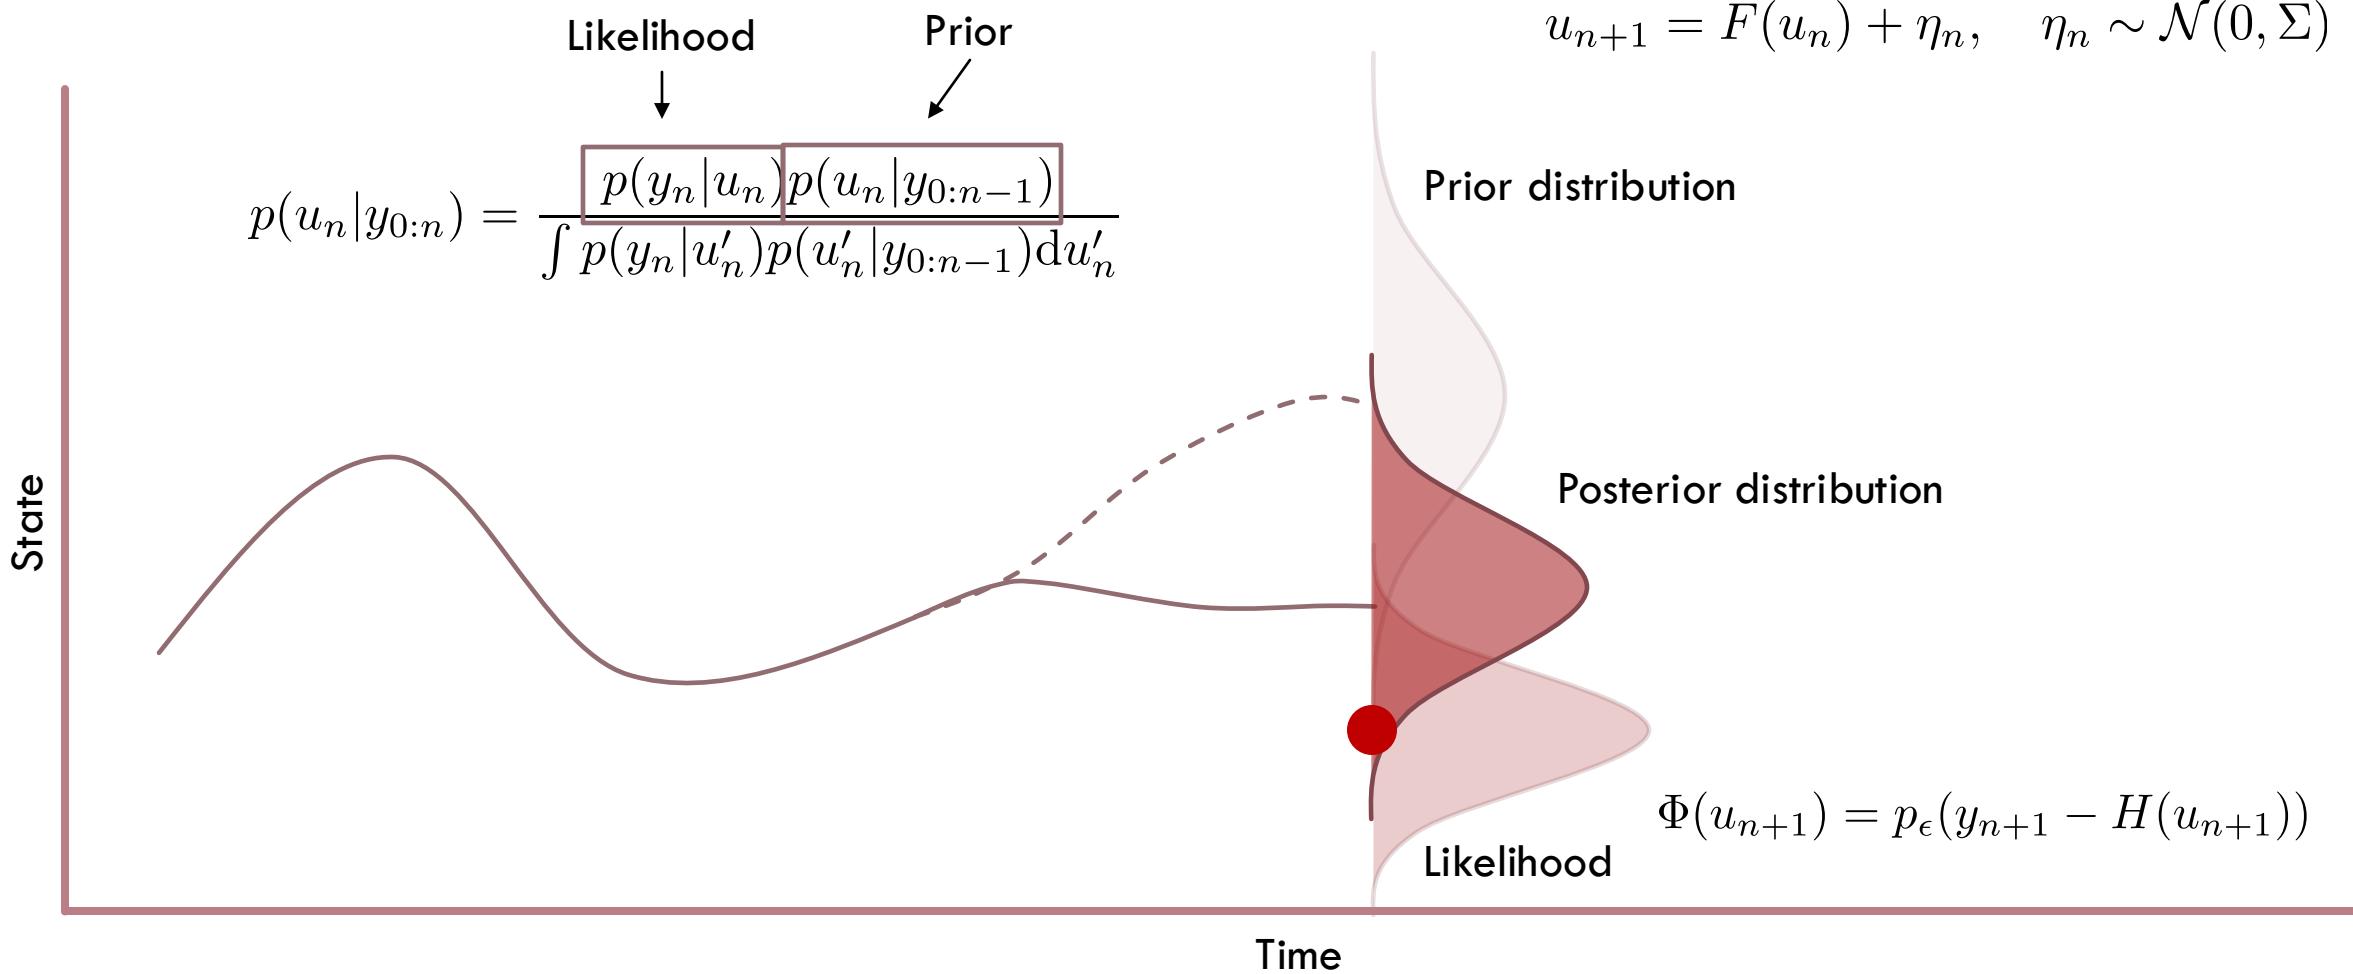


# DYNAMIC BAYESIAN INFERENCE

# BAYESIAN DATA ASSIMILATION

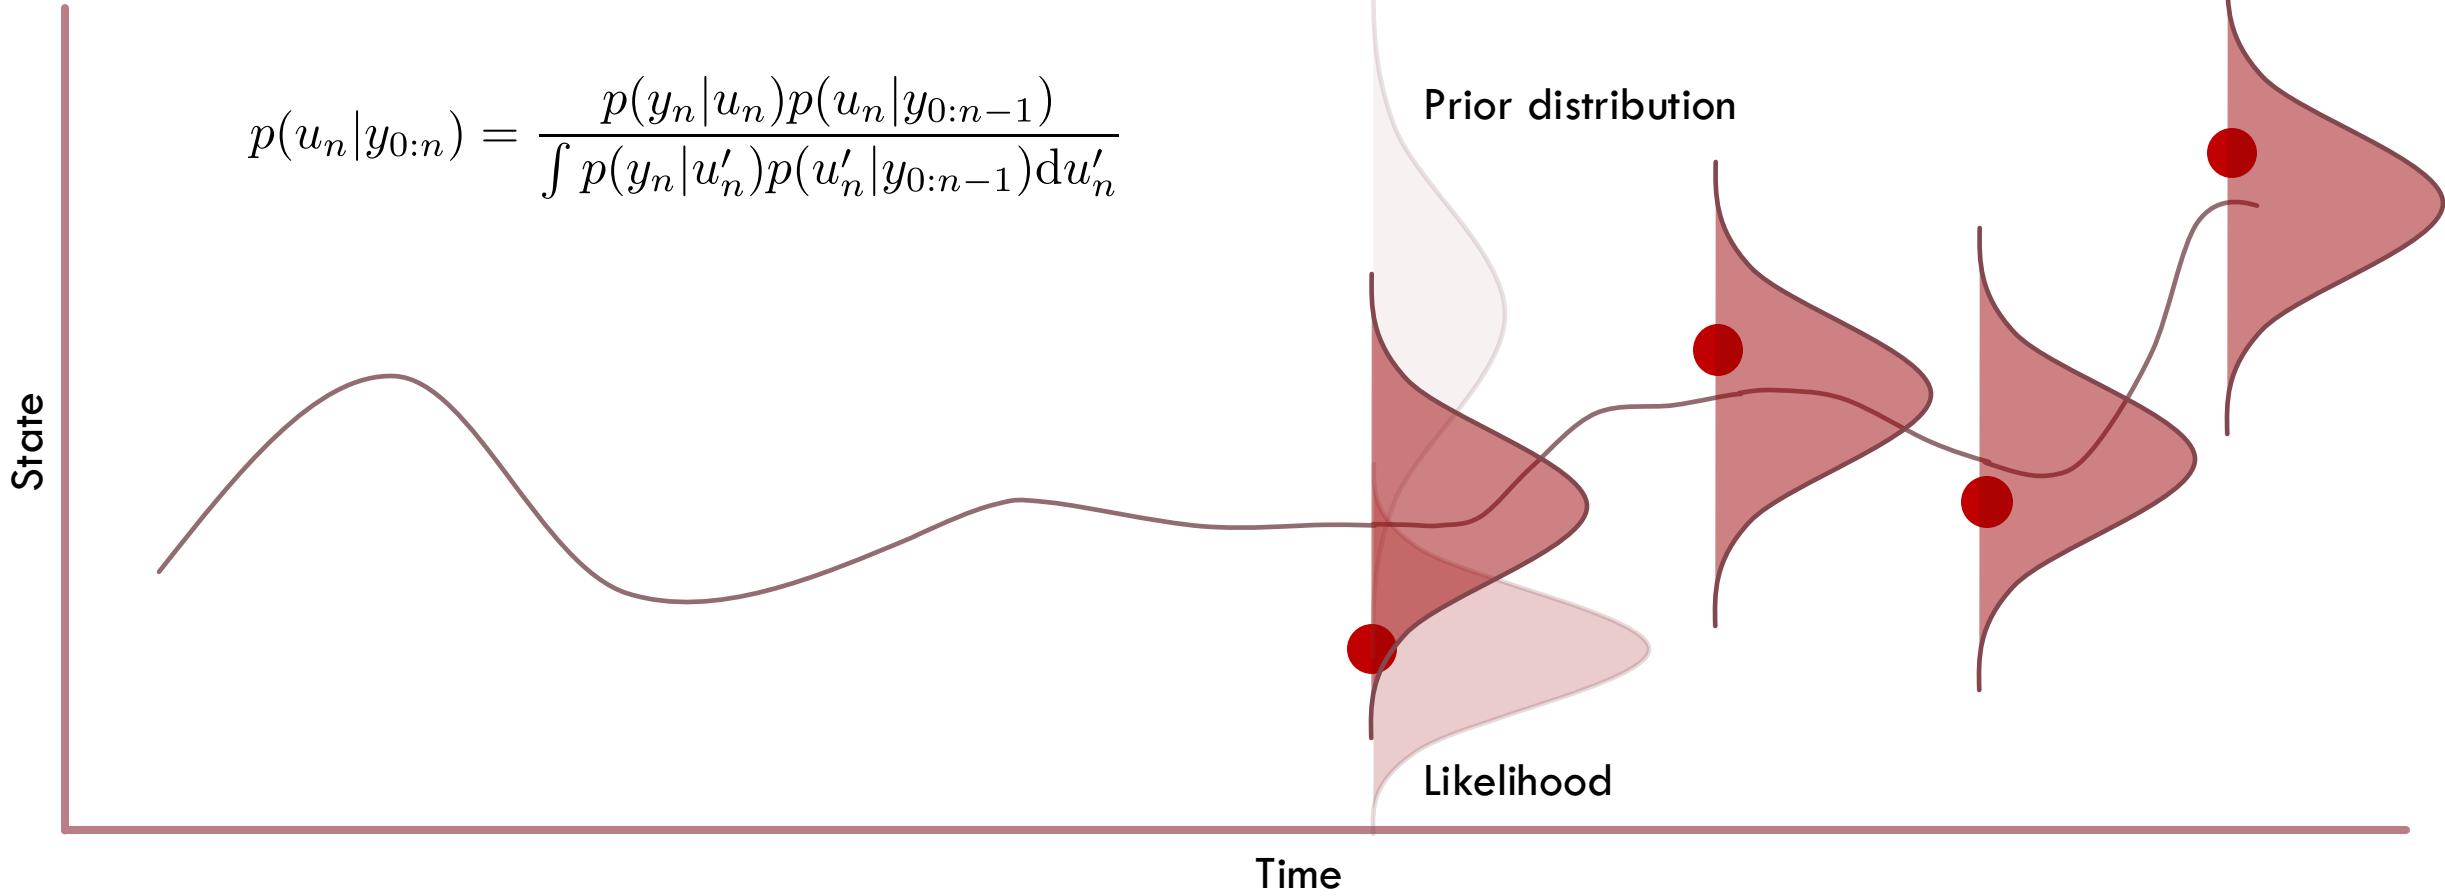


# BAYESIAN DATA ASSIMILATION

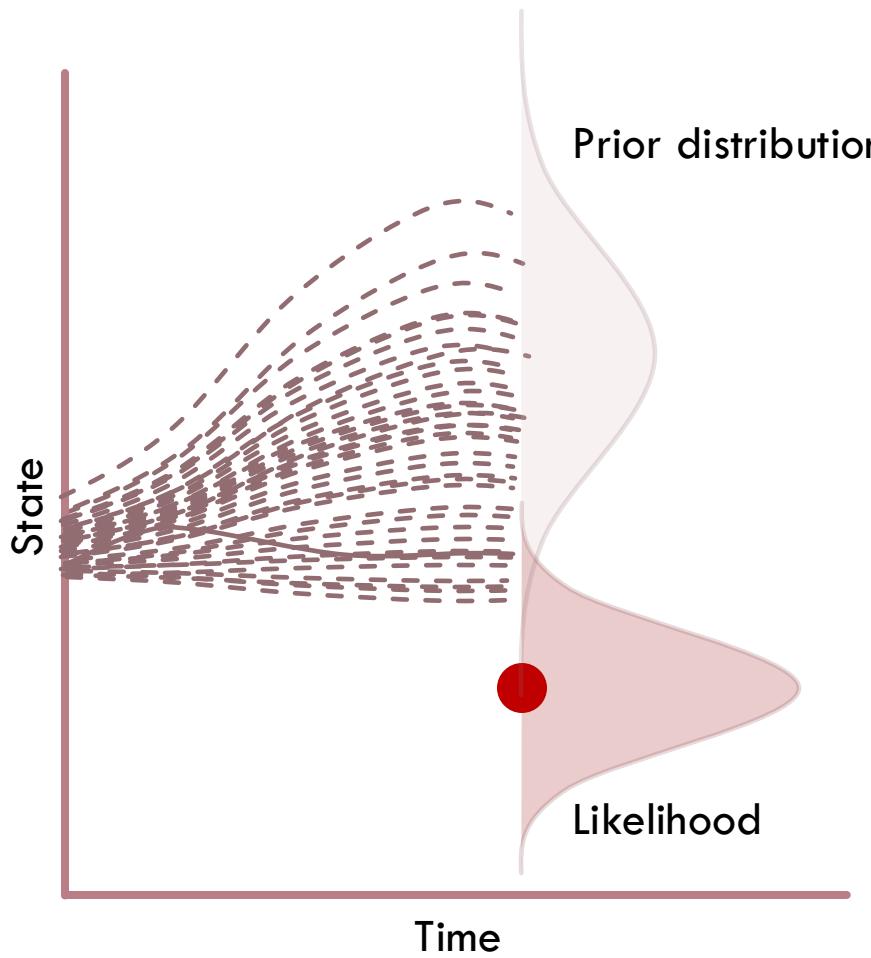


# BAYESIAN DATA ASSIMILATION

$$p(u_n|y_{0:n}) = \frac{p(y_n|u_n)p(u_n|y_{0:n-1})}{\int p(y_n|u'_n)p(u'_n|y_{0:n-1})du'_n}$$



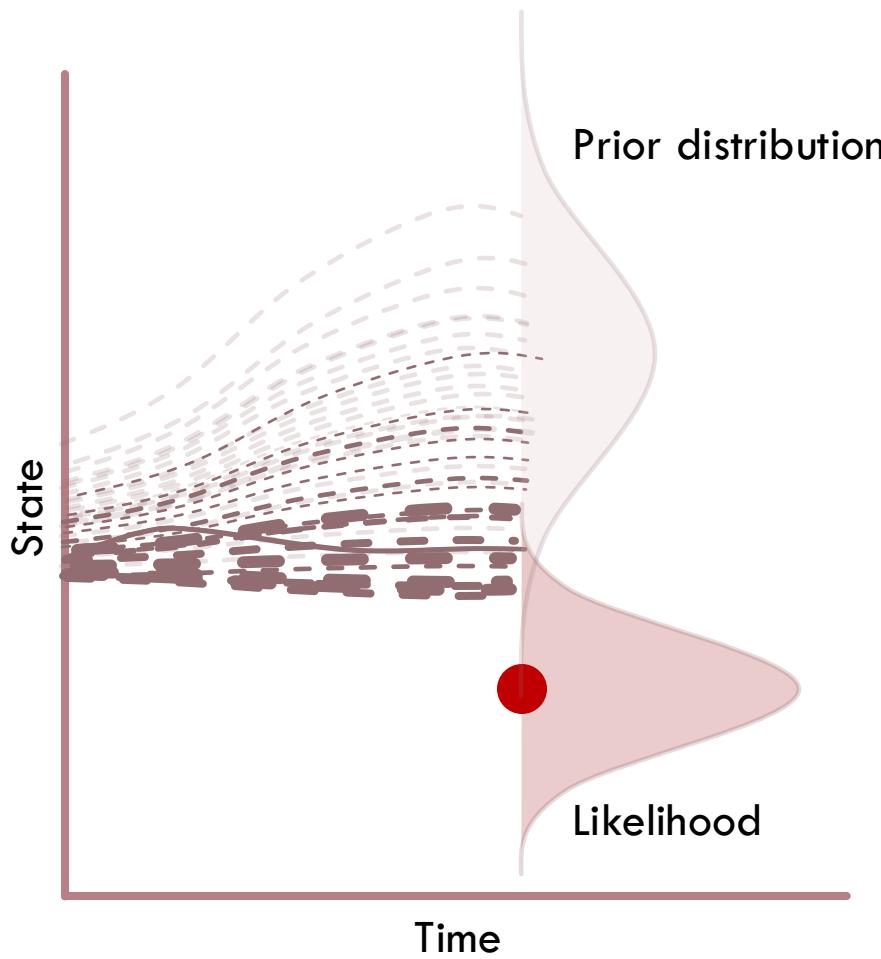
# PARTICLE FILTER



1. Compute prior particles

$$\{(w_n^i, u_n^i)\}_{i=0}^{N_p}$$

# PARTICLE FILTER



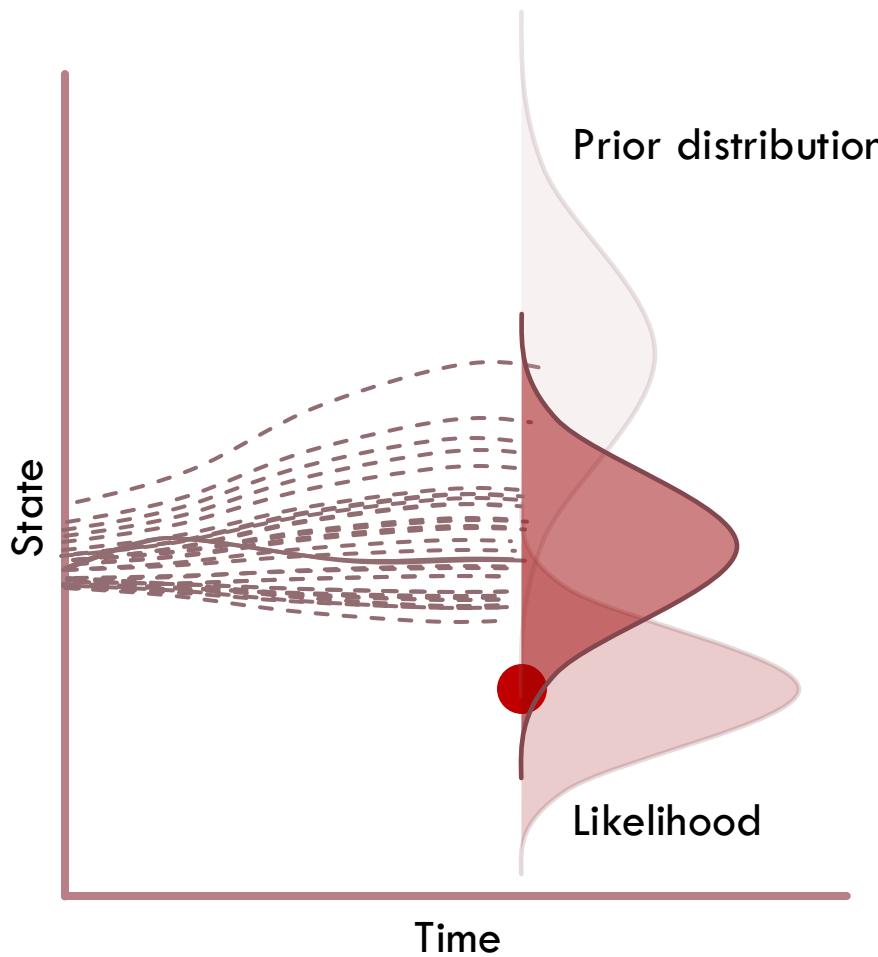
1. Compute prior particles

$$\{(w_n^i, u_n^i)\}_{i=0}^{N_p}$$

2. Update weights

$$\tilde{w}_{n+1}^i = w(u_n)p(y_{n+1}|u_{n+1}) = w(u_n) \cdot \text{likelihood}$$

# PARTICLE FILTER



1. Compute prior particles

$$\{(w_n^i, u_n^i)\}_{i=0}^{N_p}$$

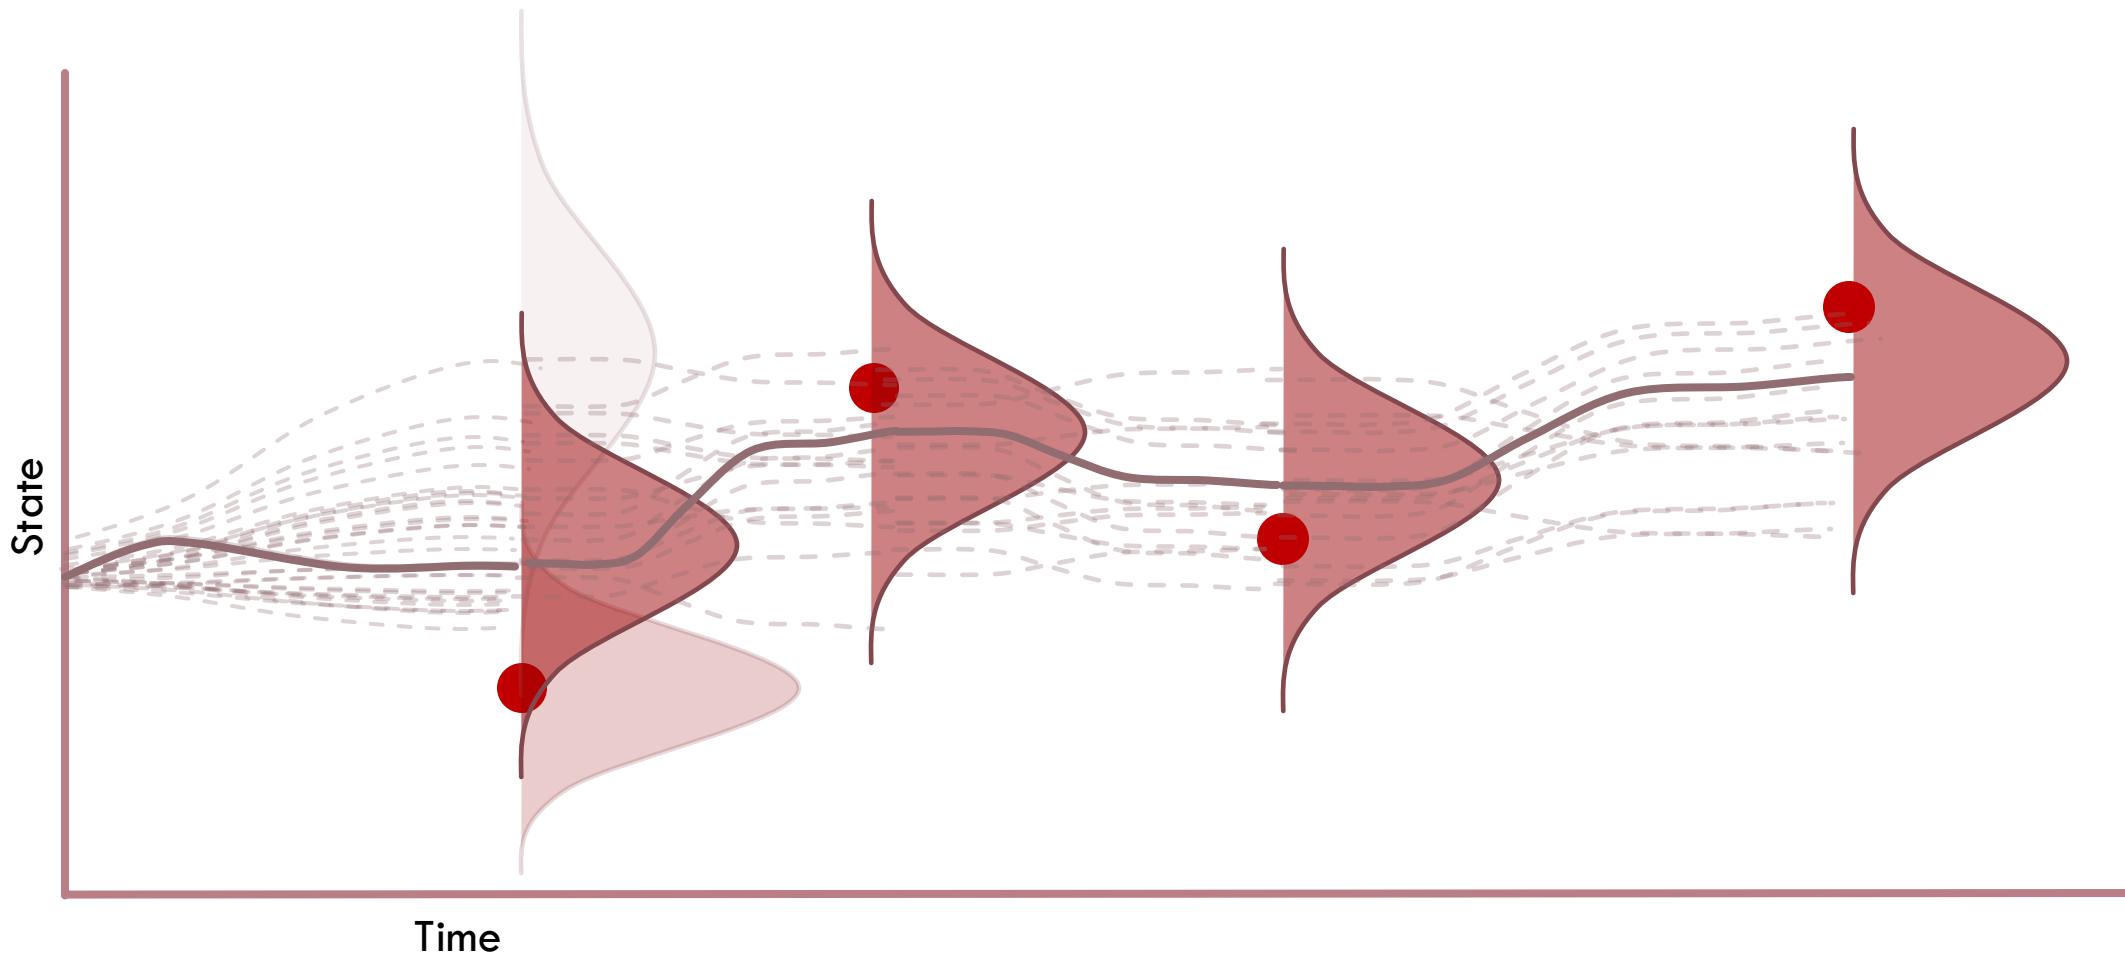
2. Update weights

$$\tilde{w}_{n+1}^i = w(u_n) p(y_{n+1} | u_{n+1}) = w(u_n) \cdot \text{likelihood}$$

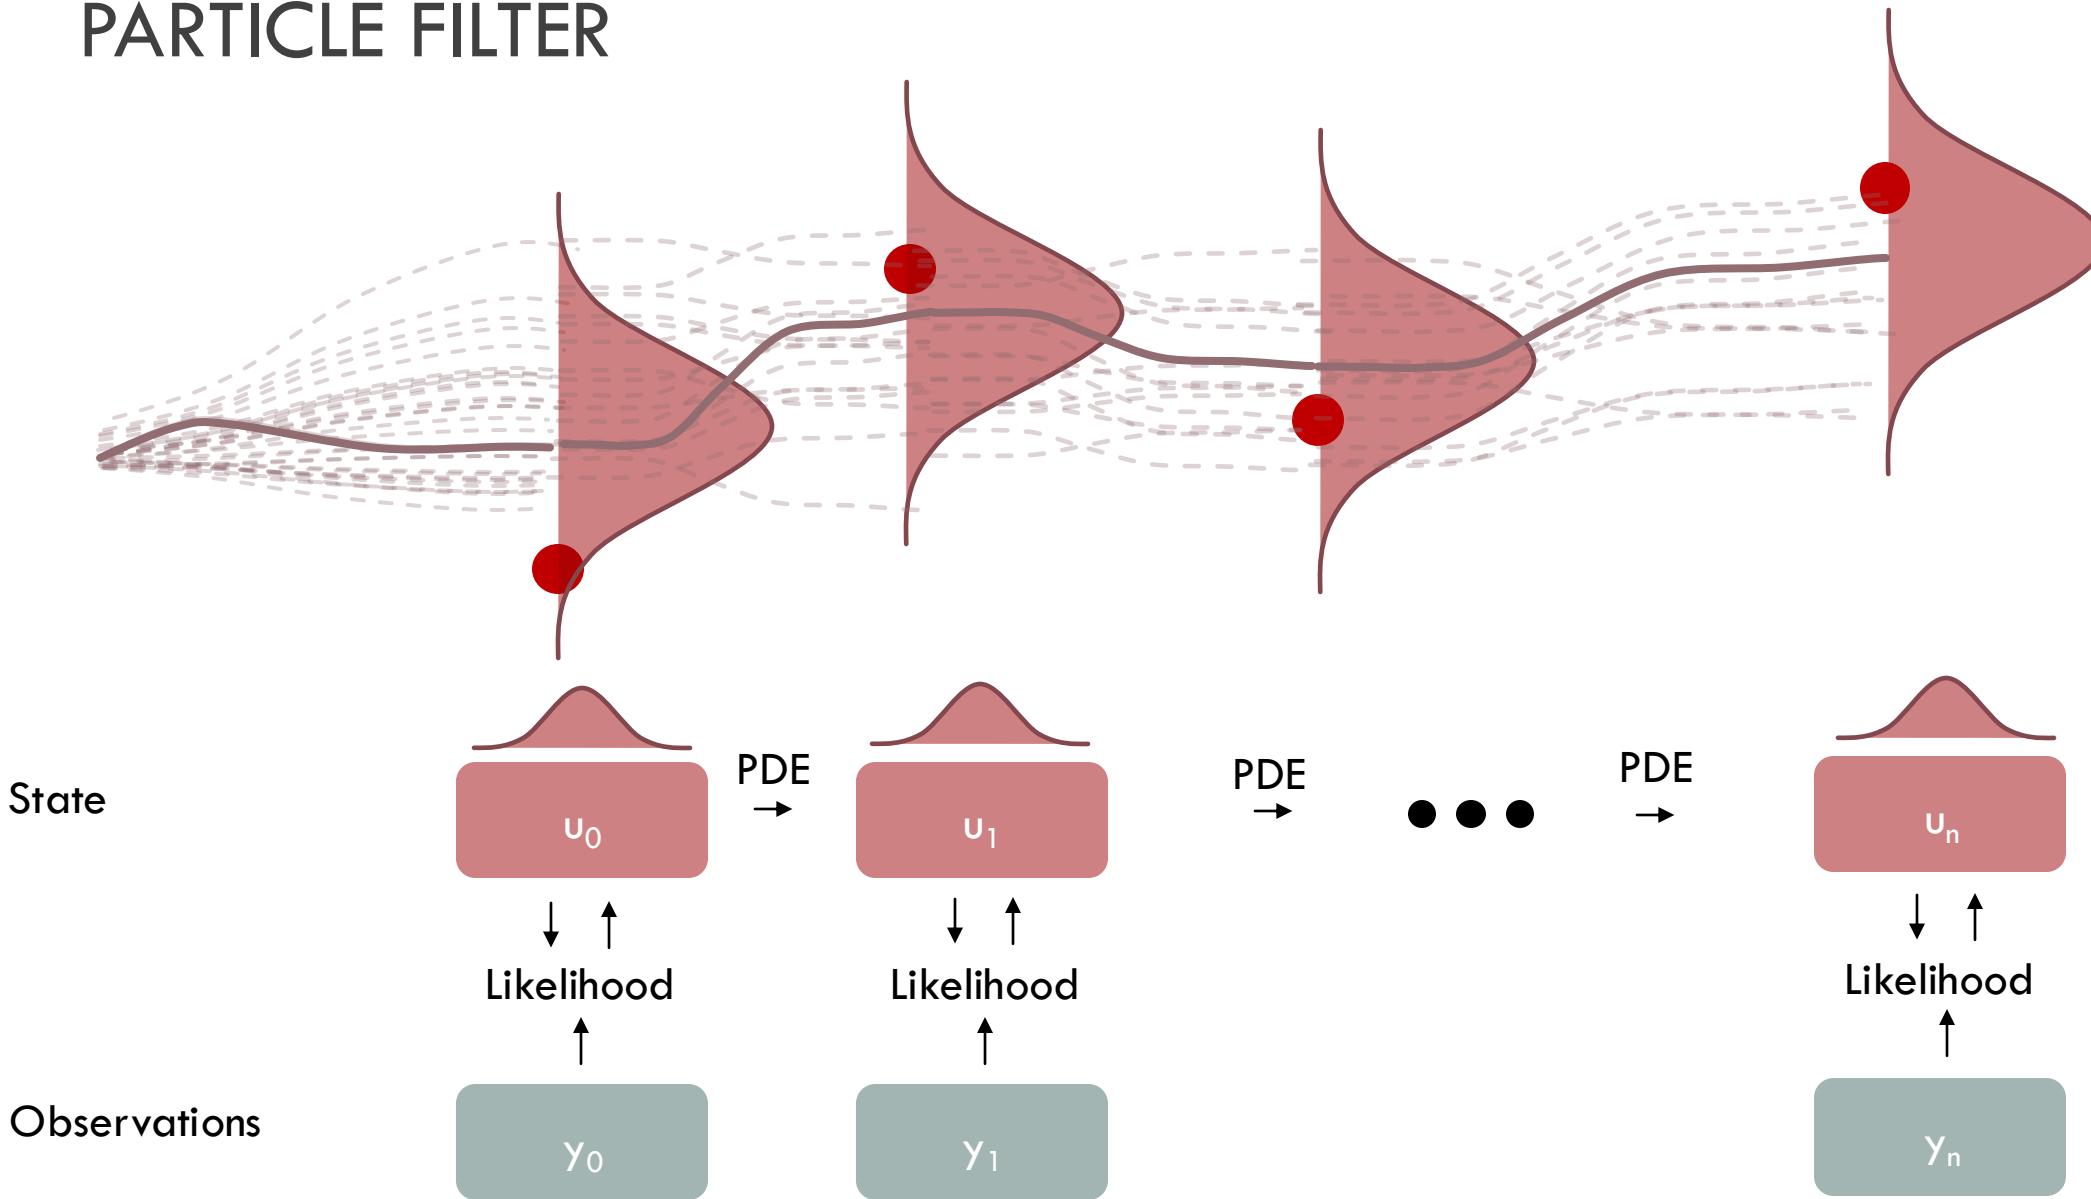
3. Resample according to weights

$$p(u_{n+1} | y_{0:n+1}) = \sum_{i=1}^{N_p} w_{n+1}^i \delta(u_{n+1} - u_{n+1}^i)$$

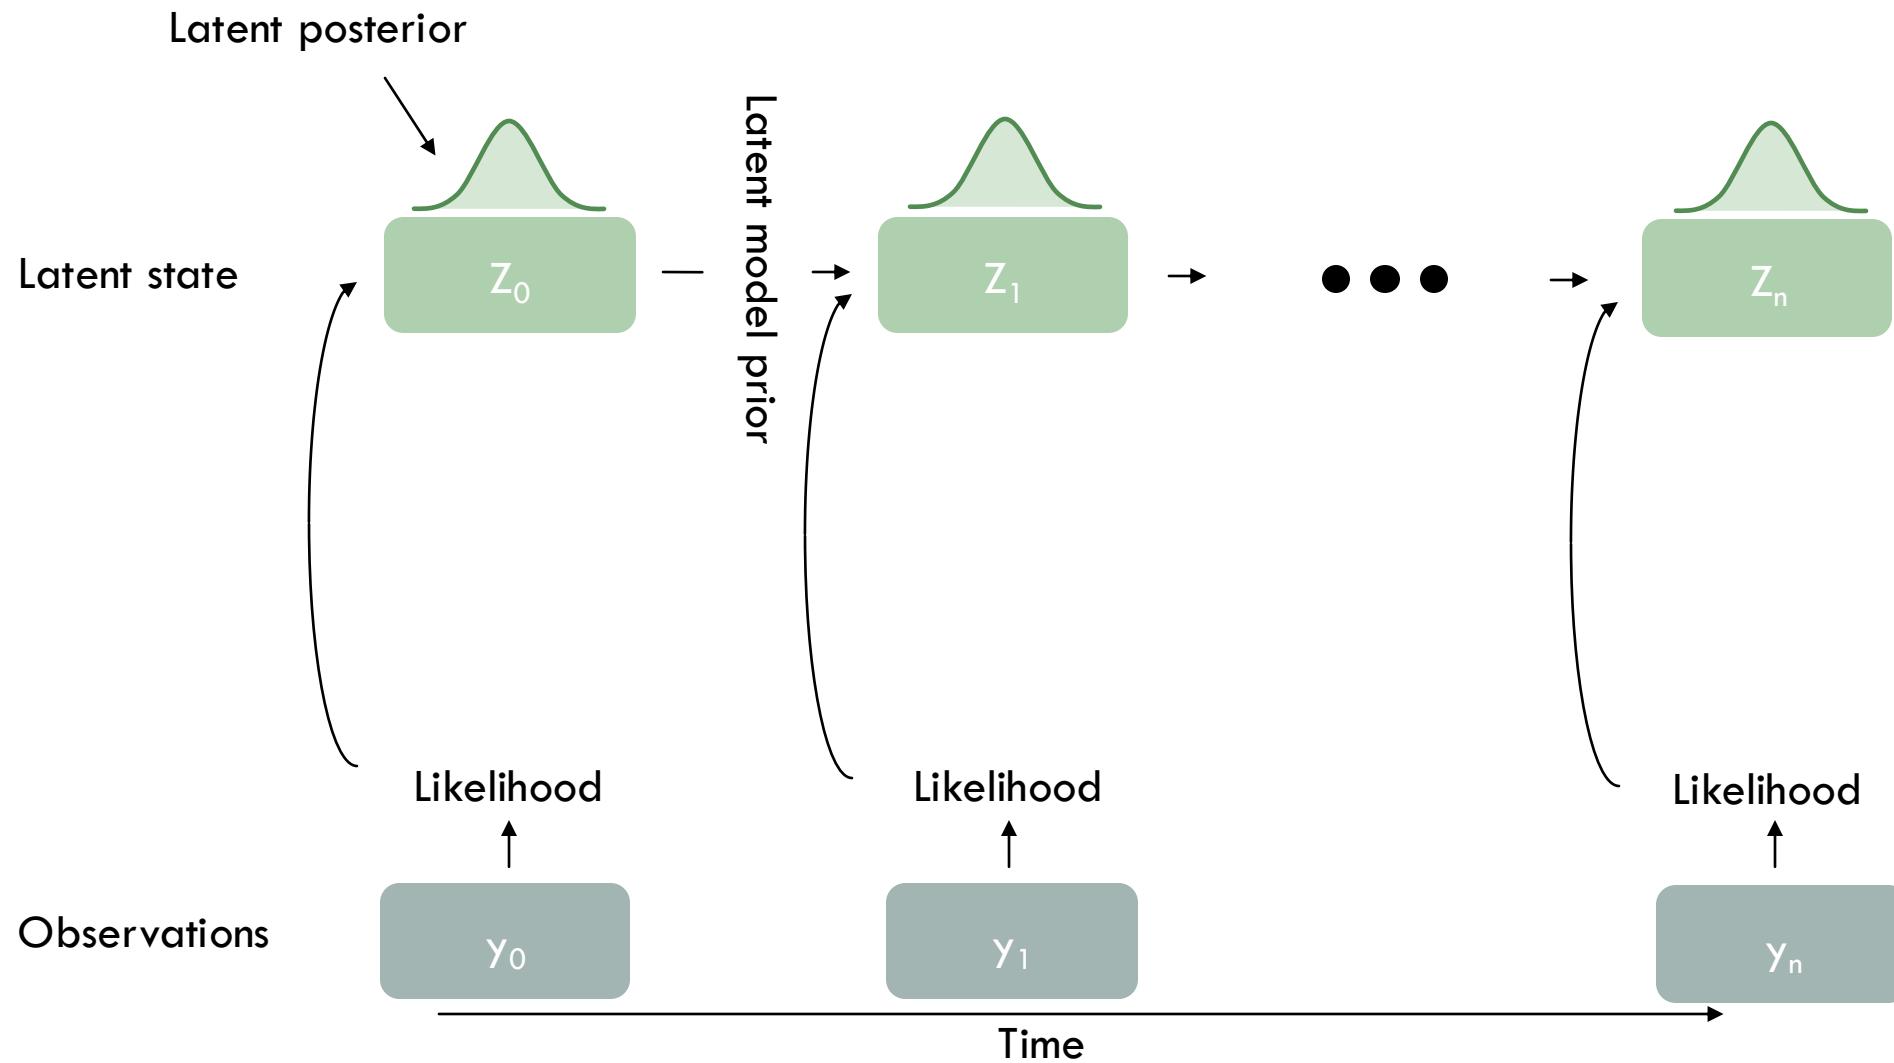
# PARTICLE FILTER



# PARTICLE FILTER



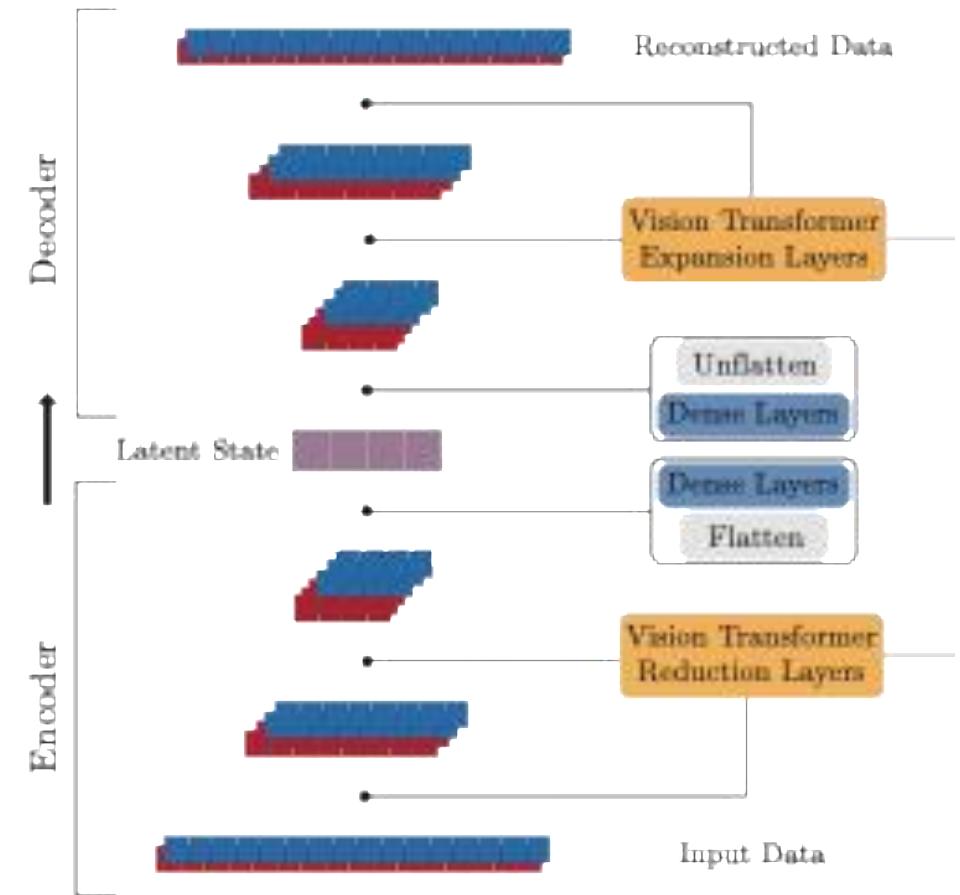
# DATA ASSIMILATION IN LATENT SPACE



# LATENT SPACE MODEL

# AUTOENCODER

- Nonlinear dimensionality reduction
- Encoder – Reduce dimension
- Decoder – Reconstruct data

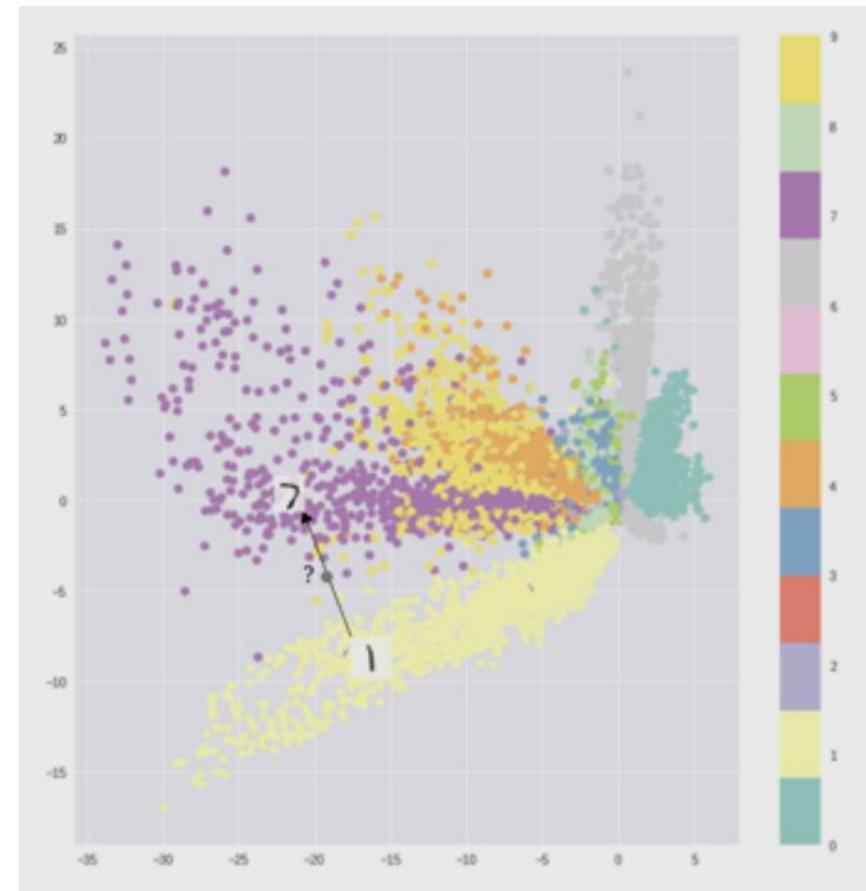


$$L_{\text{AE}}(\phi_{\text{enc}}, \phi_{\text{dec}}) = \underbrace{\frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \phi_{\text{dec}}(\phi_{\text{enc}}(\mathbf{x}_i)))^2}_{\text{reconstruction}} + \alpha \underbrace{R(\phi_{\text{enc}}, \phi_{\text{dec}})}_{\text{regularization}}.$$

# AUTOENCODER

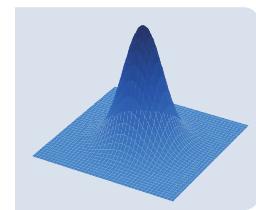
- Nonlinear dimensionality reduction
- Encoder – Reduce dimension
- Decoder – Reconstruct data

2D Latent space

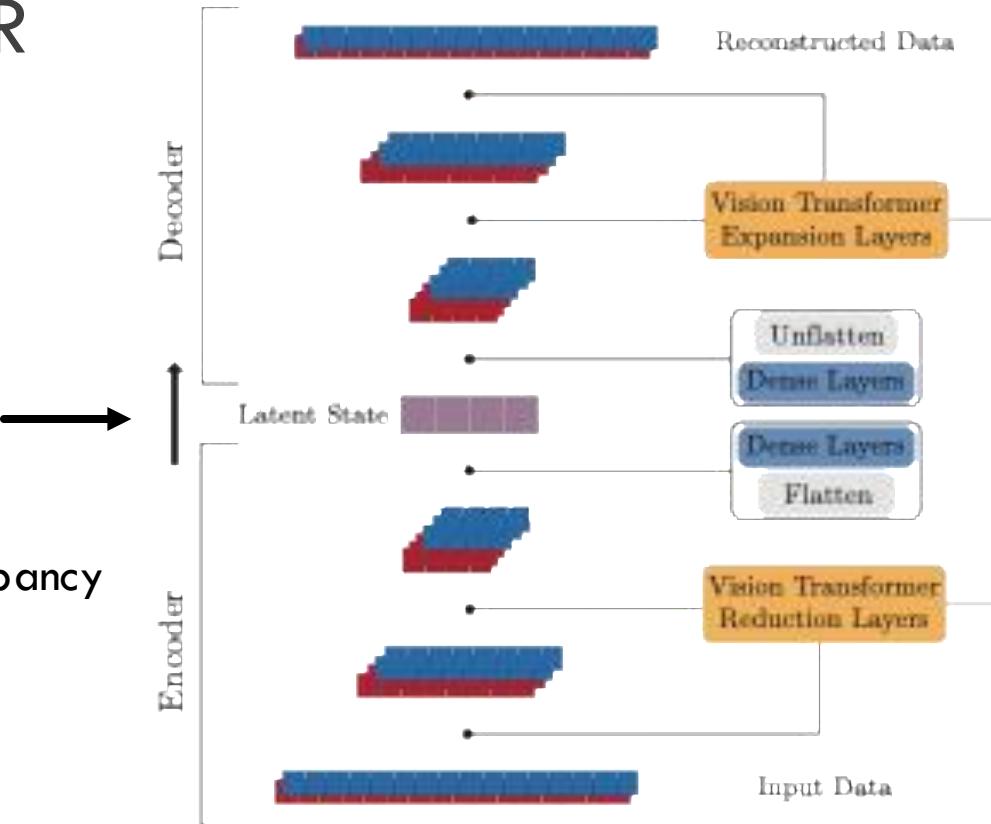


$$L_{\text{AE}}(\phi_{\text{enc}}, \phi_{\text{dec}}) = \underbrace{\frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \phi_{\text{dec}}(\phi_{\text{enc}}(\mathbf{x}_i)))^2}_{\text{reconstruction}} + \alpha \underbrace{R(\phi_{\text{enc}}, \phi_{\text{dec}})}_{\text{regularization}}.$$

# WASSERSTEIN AUTOENCODER



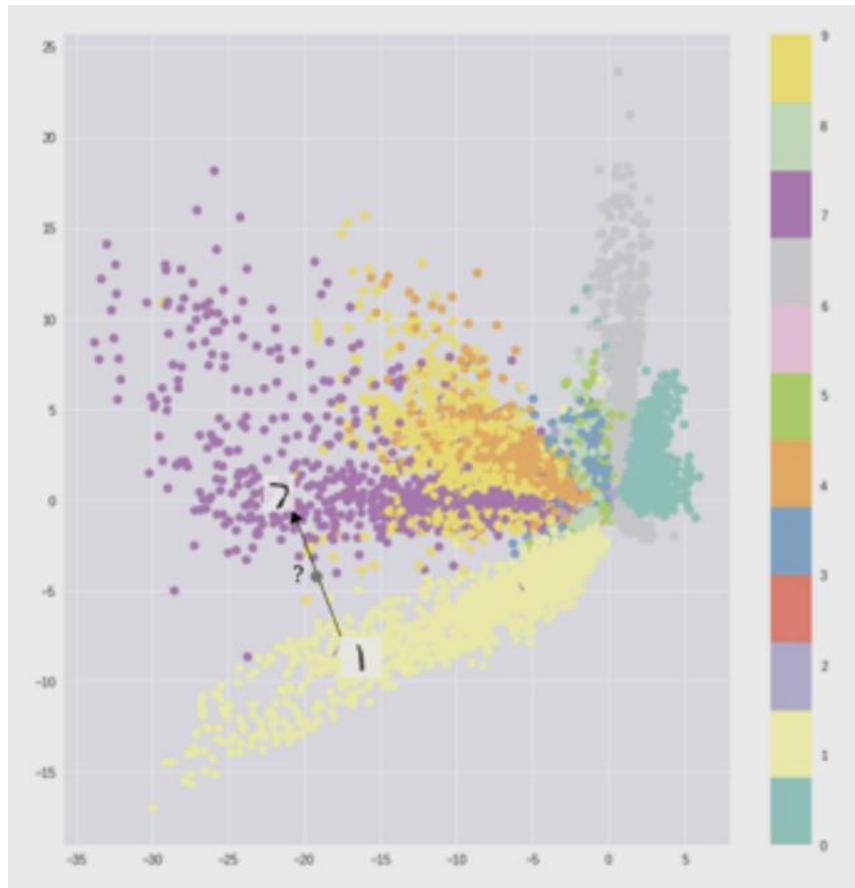
Maximum Mean Discrepancy  
Regularization



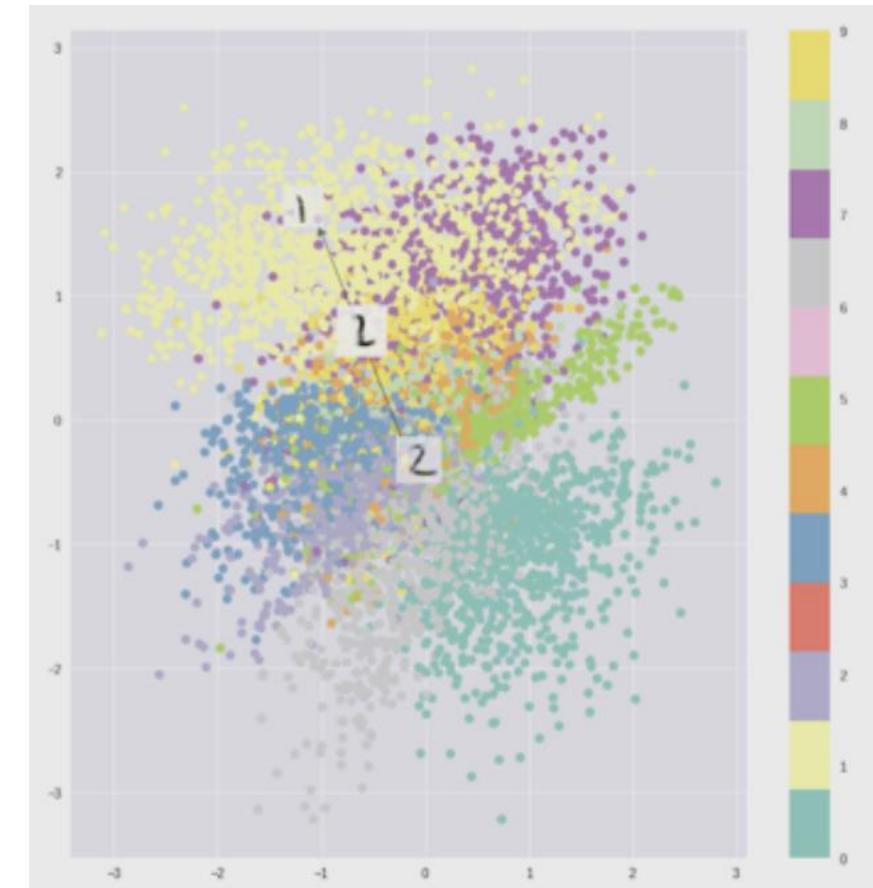
$$L_{\text{WAE}}(\phi_{\text{enc}}, \phi_{\text{dec}}) = \underbrace{\frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \phi_{\text{dec}}(\phi_{\text{enc}}(\mathbf{x}_i)))^2}_{\text{reconstruction}} + \underbrace{\lambda D(p(\mathbf{z}), p_{\text{enc}}(\mathbf{z}))}_{\text{divergence}} + \alpha \underbrace{R(\phi_{\text{enc}}, \phi_{\text{dec}})}_{\text{regularization}}.$$

# COMPARISON OF LATENT SPACE

No latent space divergence loss



With latent space divergence loss



# LEARNING TO TIME-STEP

- Encode high-fidelity training data

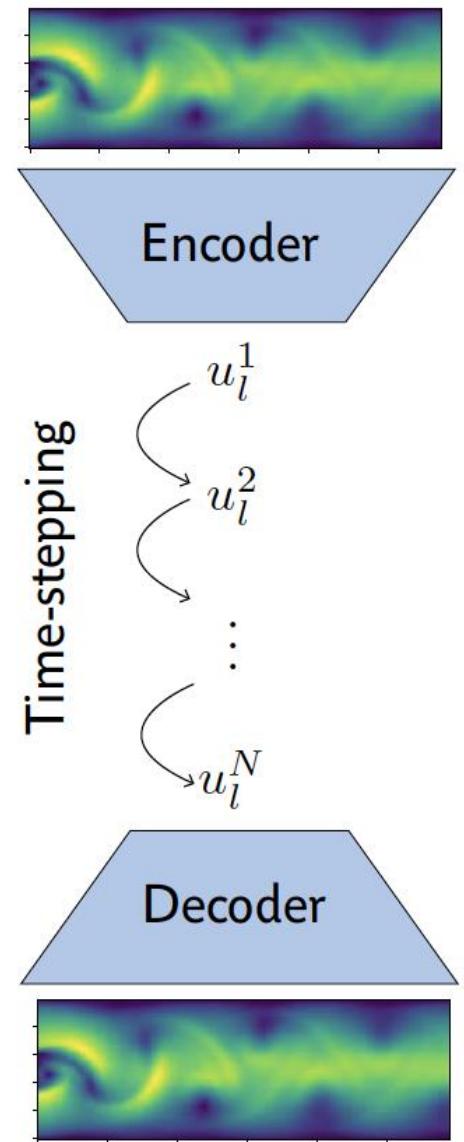
$$\{\mathbf{q}_0, \dots, \mathbf{q}_{N_t}\} \rightarrow \{\text{Enc}(\mathbf{q}_0), \dots, \text{Enc}(\mathbf{q}_{N_t})\} = \{\mathbf{z}_0, \dots, \mathbf{z}_{N_t}\}$$

- Minimize time-stepping error      Time-stepping NN

$$\sum_{i=1}^{N_t} \|\mathbf{z}_i - f(\mathbf{z}_{i-1}, \theta)\|$$

- Loop unrolling

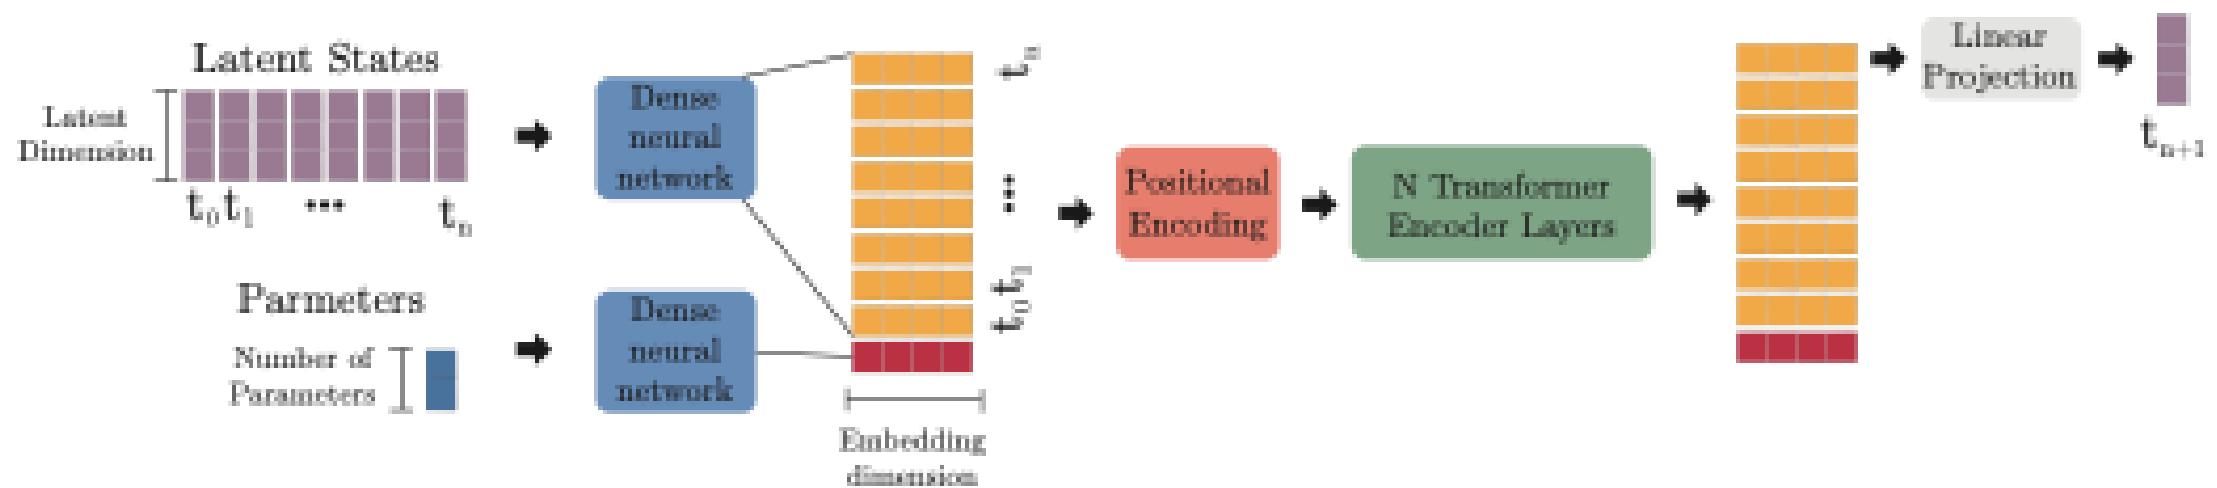
$$\sum_{i=1}^{N_t-k} \sum_{k=0}^K \|\mathbf{z}_{i+k} - f^k(\mathbf{z}_{i-1}, \theta)\|$$



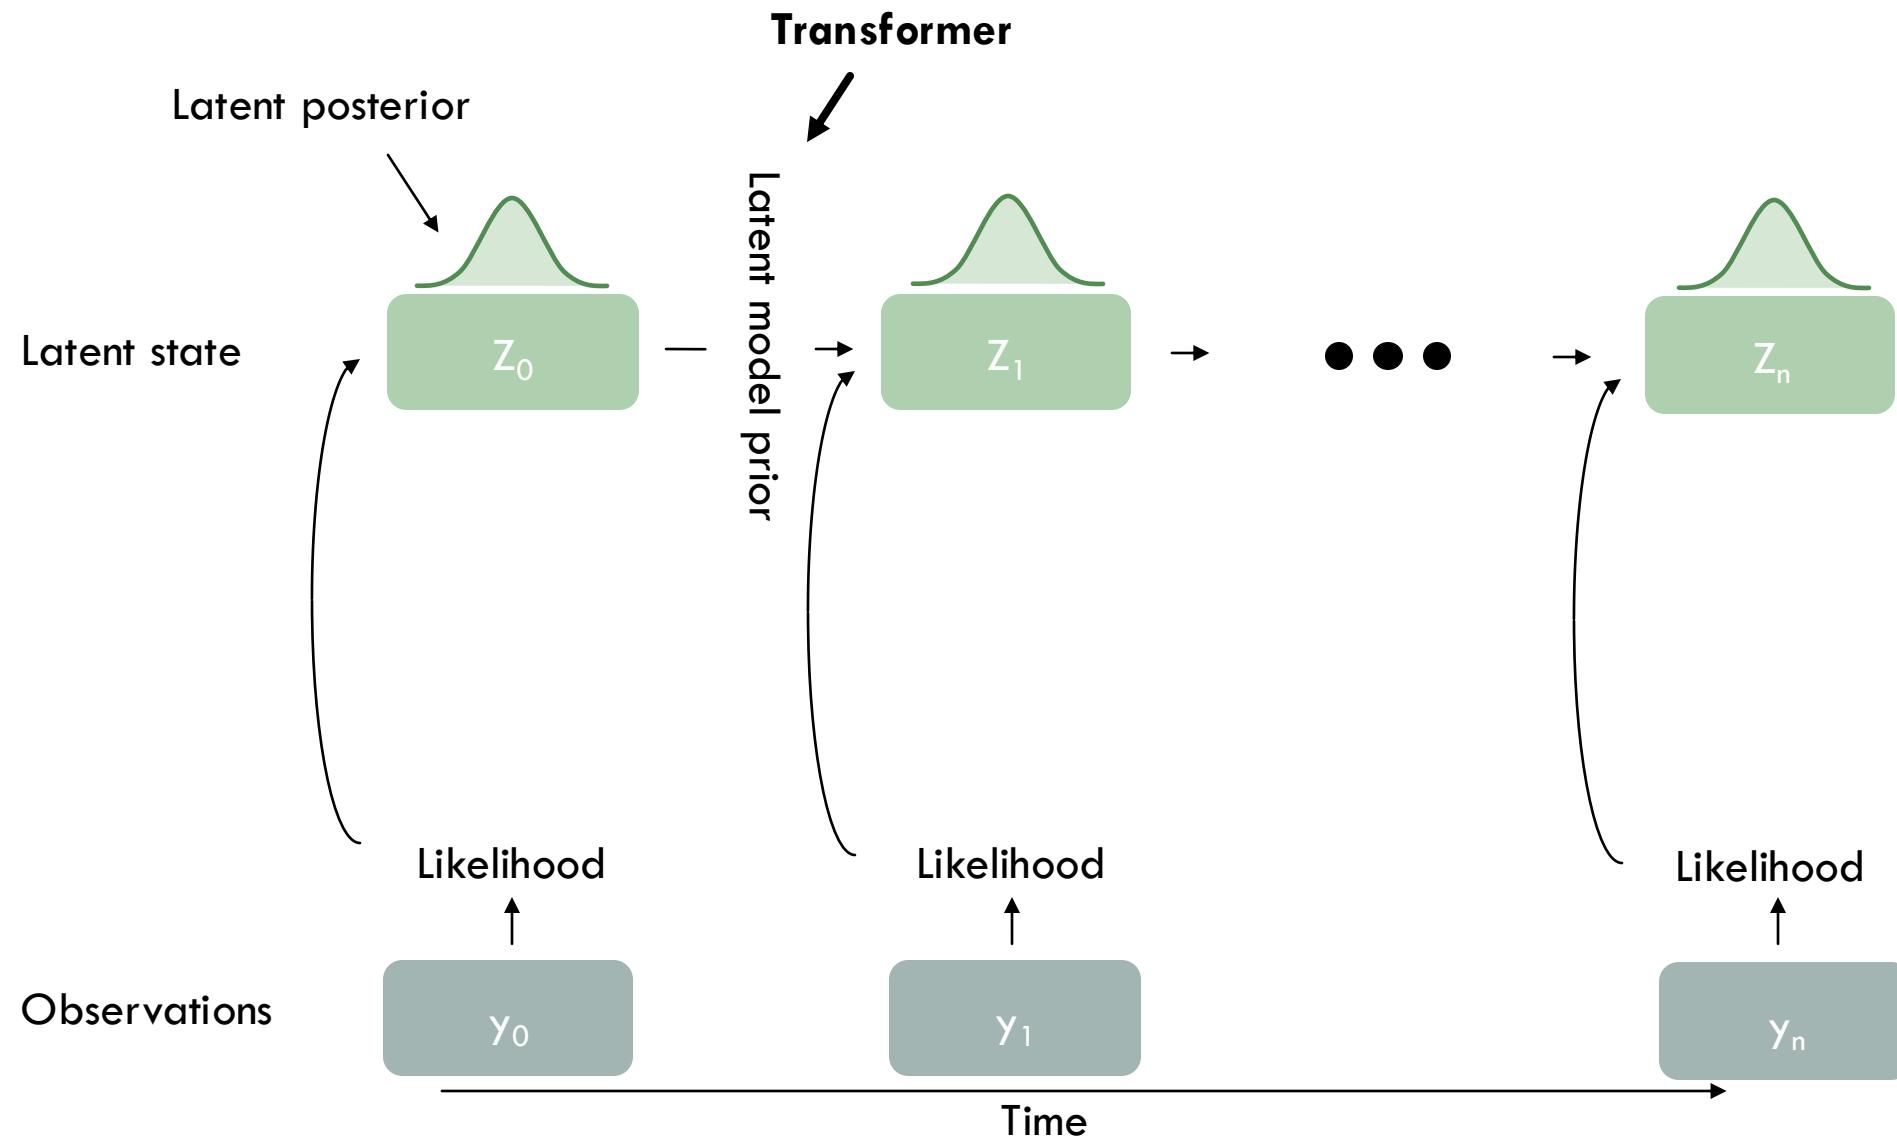
# TRANSFORMERS

- They show **state-of-the-art** performance for **time series forecasting**
  - Outperforming RNNs, LSTMs, CNNs, etc...
- Transformers "mimic" **multistep methods**
  - "The function space of a self-attention layer with a residual connection contains the space of explicit linear time-integration methods within an arbitrarily small non-zero amount of error."
  - *Geneva, Nicholas, and Nicholas Zabaras. "Transformers for modeling physical systems." Neural Networks 146 (2022): 272-289.*
- Great at incorporating **parametric dependence through attention**
  - *Han, Xu, et al. "Predicting physics in mesh-reduced space with temporal attention." arXiv preprint arXiv:2201.09113 (2022).*

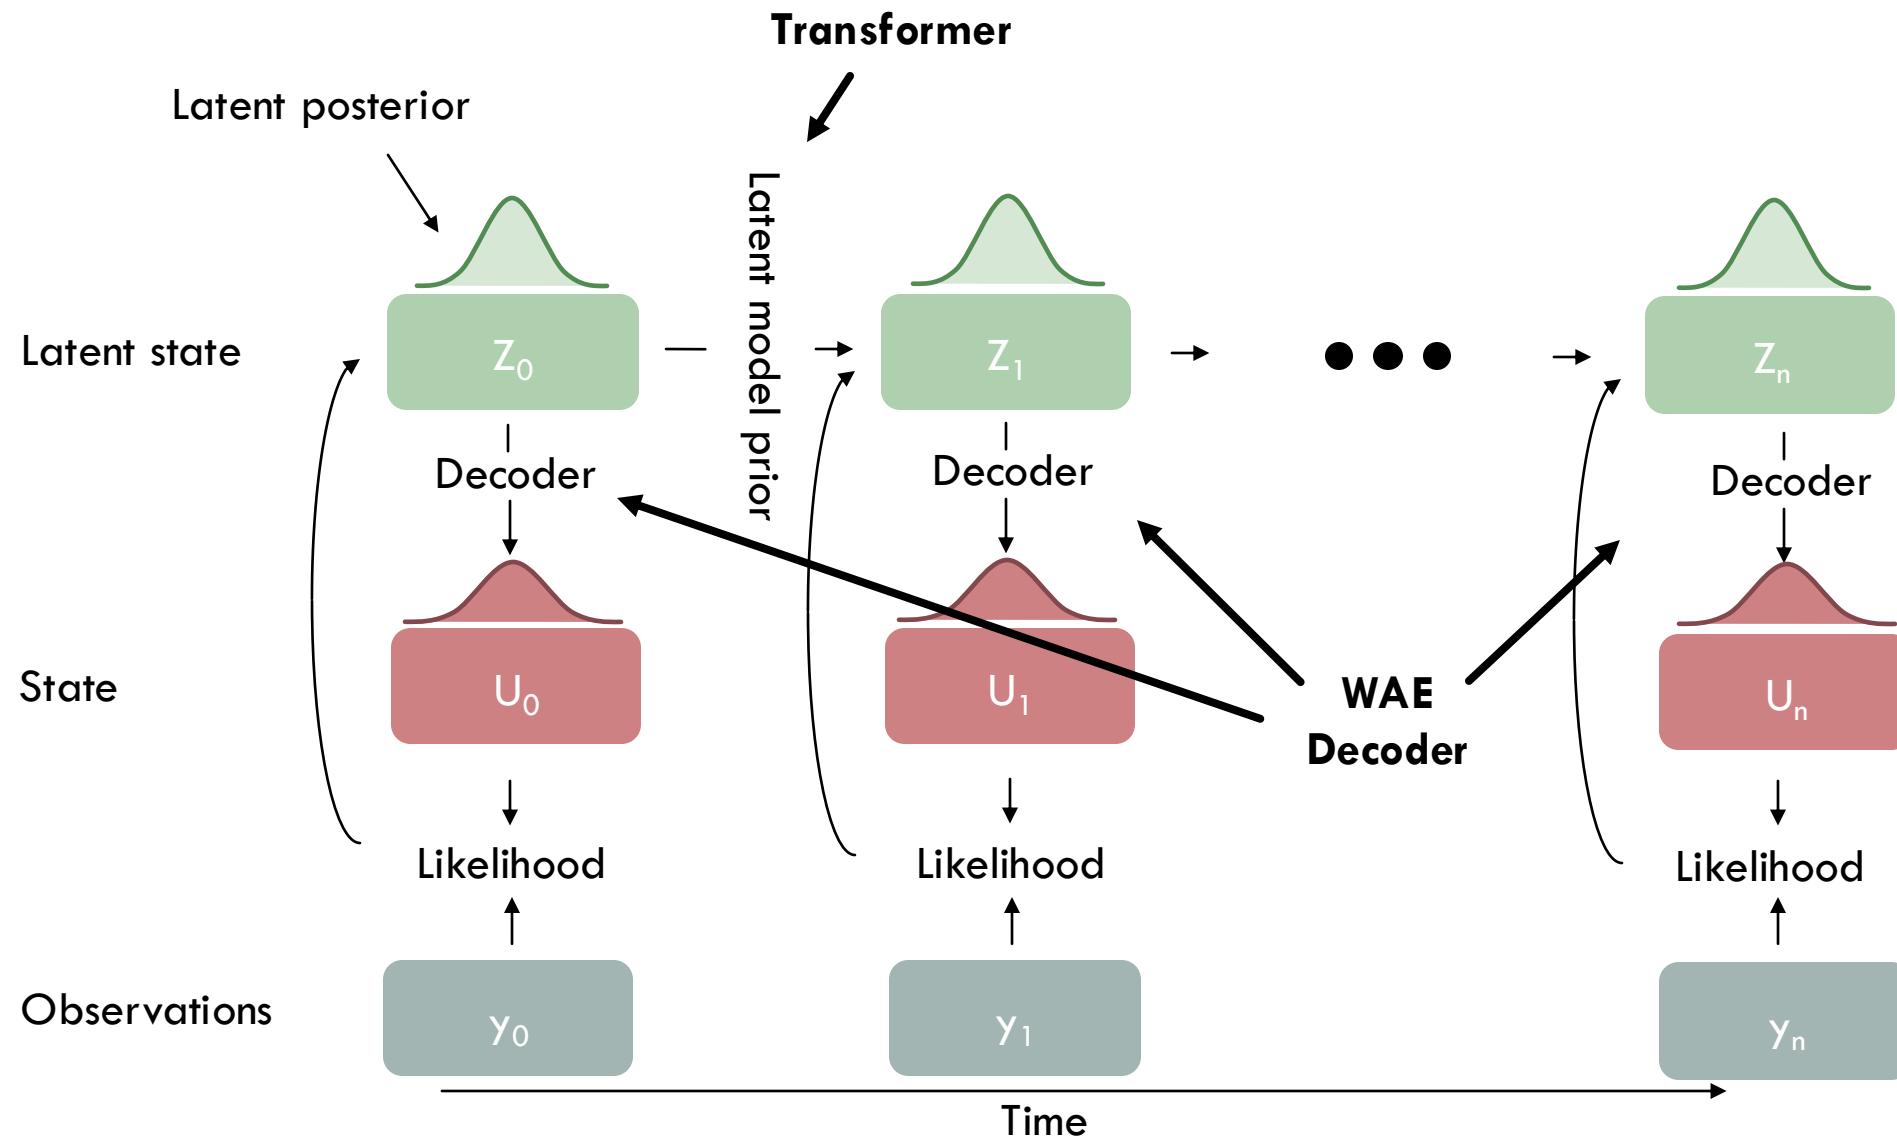
# PARAMETERIZED TIME-STEPPING

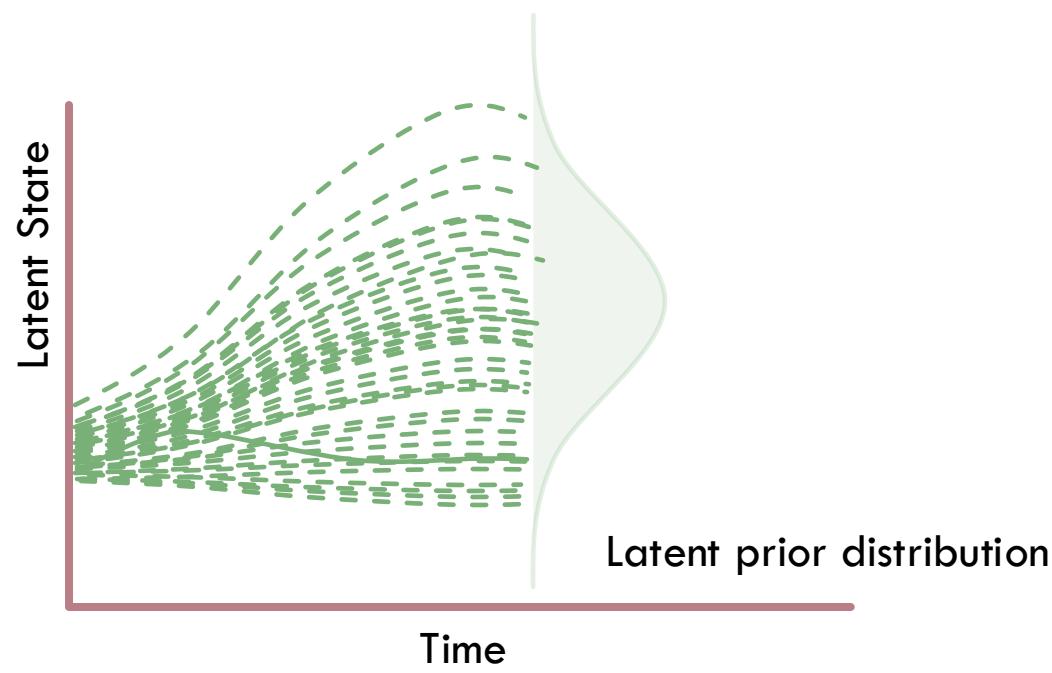
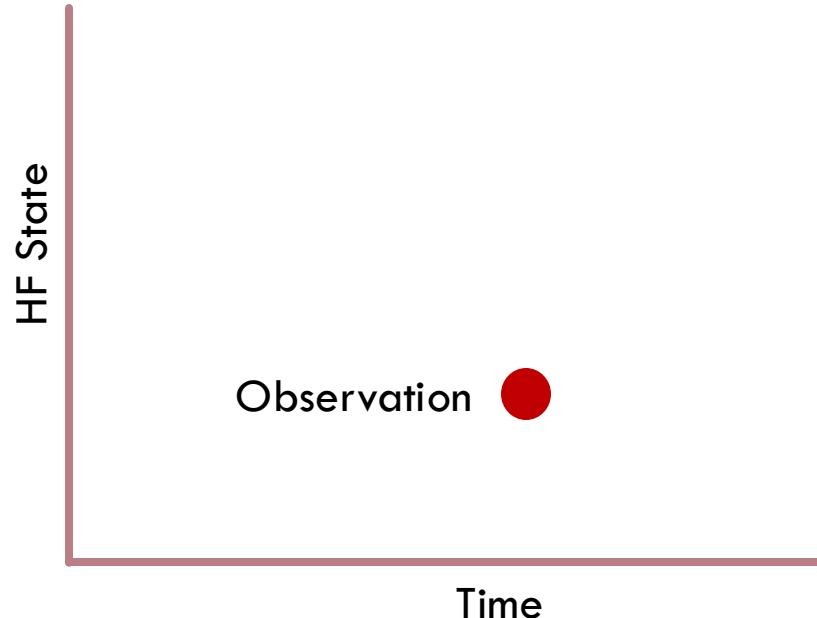


# DATA ASSIMILATION IN LATENT SPACE

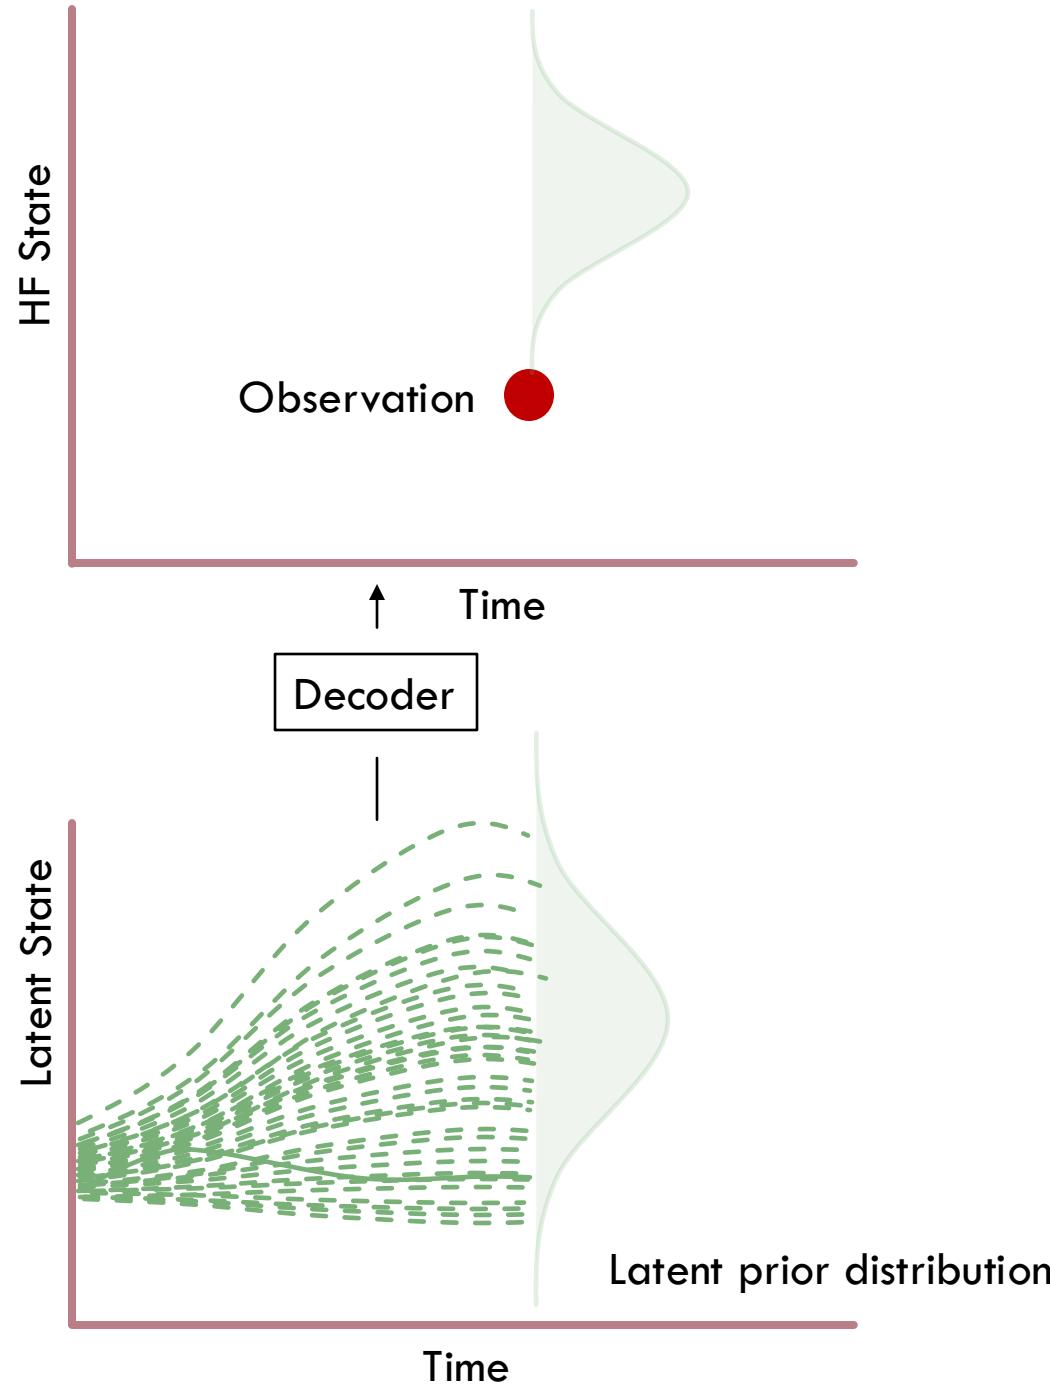


# DATA ASSIMILATION IN LATENT SPACE

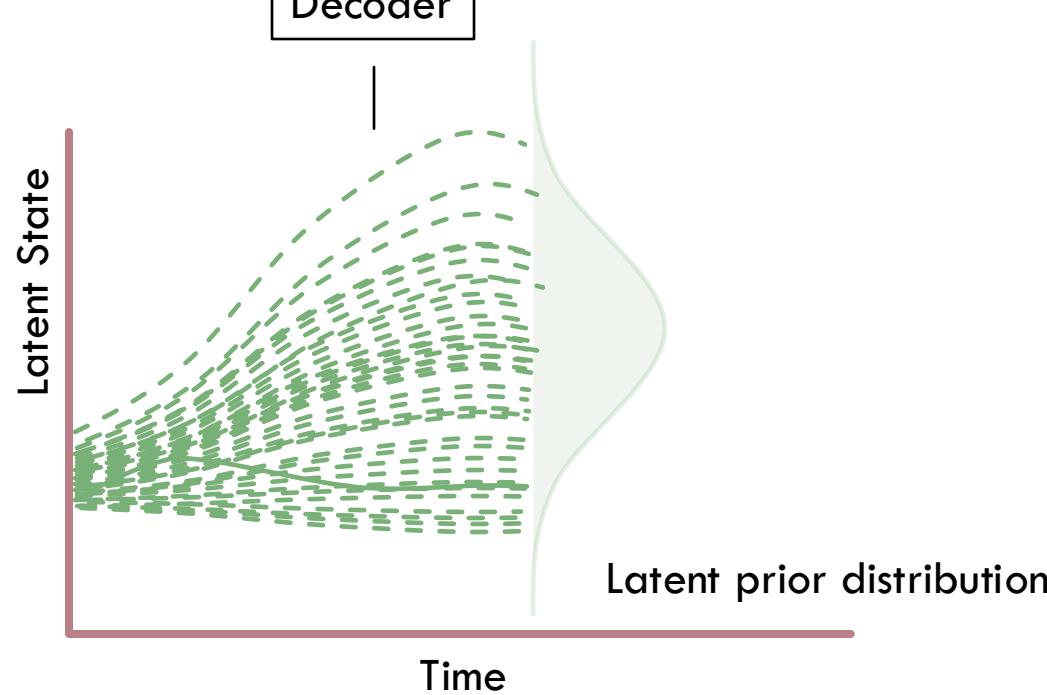
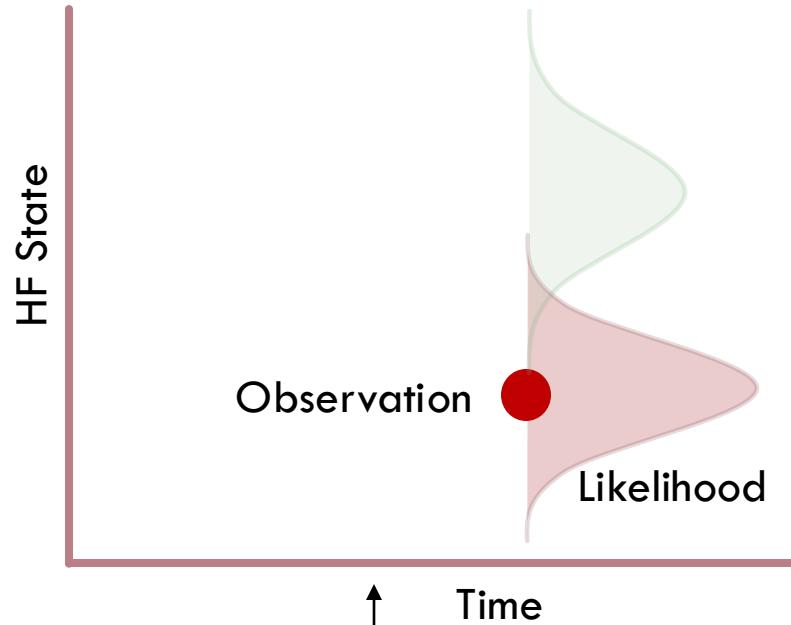




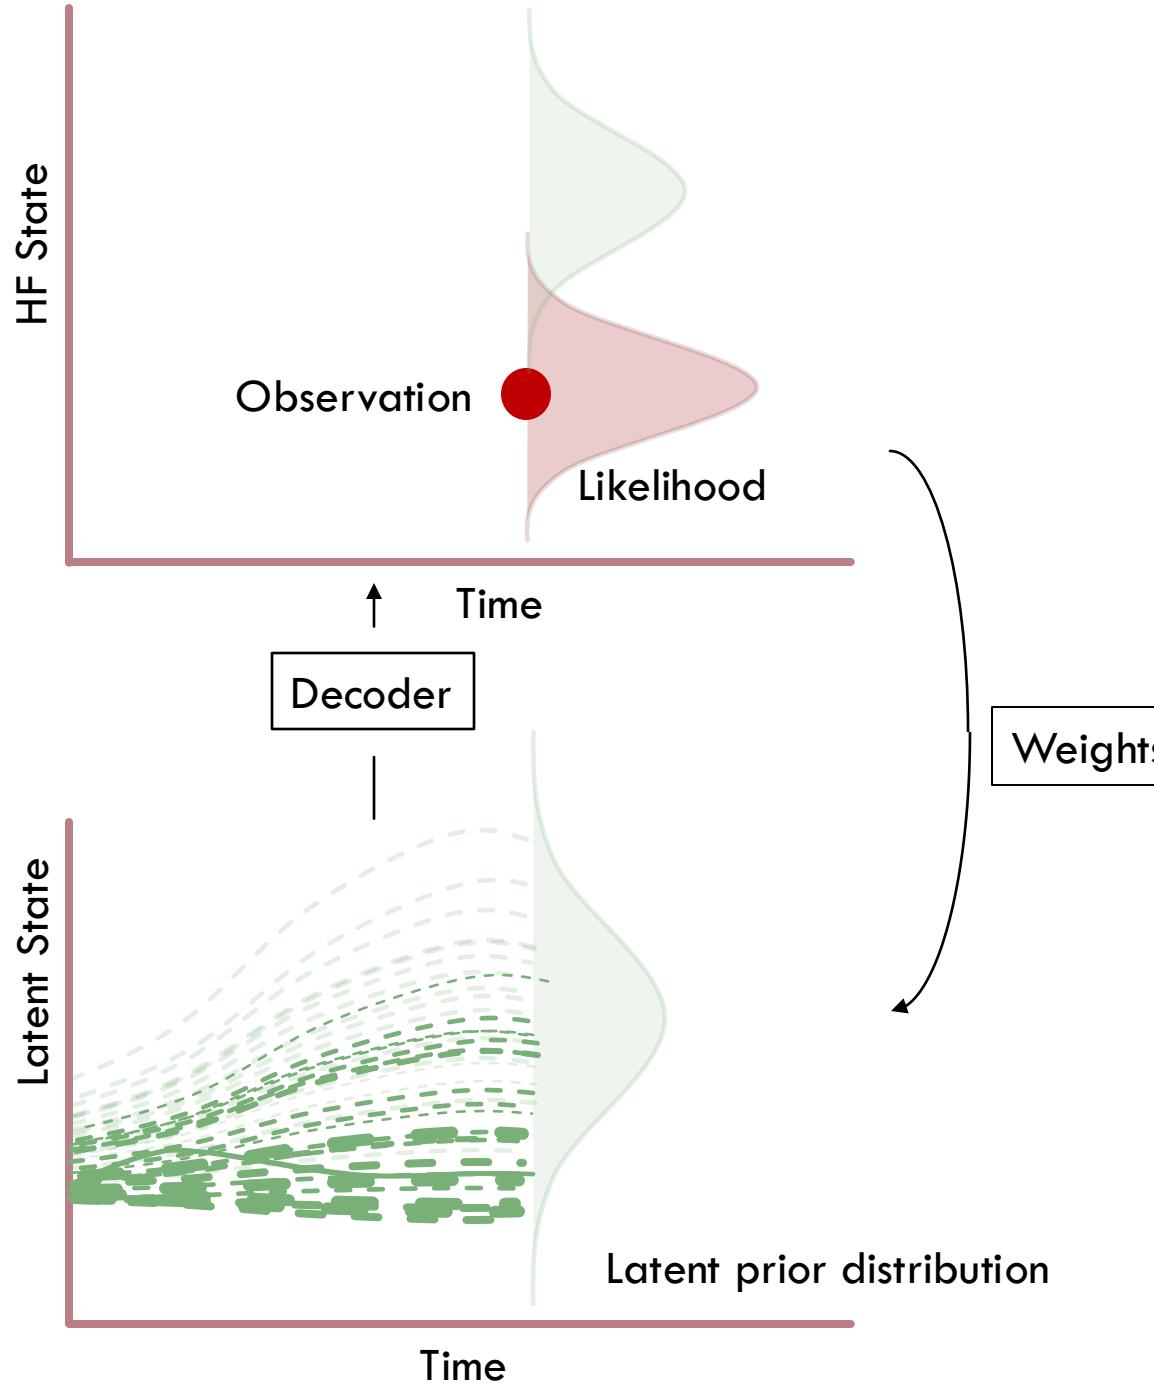
1. Compute latent prior



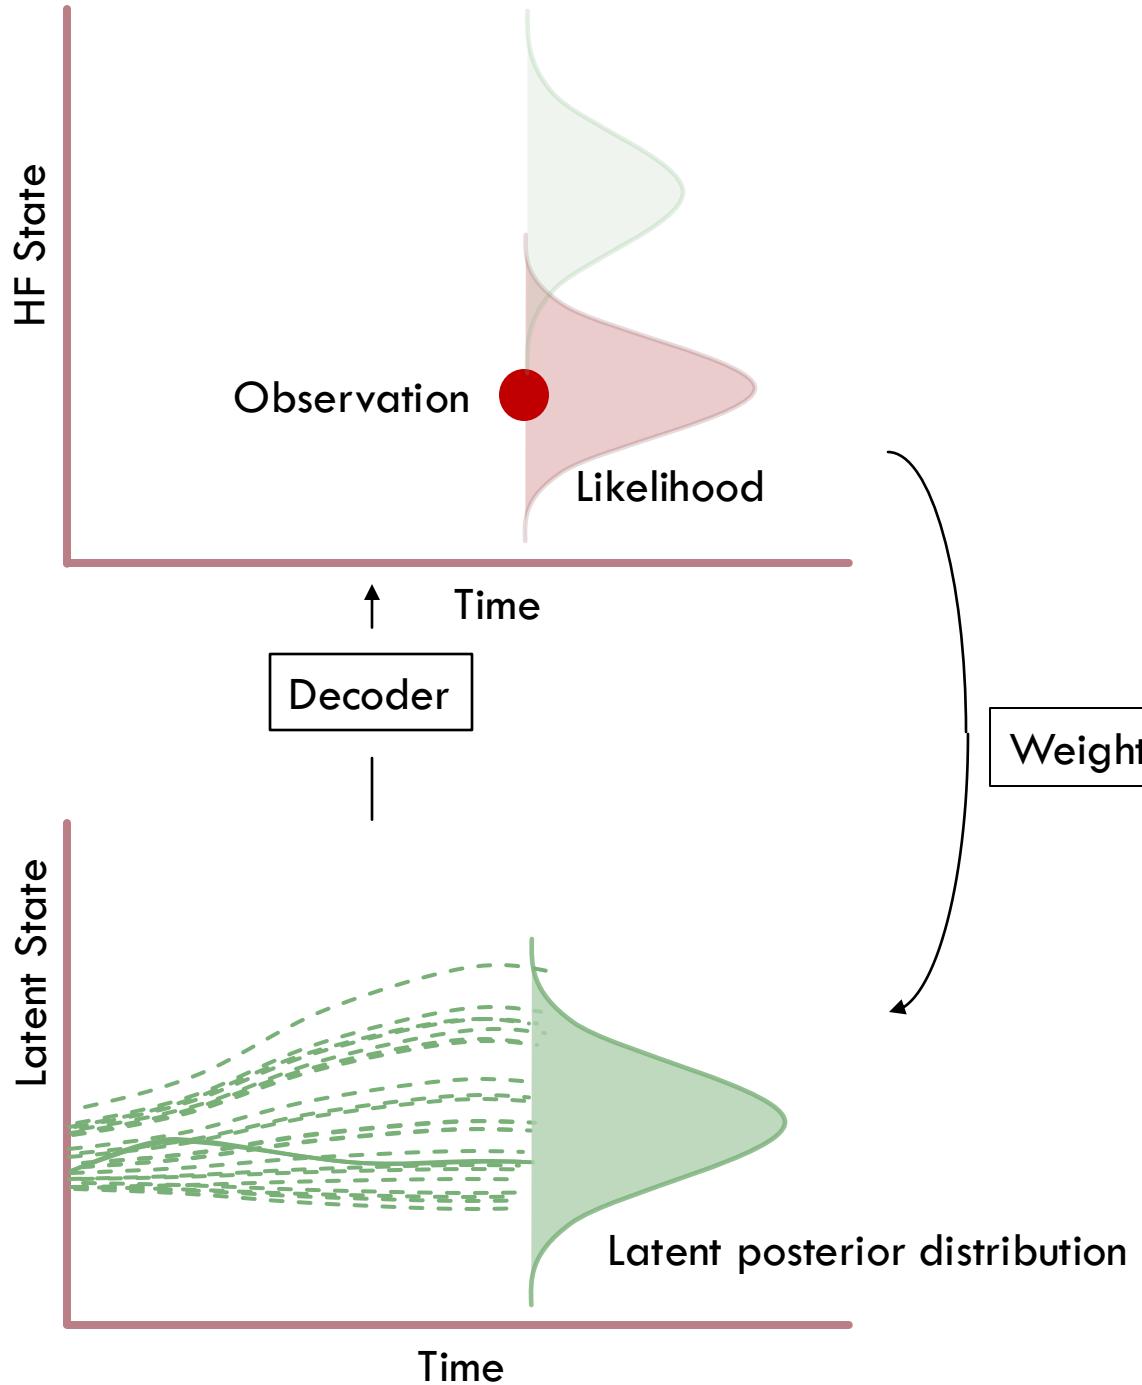
1. Compute latent prior
2. Decode latent samples



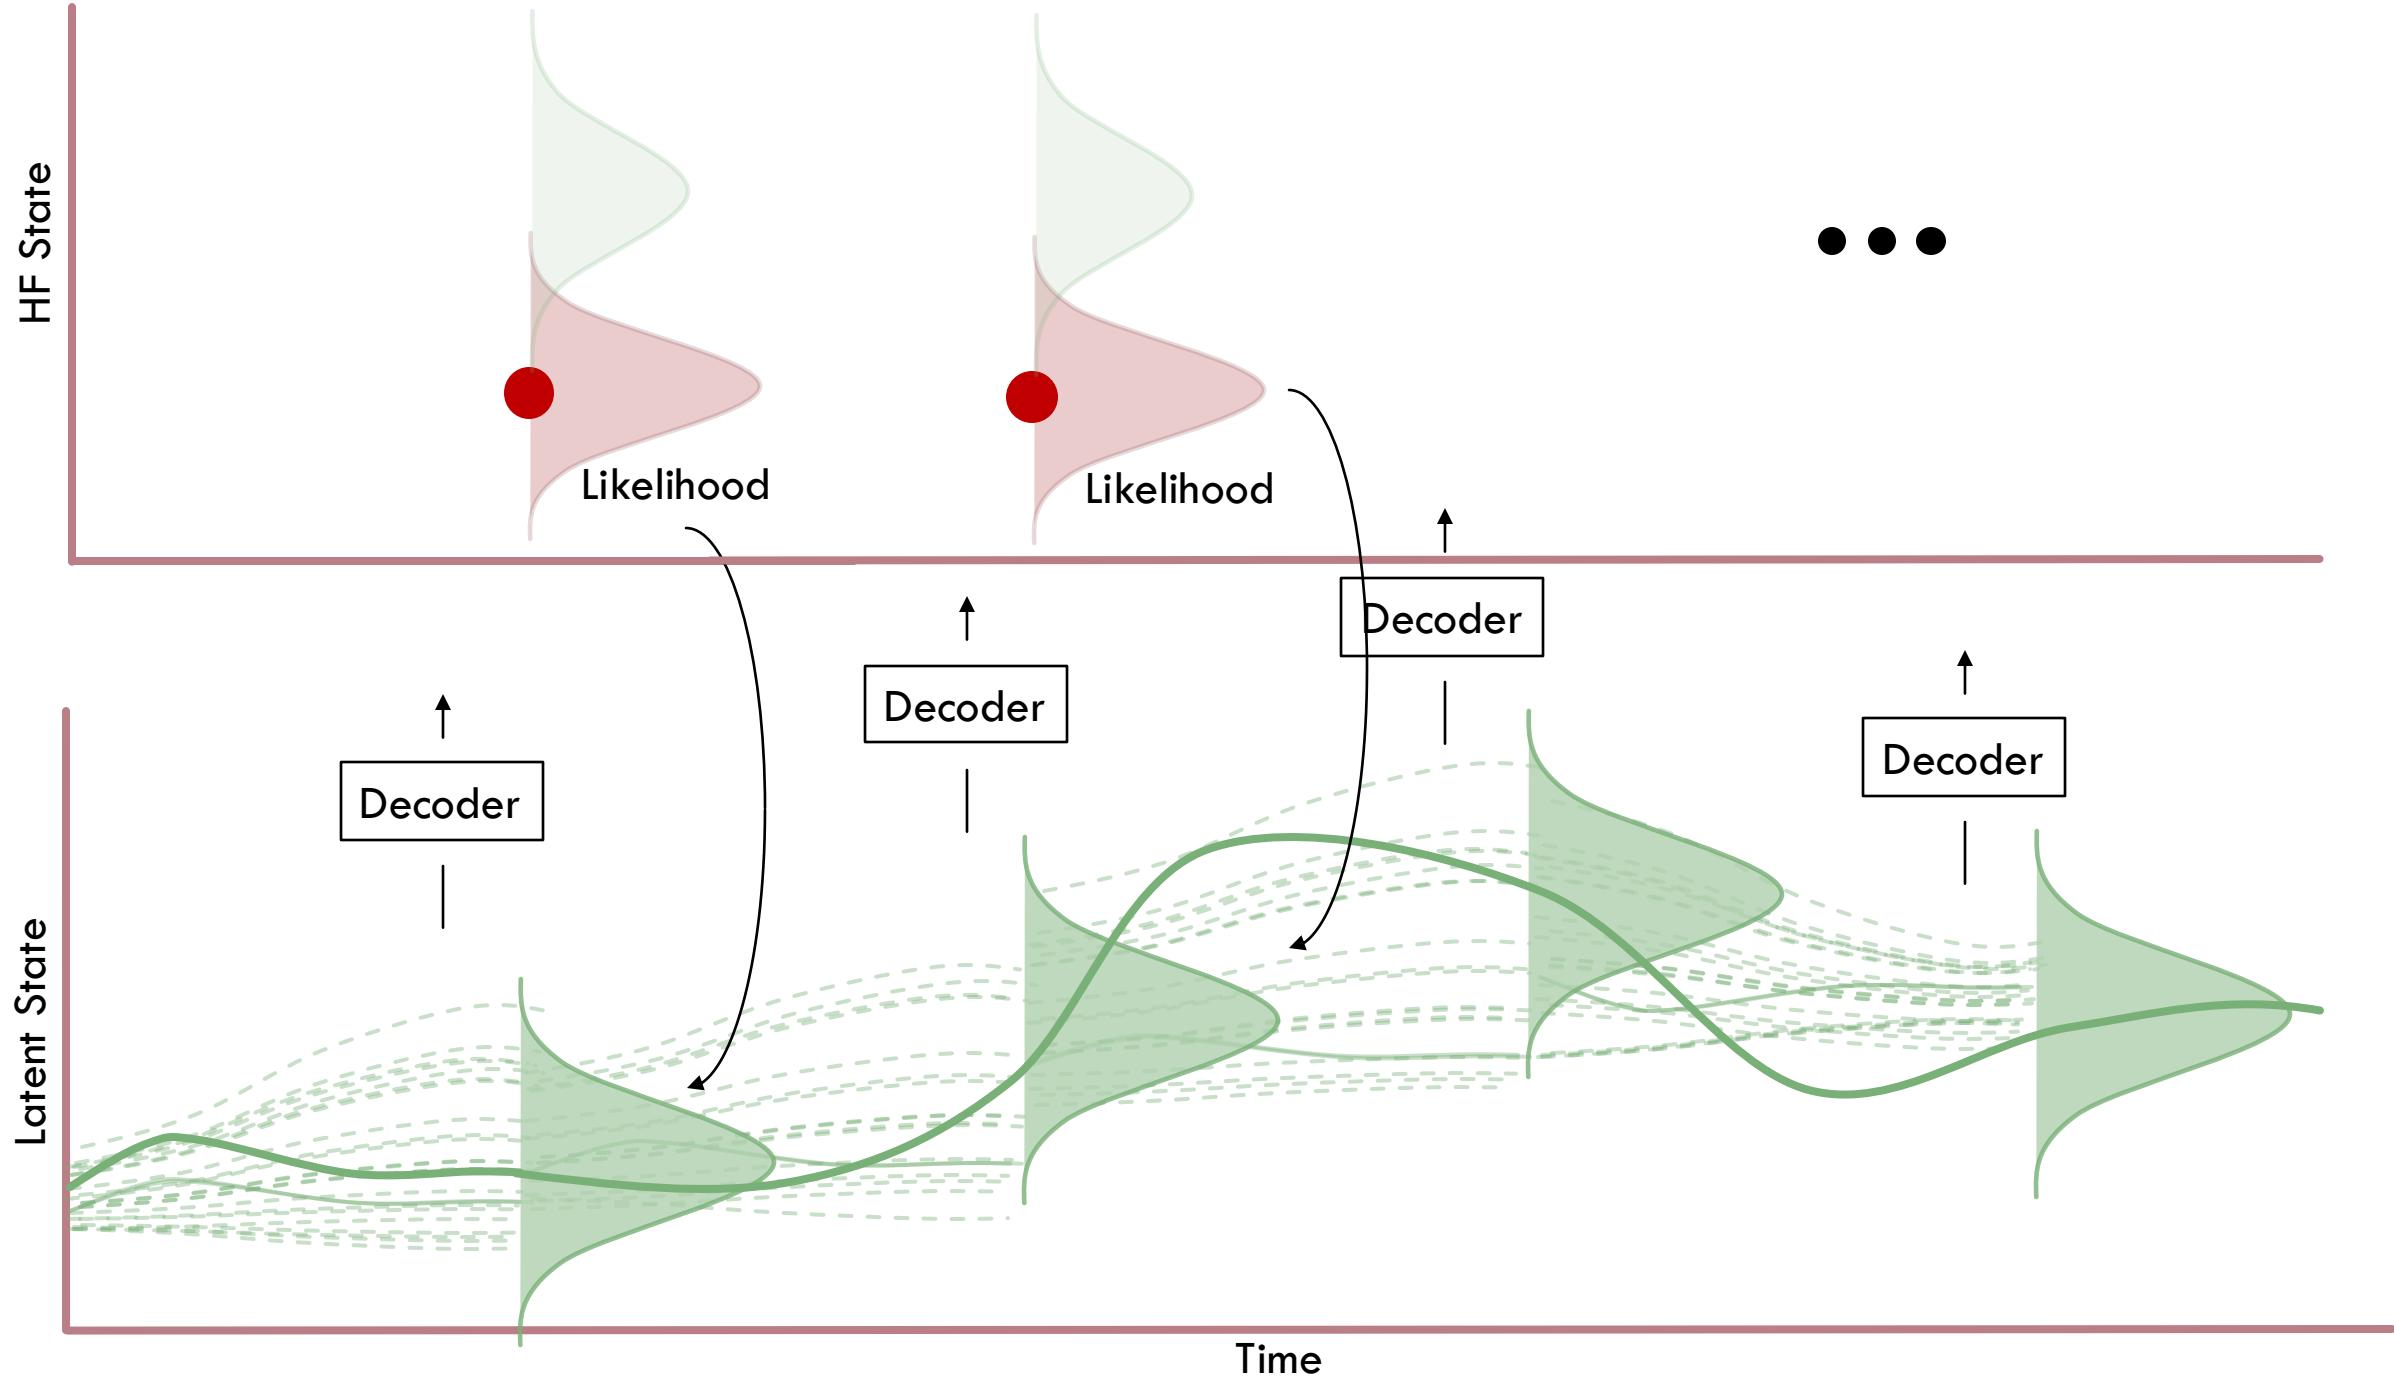
1. Compute latent prior
2. Decode latent samples
3. Compute likelihood



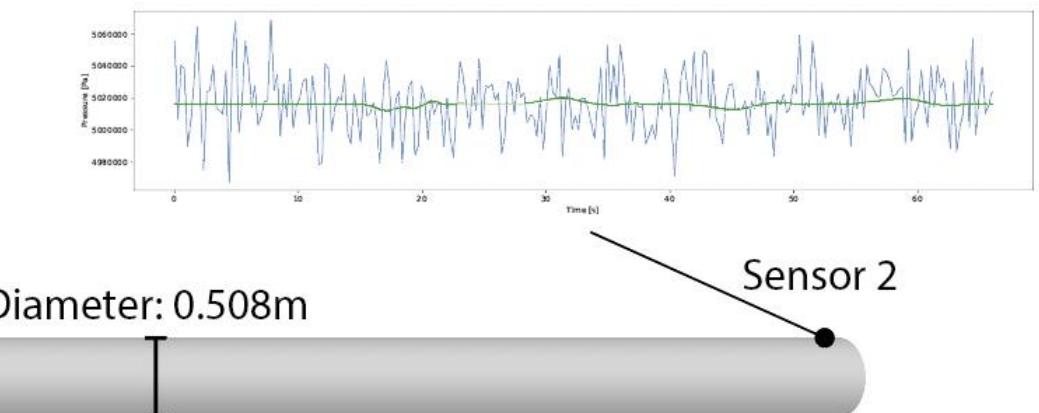
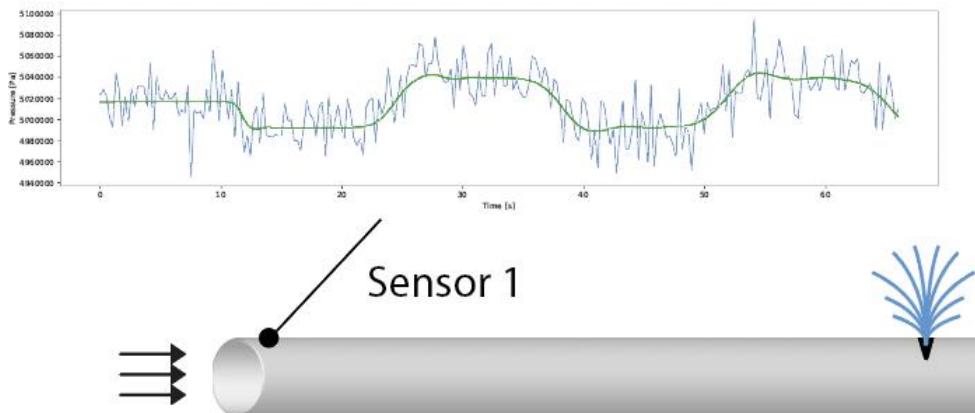
1. Compute latent prior
2. Decode latent samples
3. Compute likelihood
4. Compute weights



1. Compute latent prior
2. Decode latent samples
3. Compute likelihood
4. Compute weights
5. Resample to get latent posterior



# LEAK LOCALIZATION



Diameter: 0.508m

# PROBLEM SETTING

Leak size

Leak location

$$\partial_t q_1 + \partial_x q_2 = C_d \sqrt{\rho(p(\rho) - p_{\text{amb}})} \delta(x - x_l) H(t - t_l)$$

$$\partial_t q_2 + \partial_x \left( \frac{q_2^2}{q_1} + p(\rho) A \right) = -\frac{1}{2d} \frac{q_2^2}{q_1} f_f(q)$$

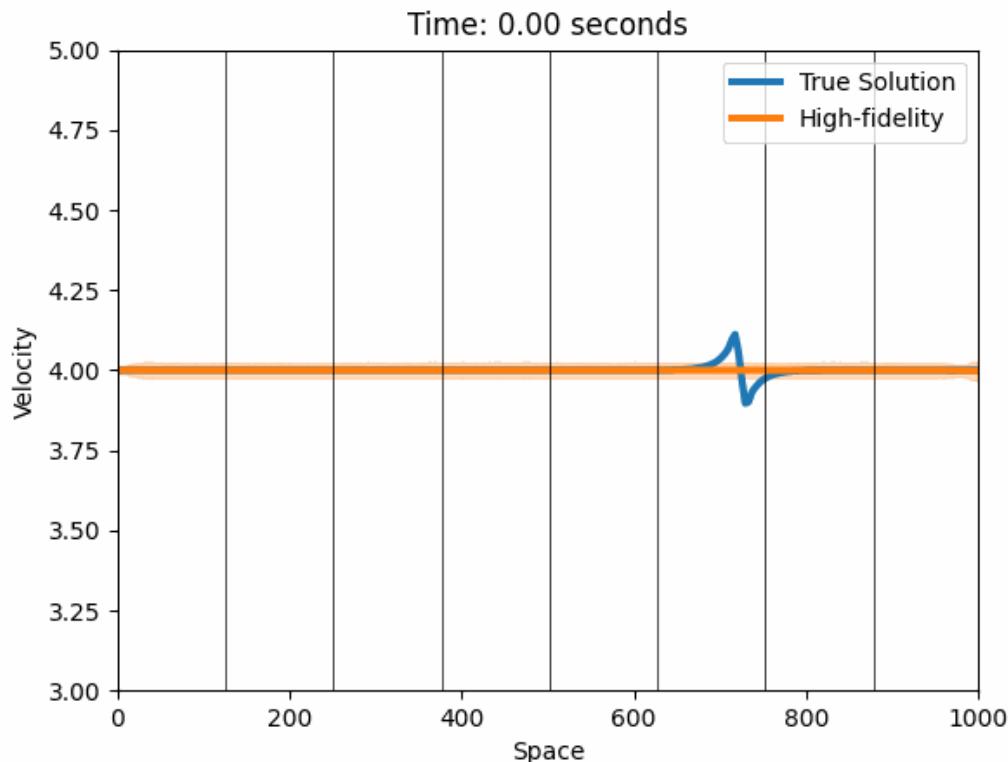
- Goal:

- State estimation
- Leak localization

- Observations:

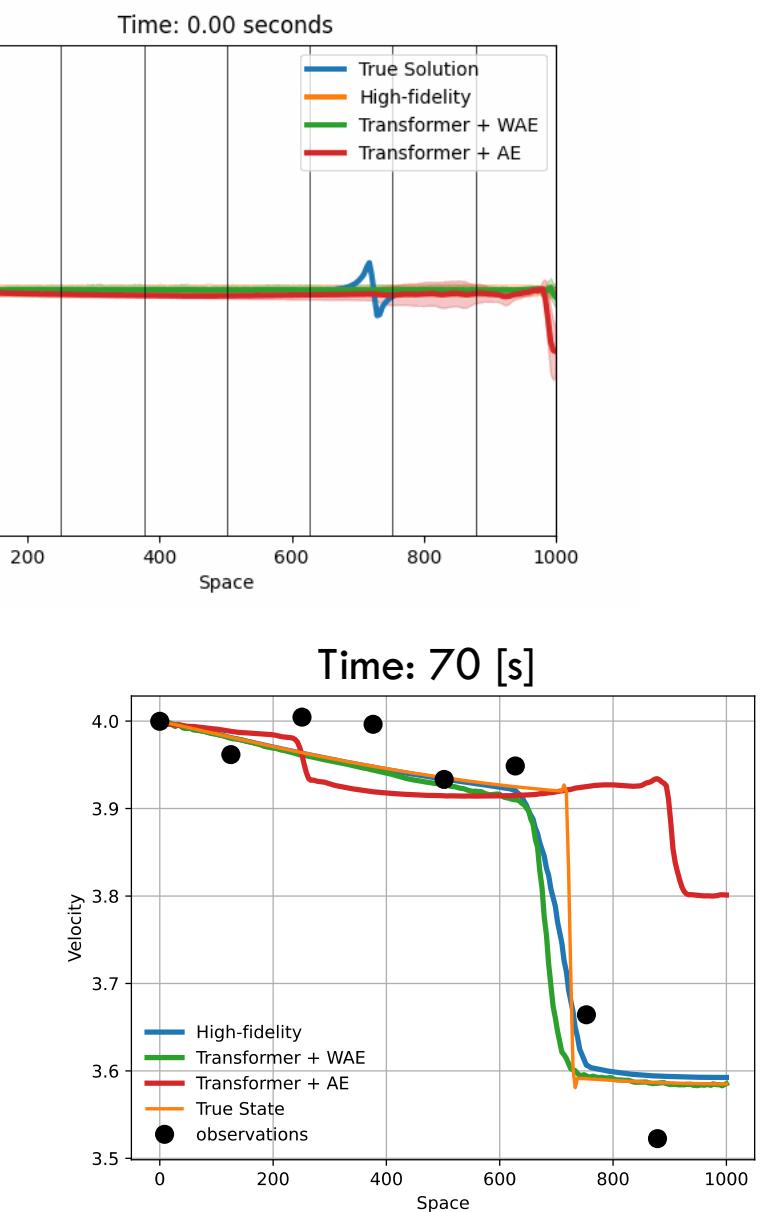
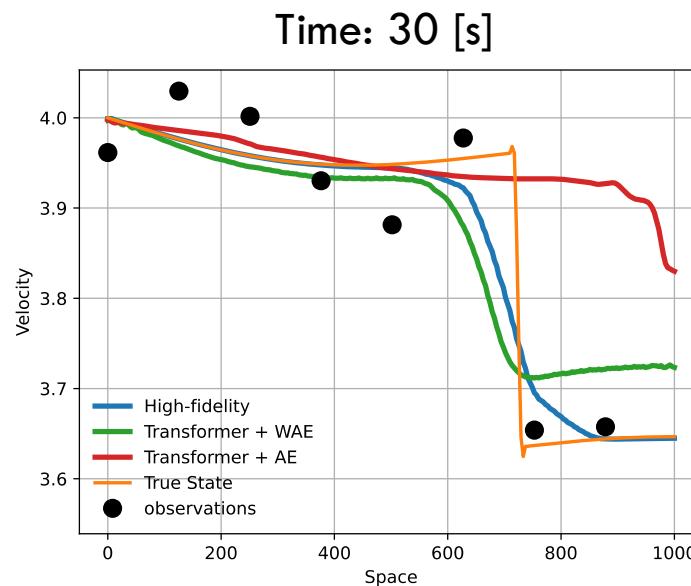
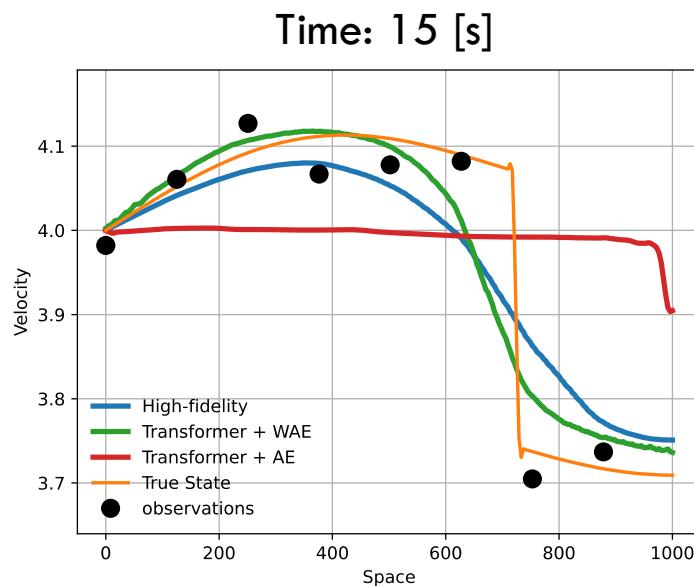
- Only velocity
- 2 seconds (40 time steps)
- 125m

- High-fidelity solution
- 10000 particles



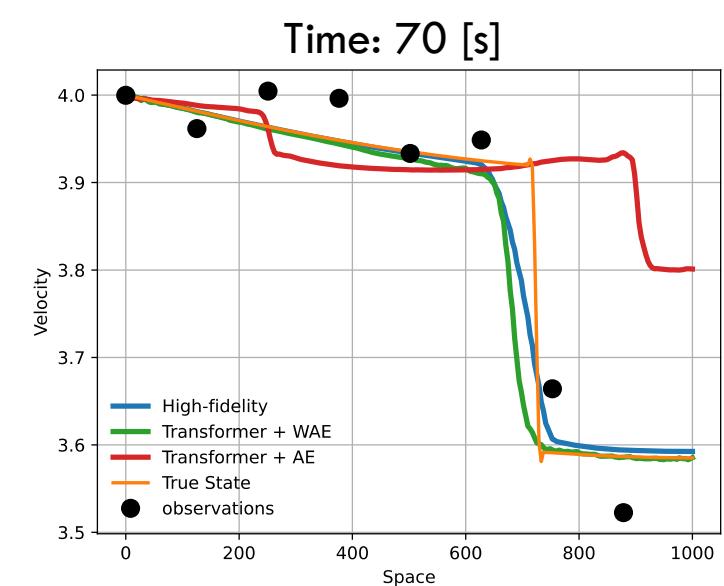
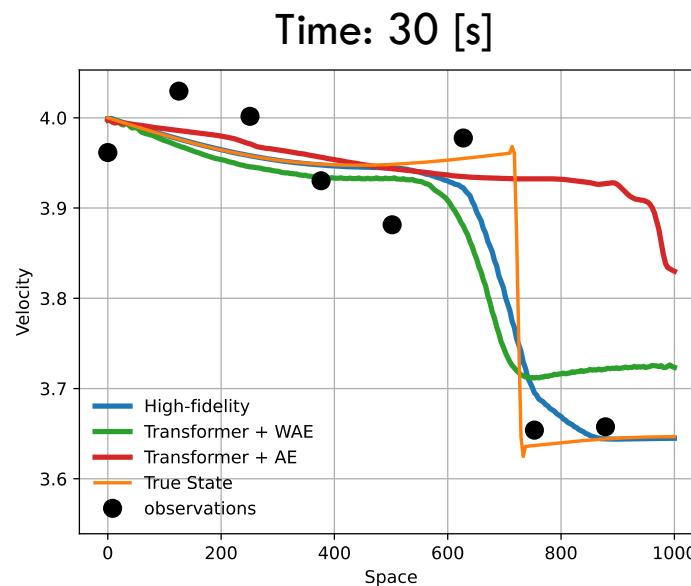
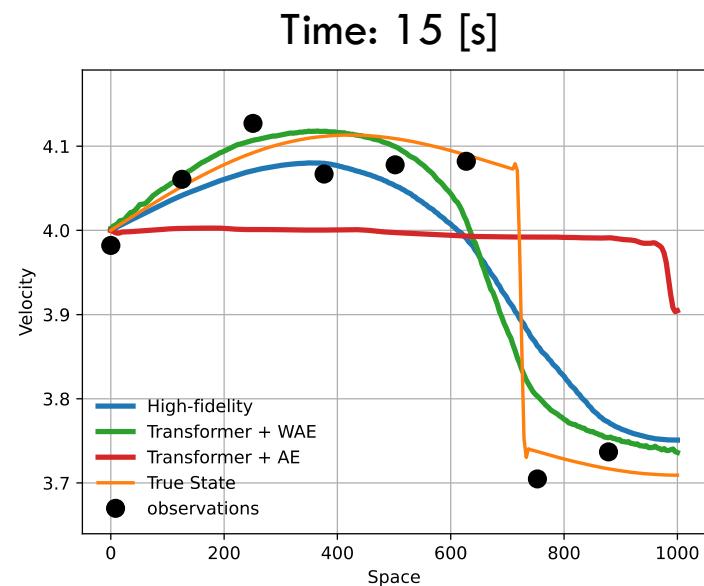
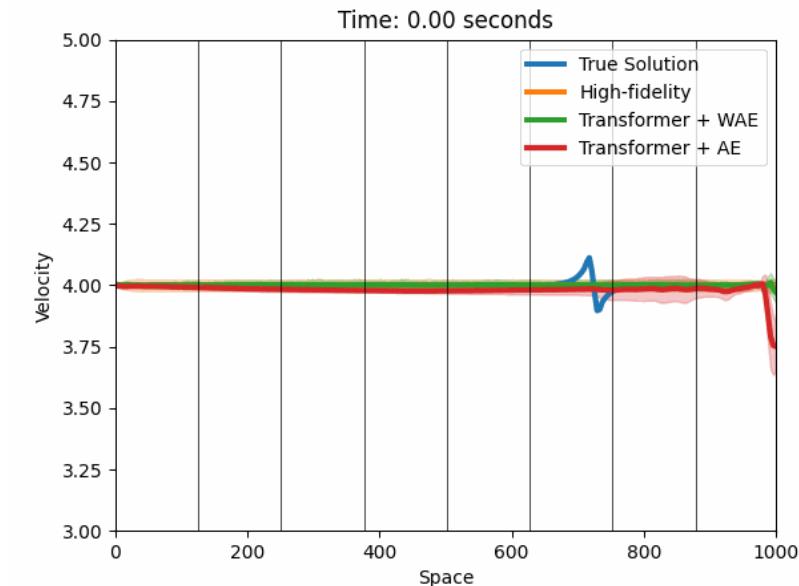
# STATE ESTIMATION

- Latent dimension = 16
- WAE and AE architectures are identical



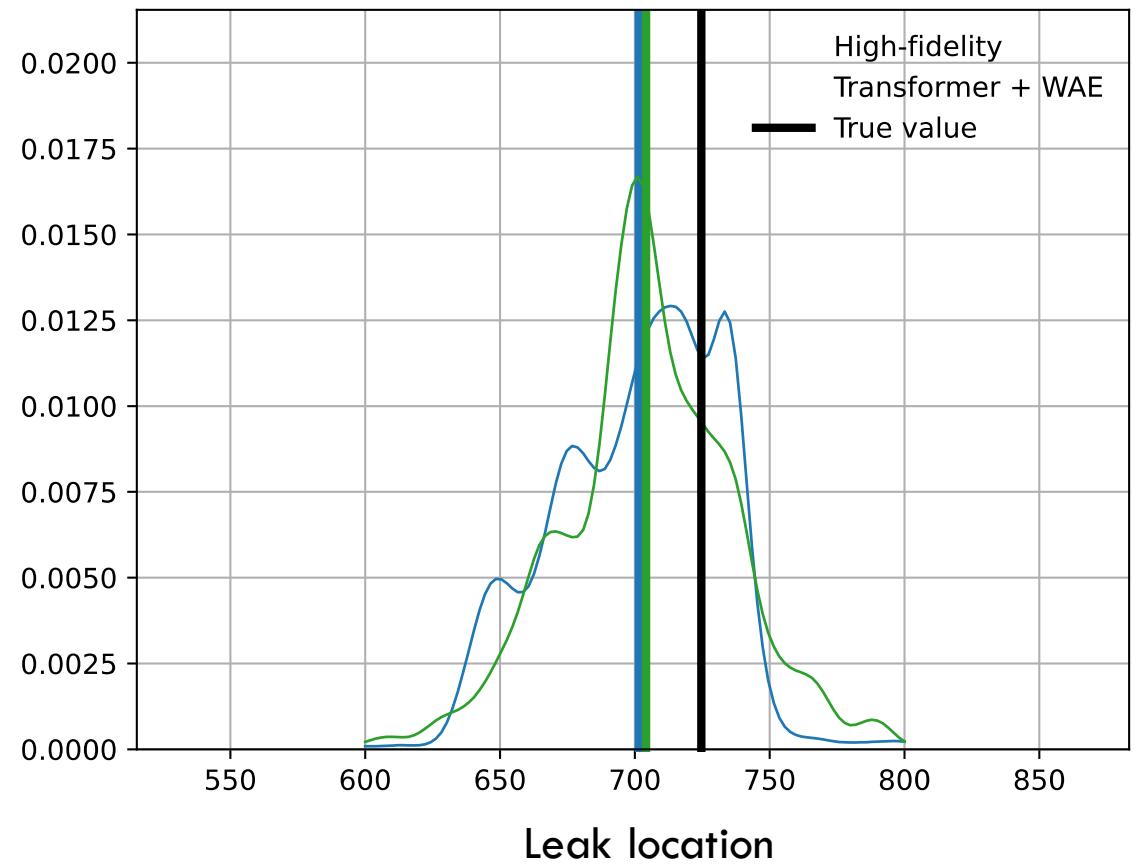
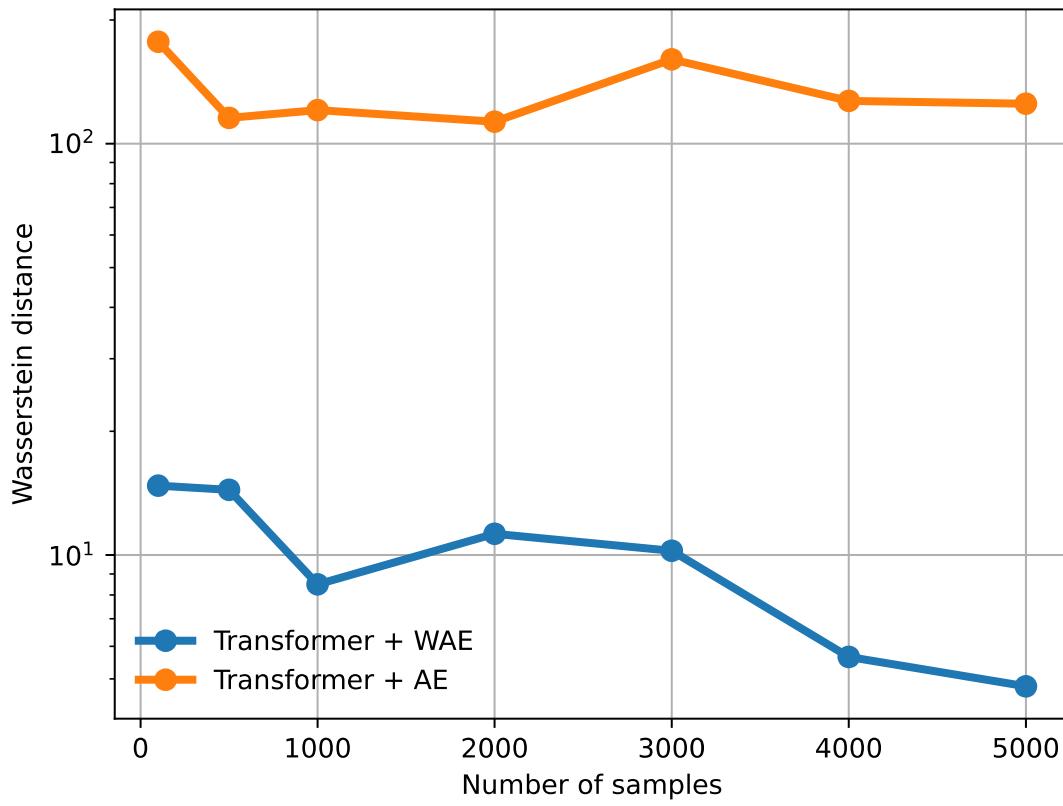
# COMPUTATION TIME

- High-fidelity(30 CPU cores): 18 hours = **64,800** seconds
- Latent:
  - 1 CPU core: **227** seconds
  - 1 GPU: **85** seconds

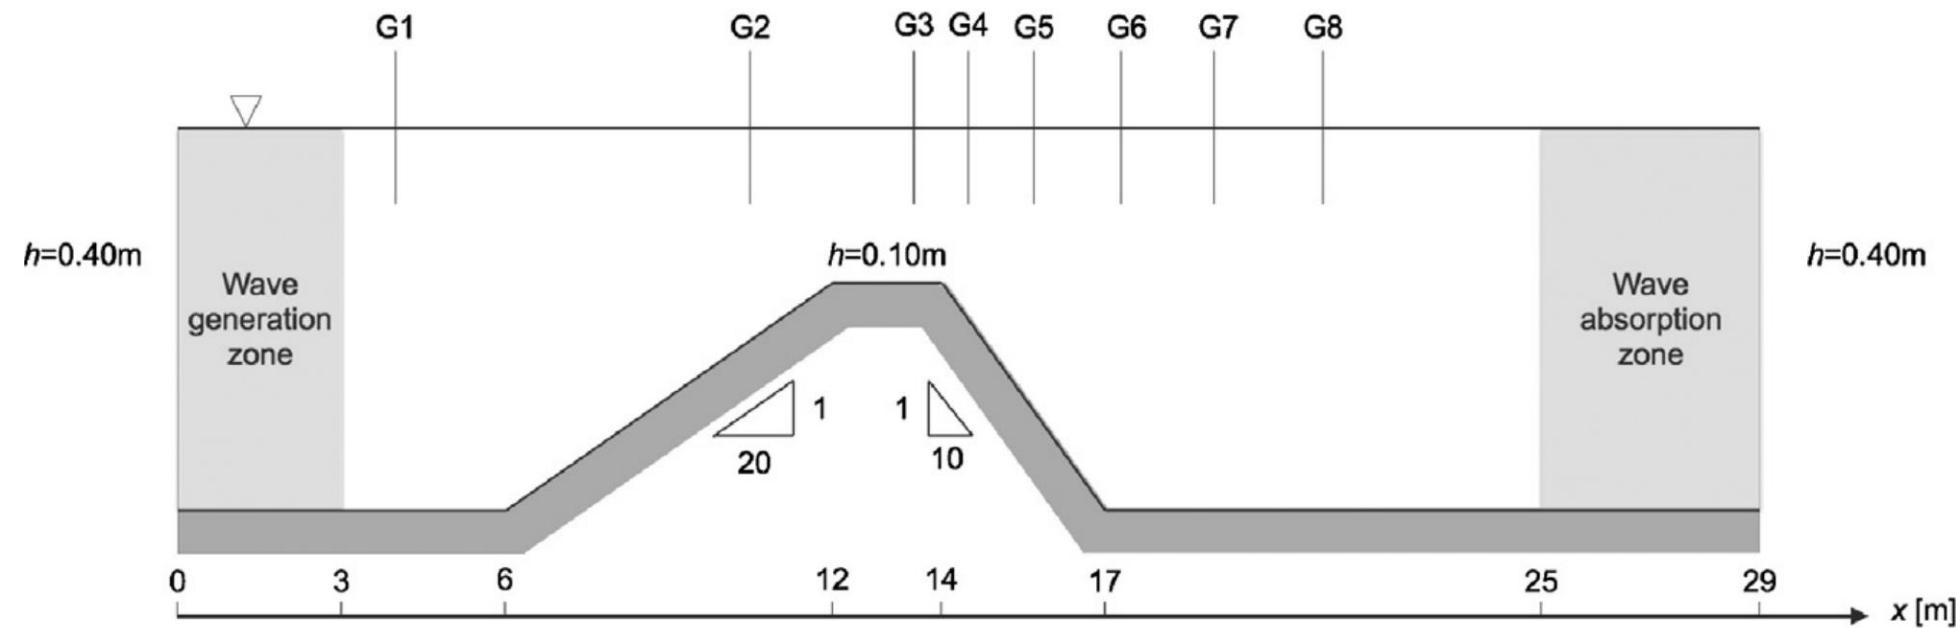




# LEAK LOCALIZATION



# HARMONIC WAVE GENERATION OVER A SUBMERGED BAR



Engsig-Karup, Allan Peter, Claes Eskilsson, and Daniele Bigoni. "A stabilised nodal spectral element method for fully nonlinear water waves." *Journal of Computational Physics* 318 (2016): 1-21.

Beji, S., and J. A. Battjes. "Experimental investigation of wave propagation over a bar." *Coastal engineering* 19.1-2 (1993): 151-162.

# HARMONIC WAVE GENERATION OVER A SUBMERGED BAR

- Fully nonlinear water wave model
- Time horizon: 42s
- Parameter: Bar height
- Surrogate model state:
  - Surface elevation
  - Surface velocity potential

Surface equations

$$\begin{aligned}\frac{\partial \eta}{\partial t} &= -\nabla \eta \cdot \nabla \tilde{\phi} + \tilde{w}(1 + \nabla \eta \cdot \nabla \eta), \\ \frac{\partial \tilde{\phi}}{\partial t} &= -g\eta - \frac{1}{2} \left( \nabla \tilde{\phi} \cdot \nabla \tilde{\phi} - \tilde{w}^2(1 + \nabla \eta \cdot \nabla \eta) \right).\end{aligned}$$

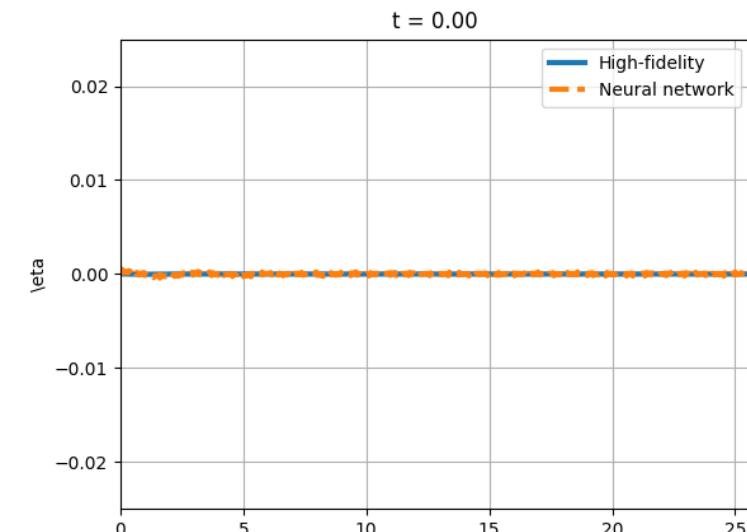
Velocity potential in full domain

$$\begin{aligned}\nabla^\sigma(K(x; t)\nabla^\sigma\phi) &= 0, \quad \text{in } \Omega^c, \\ \phi &= \tilde{\phi}, \quad z = \eta \quad \text{on } \Gamma^{FS} \\ \mathbf{n} \cdot \nabla\phi &= 0, \quad z = -h(x, y) \quad \text{on } \Gamma^b.\end{aligned}$$

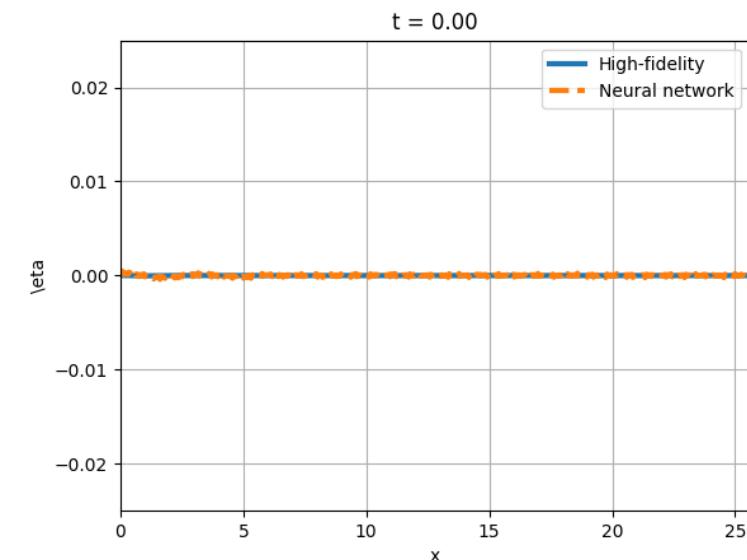
# HARMONIC WAVE GENERATION OVER A SUBMERGED BAR

- Fully nonlinear water wave model
- Time horizon: 42s
- Parameter: Bar height
- Surrogate model state:
  - Surface elevation
  - Surface velocity potential

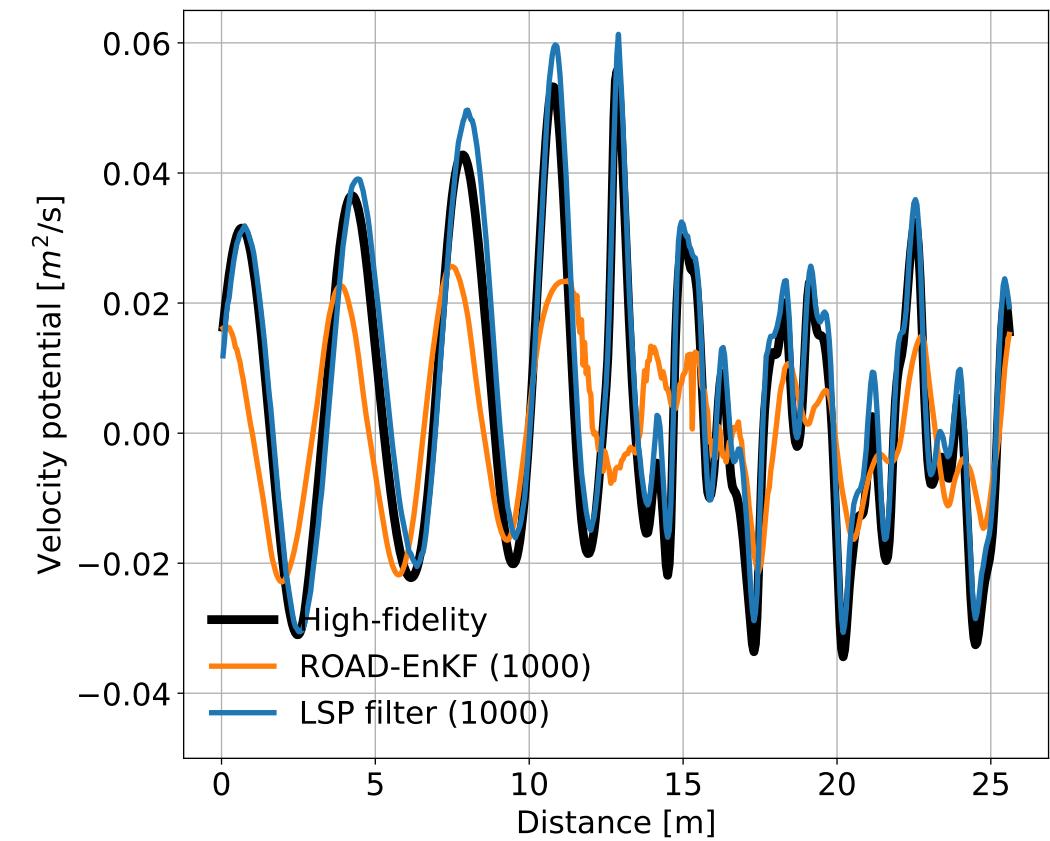
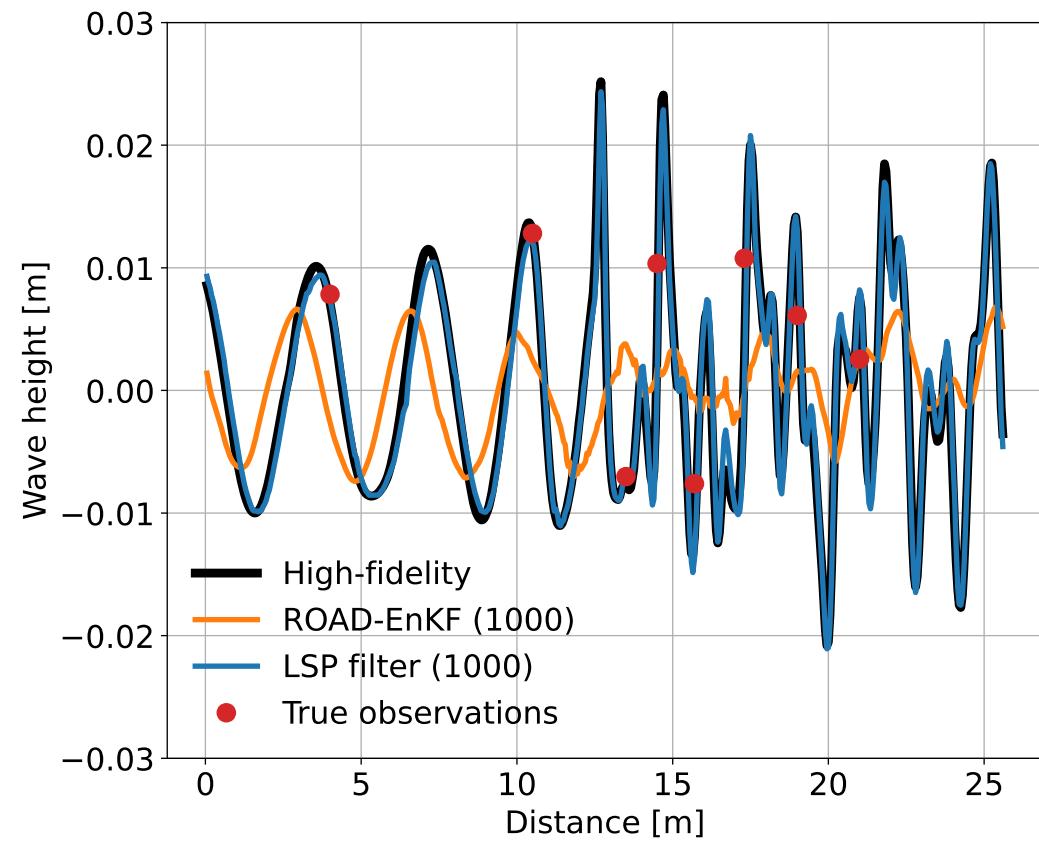
Tall bar



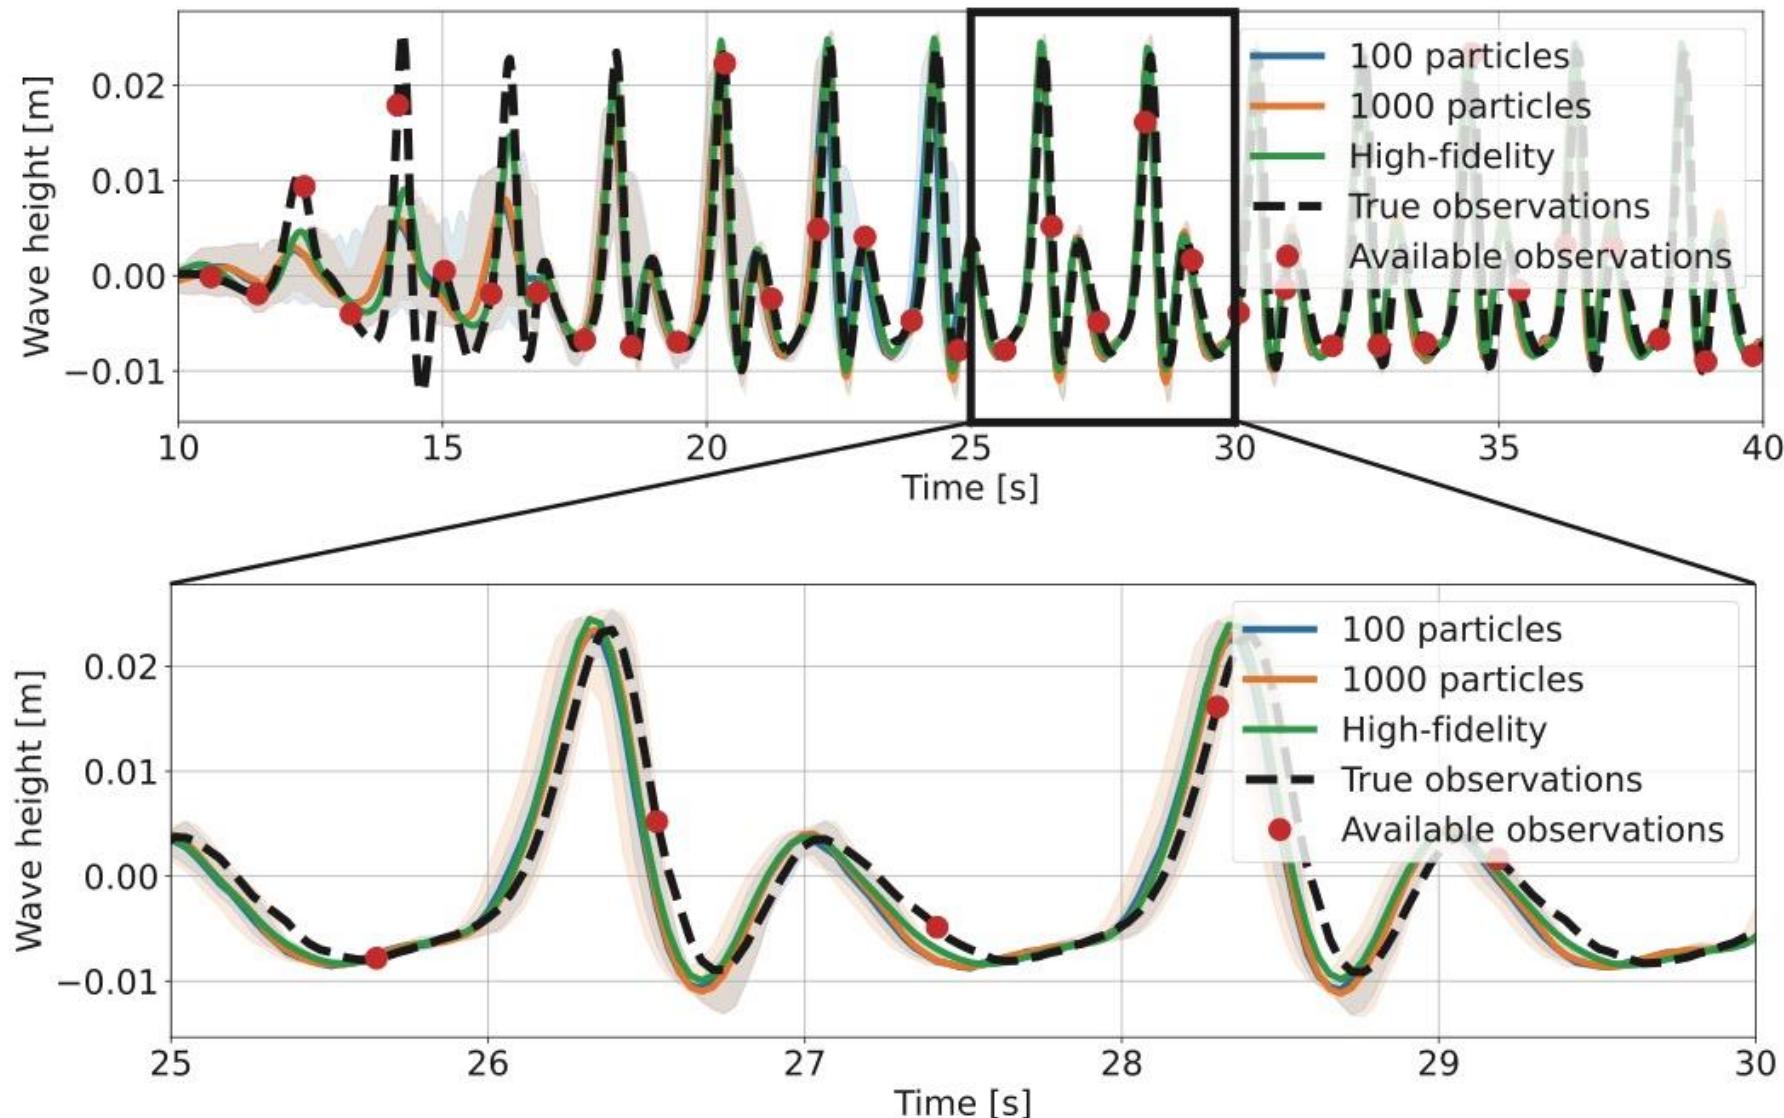
Low bar



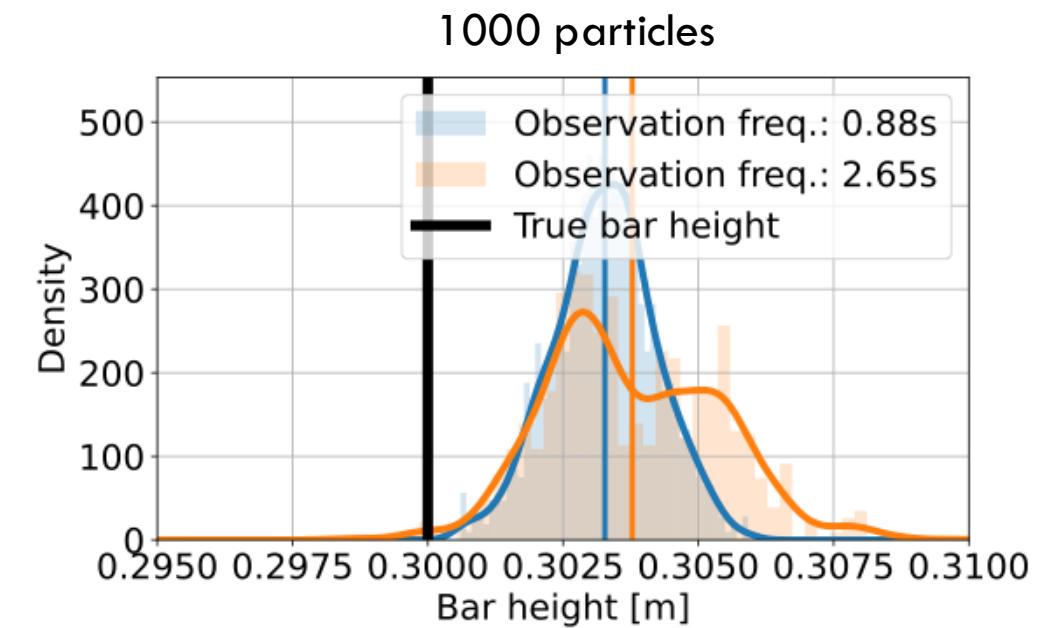
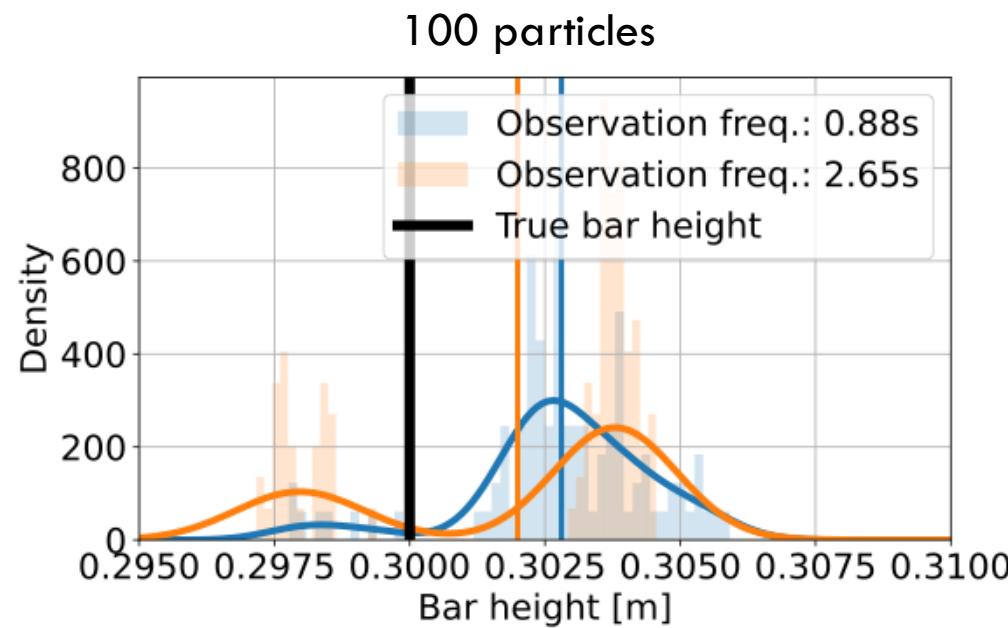
# STATE ESTIMATION



# OBSERVATIONS AT SENSOR



# BAR HEIGHT DISTRIBUTION





# CONCLUSION

- Latent space sampling can speed up Bayesian inference
- Works for static and dynamic problems

---

WANT TO KNOW MORE?

[nikolaj.mucke@cwi.nl](mailto:nikolaj.mucke@cwi.nl)

Mücke, N. T., Sanderse, B., Bohté, S. M., & Oosterlee, C. W. (2023). Markov chain generative adversarial neural networks for solving Bayesian inverse problems in physics applications. *Computers & Mathematics with Applications*, 147, 278-299.

Mücke, N. T., Bohté, S. M., & Oosterlee, C. W. (2024). The deep latent space particle filter for real-time data assimilation with uncertainty quantification. *Scientific Reports*, 14(1), 19447.